



CS 410/510

Languages & Low-Level Programming

Mark P Jones
Portland State University

Spring 2016

Week I: Introduction, Assembly Language


1

Introduction and Goals

2

Origins

- For a long time, a group of us at PSU have been looking at the role that high-level programming languages can play in the construction of (very) low-level software.
- By using **high-level languages**, we can hope to increase programmer productivity, and improve software quality
- By focussing on very **low-level software**, we hope to provide strong foundations for the complete software stack



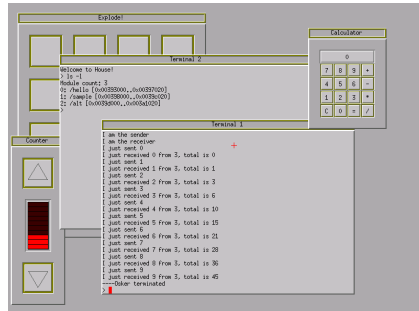
3

House (2005)

Kernel, GUI, drivers, network stack, and apps

Boots and runs in a bare metal environment

... all written in Haskell, a “purely functional” programming language



4

Why “House”?

“The Haskell User’s Operating System Environment”

You are more secure in a house ...



than if you only have Windows ...

5


Performance concerns

- By design, higher-level languages abstract away from the details of how the underlying machine works
- Can we obtain the levels of performance and predictability that are typically required/expected in the systems programming domain?
- Can we write good systems software in a language that intentionally distances users from details of memory layout, representation, instruction selection, alignment, caching, etc.?
- Traditional approaches to building system software resort to using old, low-level languages like assembly and C
- Do “modern” languages have anything to offer in this area?

6


The Habit programming language

- “a dialect of Haskell that is designed to meet the needs of high assurance systems programming”
- How do you design a programming language for a specific domain?
- Experiment with existing languages
- Understand the domain ...



The seL4 experience

- In 2009, a group from NICTA, UNSW, and OK Labs in Australia announced seL4, as “the world's first operating system kernel with an end-to-end proof of implementation correctness and security enforcement.”



- A landmark achievement for formal verification, and a strong foundation for building trustworthy systems

8

seL4 and capabilities

- Even without the verification result, the design of seL4 is interesting in its own right:
 - seL4 is a “capability enhanced” version of an earlier microkernel design called L4
 - The “capability” abstraction in seL4 provides facilities for implementing “least privilege” security policies and novel mechanisms for controlling resource usage

9

Proof vs performance?

- Security properties established in the seL4 verification include:
 - Absence of code injection attacks
 - Absence of buffer overflows
 - Absence of null pointer dereferences
 - ...
- Many of these properties could be established “for free” if the implementation had been written in a “safer” language
- How might things be different if we built something like seL4 in Habit?

10

The CEMLaBS project

- “Using a Capability-Enhanced Microkernel as a Testbed for Language-Based Security”
- Started October 2014, funded by The National Science Foundation
- Three main questions:
 - **Feasibility:** Is it possible to build an inherently “unsafe” system like seL4 in a “safe” language like Habit?
 - **Benefit:** What benefits might this have, for example, in reducing verification costs?
 - **Performance:** Is it possible to meet reasonable performance goals for this kind of system?

11

CEMLaBS workplan

- The original plan was to use Habit to design and build HAL4, a “capability-enhanced” microkernel in the style of seL4 ...
- At the time, neither the specification or the implementation of seL4 was publicly available ... but some key aspects of its design had been revealed in published papers, dissertations, and talks
- In July 2014, seL4 was released as open source software!
- Now we are free to build HAL4 as an implementation of seL4, which will allow use to make more direct comparisons between the two implementations ...

12

CEMLaBS status

- The original HAL4 implementation is incomplete (and will need to be reworked to match the seL4 specification)
- The Habit compiler is not complete (a full compiler pipeline has been implemented, but it is not feature complete, and it is not robust)
- We'll be working on addressing both of these shortcomings this coming summer
- In the meantime, there is still much to study ...

13

Course description

- **An overview of conventional low-level programming techniques:**
 - Bare metal programming
 - Fundamental programmable hardware components
- **Case studies of practical microkernel implementations:**
 - OS abstractions (address spaces, threads, capabilities, ...)
 - The L4 and seL4 microkernels
- **Reflections on the design of programming languages for use in this application domain:**
 - Assembly, C, Rust, Habit, domain specific languages, ...

14

Course learning objectives

Upon the successful completion of this course, students will be able to:

1. Write simple programs that can run in a bare-metal environment using low-level programming languages.
2. Discuss common challenges in low-level systems software development, including debugging in a bare-metal environment.
3. Explain how conventional operating system features (multiple address spaces, context switching, protection, etc.) motivate the desire for (and benefit from) hardware support.

15

Course learning objectives, continued

4. Develop code to configure and use programmable hardware components such as a memory management unit (MMU), interrupt controller (PIC), and interval timer (PIT).
5. Describe the key steps in a typical boot process, including the role of a bootloader.
6. Describe the motivation, implementation, and application of microkernel abstractions for managing address spaces, threads, and interprocess communication (IPC).
7. Explain the use and implementation of capabilities in access control and resource management.
8. Develop programs using the capability abstraction provided by the seL4 microkernel.

16

Course learning objectives, continued

9. Illustrate the use of a range of domain specific languages in the development of systems software.
10. Use practical case studies to evaluate and compare language design proposals.
11. Describe features of modern, high-level programming languages—including abstract datatypes and higher-order functions—and show how they can be leveraged in the construction of low-level software.
12. Explain how the requirements of low-level systems programming motivate the desire for (and benefit from) language-based support.

17

The “programming languages” perspective

- We will survey and evaluate a range of programming languages during this course:
 - Low-level machine and assembly languages
 - Systems programming languages (e.g., C, Rust, ...)
 - Object-oriented languages (e.g., the seL4 API)
 - Domain specific languages
 - Functional languages (e.g., Habit, Haskell, ...)
- What are the driving needs of the systems domain?
- How can a programming language design best meet those needs?

18

Context

- Basic Platform: Generic “IBM PC” compatible
 - 32 bits ... not 64
 - IA32 ... not x86_64 or ARM
 - BIOS ... not EFI or UEFI
 - `int` and `iret` ... not `sysenter`/`sysexit`
 - PIC ... not APIC
 - No PAE, PCI, ACPI, MMX, SSE, SMM, SMP, VTx, ...
 - etc., ...
- Already complicated enough for our purposes!
- Well supported by current hardware, emulators, and tools
- Underlying concepts still very broadly applicable

19

Development environment

- Ubuntu Linux
 - Week 1: using the LinuxLab machines (Mac OS also an option)
 - Weeks 2+: using a VirtualBox virtual machine, preconfigured with appropriate development tools (can be used on Linux, Mac OS, or Windows)
- Bare metal emulation using the QEMU emulator

20

Rough schedule

Week	Topic
1	Assembly language programming
2	Bare metal programming
3	Hardware support for OS abstractions
4	Memory management & protection
5	Case Study: L4 use & implementation
6	
7	Case Study 2: seL4 use & implementation
8	
9	Language design for low-level programming
10	

21

Lectures and labs

- Monday Lectures (UTS 210), Wednesday Labs (FAB 88-09)
 - Students are expected to attend all lectures and labs
 - Sets of guided practical exercises will be provided for each lab
 - Students will work on labs during lab sessions with instructor supervision
 - Students are expected to complete each set of exercises before the following week's lab session
 - We will use D2L forums for discussion of exercises outside lab sessions
 - Lab sessions may also be used for direct instruction, as necessary

22

Exams

- No midterm or final exams
 - Students are still required to attend the “final” on 6/9
 - This will be used for student presentations or demos, or else for other instruction

23

Portfolio assessment

- Final grades will be assigned on the basis of a “Portfolio” comprising the results of lab work and independent projects
- Based on previously listed collection of 12 course objectives
- General process:
 - Student uploads work item (typically a tar ball or zip file), accompanied by a write up explaining how the submitted work addresses the expectations of one or more objectives
 - Grading assigns a score of N/A, Below, Meets or Exceeds for each objective (with +/- versions)
 - Different final letter grades awarded depending on the combination of individual scores

24

Portfolio assessment, continued

- Students may request consideration of at most two objectives in any given week (Saturday-Friday)
 - discourages procrastination, balances grading load
- I will attempt to provide feedback on submitted items within a week
- Students may resubmit in an attempt to improve score for any given objective
 - (Each resubmission counts as a new submission)
- Completion of lab exercises for a given week can typically be used to satisfy the “Meets” condition for specific objectives
- Process starts in Week 2, runs until end of finals week

25

General approach and caveats

- This course will be experimental, open-ended, hands-on, and interactive
- Your flexibility, patience, and tolerance of informal or underspecified aspects of the course will be very much appreciated
- Questions, suggestions, and contributions are welcome at any point

26

An introduction to IA32 assembly language programming

27

What is IA32?

- We'll be using the IA32 (x86) architecture as our main target:
 - A "32-bit" instruction set
 - Broadly adopted by:
 - processors from Intel, AMD, Via, ...
 - laptops, desktops, servers, gaming consoles, ...
 - Linux, Mac OS X, Windows, ...
 - Arguably, a bit dated ... but still very relevant, and a good platform for learning and exploration
 - (... and one of the two architectures supported by sel4)

28

Other architectures:

- Not to be confused with:
 - x86-64/AMD64: a 64 bit architecture supported (in addition to IA32) by more recent AMD/Intel designs
 - IA64: a completely different 64-bit Intel architecture (Itanium)
 - ARM: widely used in phones, tablets, and more
 - IBM Power: used in Xbox 360, PS3, Wii, servers, and more
 - SPARC: used by some of the college's Unix servers
- Except for x86-64, you can't run IA32 code directly on a machine that uses one of these alternatives instruction sets!

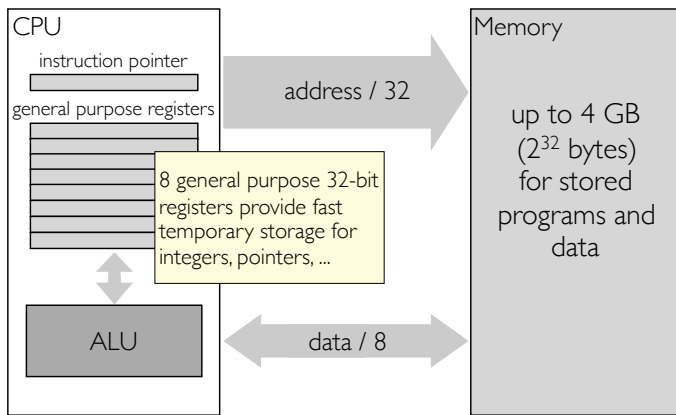
29

Notes

- No prior or in-depth knowledge of IA32 programming will be assumed
- We will only use a small subset of the full instruction set
- If you're looking to become an expert on IA32 programming, you'll want to look for another class!
- We'll be using the *AT&T syntax* for IA32 assembly language rather than the *Intel syntax*. This is the default syntax used by the free GNU tools in Linux, MacOS, and DJGPP or Cygwin on Windows, and others
- For simplicity, I recommend: [ssh yourid@linuxlab.cs.pdx.edu](ssh://yourid@linuxlab.cs.pdx.edu) (or, on Windows, the equivalent using PuTTY)

30

A (greatly) simplified view of the IA32



31

Programming for IA32

- In concrete terms, an IA32 program is just a collection of byte values (*machine code*)
- Once it has been loaded in to memory, the processor can *execute* a program by interpreting the byte values as *instructions* for the processor to act on
- For practical purposes, we will usually write IA32 programs in a textual format called *assembly language* that is easier to read than the raw byte values
- The program that translates assembly language programs in to machine code is called an *assembler*

32

The GNU assembler, as

- Assembly code goes in files with a `.s` suffix
- We will typically use `gcc` to invoke the assembler

```
gcc -m32 -o output assemblyCode.s extras.c
```
- You can also invoke the assembler directly: detailed documentation is available from:
<http://sourceware.org/binutils/docs/as/>
For IA32 programming, look in particular at the section on “80386 Dependent Features”

33

An assembly code listing

```

.global f
f:
    pushl   %ebp
    movl    %esp,%ebp
    pushl   %ebx
    movl    8(%ebp), %ebx

    movl    $0, %eax    # initialize length count in eax

loop:
    jmp     test
    incl    %eax        # increment count
    addl    $4, %ebx     # and move to next array element

test:
    movl    (%ebx), %ecx # load array element
    cmpl    $0, %ecx     # test for end of array
    jne     loop         # repeat if we're not done ...

    popl    %ebx
    movl    %ebp,%esp
    popl    %ebp
    ret
    
```

Assembly code

34

An assembly code listing

<pre> 0000 55 0001 89E5 0003 53 0004 8B5D08 0007 B8000000 0008 00 000c EB04 000e 40 000f 83C304 0012 8B0B 0014 83F900 0017 75F5 0019 5B 001a 89EC 001c 5D 001d C3 </pre>	<pre> .global f f: pushl %ebp movl %esp,%ebp pushl %ebx movl 8(%ebp), %ebx movl \$0, %eax # initialize length count in eax loop: jmp test incl %eax # increment count addl \$4, %ebx # and move to next array element test: movl (%ebx), %ecx # load array element cmpl \$0, %ecx # test for end of array jne loop # repeat if we're not done ... popl %ebx movl %ebp,%esp popl %ebp ret </pre>
---	---

Machine code

Assembly code

35

addresses

/offsets

labels

<pre> 0000 55 0001 89E5 0003 53 0004 8B5D08 0007 B8000000 0008 00 000c EB04 000e 40 000f 83C304 0012 8B0B 0014 83F900 0017 75F5 0019 5B 001a 89EC 001c 5D 001d C3 </pre>	<pre> .global f f: pushl %ebp movl %esp,%ebp pushl %ebx movl 8(%ebp), %ebx movl \$0, %eax # initialize length count in eax loop: jmp test incl %eax # increment count addl \$4, %ebx # and move to next array element test: movl (%ebx), %ecx # load array element cmpl \$0, %ecx # test for end of array jne loop # repeat if we're not done ... popl %ebx movl %ebp,%esp popl %ebp ret </pre>	<pre> .directive comments </pre>
---	---	--

machine
code


instructions

36

IA32 registers

37


8-bit registers (holding a single byte, 0-255)



Introduced in 1978 as part of the 8086 architecture

38

16-bit registers (“word”)



Introduced in 1978 as part of the 8086 architecture

39

32-bit registers (“double word”)

accumulator	_____	eax	ax	ah	al
base	_____	ebx	bx	bh	bl
count	_____	ecx	cx	ch	cl
data	_____	edx	dx	dh	dl
source index	_____	esi	si		
destination index	_____	edi	di		
base pointer	_____	ebp	bp		
stack pointer	_____	esp	sp		

“e” for extended

sometimes
referred to as
“long word”s

Introduced in 1985 as part
of the 80386 architecture

40

Special vs. general purpose registers

- `eip`: the instruction pointer register
- `esp`: the stack pointer register
- `eflags`: the flags register, stores information about the results of the most recent arithmetic or logic instruction
- Other registers can typically be used for any purpose (although some instructions—division, for example—work only with specific registers)

41

IA32 instructions

42

Instruction format

- A typical IA32 instruction has the form:



- A suffix on the opcode indicates the size of the data that is being operated on:
 - 32-bit values use the suffix **l**(ong)
 - 16-bit values use the suffix **w**(ord)
 - 8-bit values use the suffix **b**(yte)

43

Addressing modes

- **Register access**, reg:
 - `%eax`: the value in register `eax`
 - Can typically use any registers except `eip` and `eflags`
- **Memory access**, mem:
 - `var`: the value in the memory location at address `var`
 - `(%eax)`: the value in memory at the address in `eax`
 - `8 (%eax)`: the value in memory at the address given by adding 8 to the value in `eax`
- **Immediate**, immed:
 - `$42`: the constant value 42 (decimal; use `$0x2A` for hex)
 - `$var`: the address of memory location `var`

44

Directives for “declaring” variables

```
.data                # put variables in the "data" section
                    # (code usually goes in .text)

myvar: .align 4      # make sure address is multiple of 4
      .long 42       # Simple variable, initialized to 42

days: .global days  # A globally accessible array of ints
      .long 31, 28, 31, 30, 30, 30
      .long 31, 31, 30, 31, 30, 31


scratch: .space 4*100 # reserve uninitialized space

medium: .long 123     # a 32-bit integer (takes 4 bytes)
regular: .short 123   # a 16-bit integer (takes 2 bytes)
small: .byte 123      # an 8-bit integer (takes 1 byte)
```

45

How values are stored in memory

- A double word holds 32 binary digits (“bits”) (i.e., 4 bytes)



- 0xBE1A3910 can be interpreted as -1,105,577,712 (signed) or 3,189,389,584 (unsigned)
- Stored in memory with the least significant byte at the lowest address (“little endian”):

stored byte	0x10	0x39	0x1A	0xBE
address	400	401	402	403

46

IA32 instructions: data movement

47

Move instructions

- Copy data from a source to a destination (where X is one of the size suffixes: b,w,l):

`movX src, dst`

- Any of the following combinations of arguments is allowed:

`movX reg, (reg | mem)`

`movX mem, reg`

`movX imm, (reg | mem)`

- Note that you can't move mem to mem in one instruction

48

Examples

Suppose that the memory (starting at address 0) contains the following (four byte) values:

8	6	2	8	0	2	4	1	7	3	4	5	6
0	4	8	12	16	20	24	28	32	36	40	44	48


Then

instruction	contents of eax
movl \$12, %eax	12
movl (%eax), %eax	8
movl 8(%eax), %eax	0

49


Zero and sign-extension

- Suppose we want to copy a value from a 16-bit register in to a 32-bit register



- Two common strategies:

- Zero extension: for unsigned values



- Sign extension: for signed values



50

Move with sign, move with zero extension

- Copy from source to larger destination with sign extension:

`movsFT src, dst`

- Copy from source to larger destination with zero extension:

`movzFT src, dst`

- F and T are the “from” and “to” sizes (either b, w, or l)

- Valid combinations: bw, bl, or wl

- Examples:

```
movsbw %al, %dx    # byte to word
movzwl %ax, %edx    # word to long
```

51

Scaled indexed addressing

• $[base]([reg_1], reg_2[, index])$

a memory operand whose address is the value in `reg1`,
plus the specified base constant, plus the value of `reg2`
times the index (which must be 1, 2, 4, or 8)

• Any of the parts in [...] can be omitted

• Examples:

`(eax, ebx, 4)` the `ebx`th element in the array of 32-bit words starting at the address in `eax`

`days(, ebx, 4)` the `ebx`th element in the array of 32-bit words starting at the address `days`

52

More examples

Suppose that the memory (starting at address 0) contains the following (four byte) values:

8	6	2	8	0	2	4	1	7	3	4	5	6
0	4	8	12	16	20	24	28	32	36	40	44	48

Then

instruction	eax	ebx
<code>movl \$12, %eax</code>	12	
<code>movl 8(%eax), %ebx</code>	12	2
<code>movl 12(%eax, %ebx, 4), %eax</code>	7	2

53

The lea (load effective address) instruction

• Load the address of the source operand (must be memory) to a destination (where X is one of the size suffixes: b,w,l):

`leaX src, dst`

• Can also be used to co-opt the addressing mode circuitry into performing arithmetic operations:

```
leal 4(%eax), %eax    # eax += 4
leal 1(%eax, %eax, 2), %eax # eax = 3*eax + 1
leal 1(%eax, %eax), %eax # eax = 2*eax + 1
leal 4(%eax, 8), %eax   # eax = 8*eax + 4
```

• These instructions just do an address calculation and do not attempt to read the data at that address.

54

The exchange instruction

- Exchange data between two locations

`xchgX (reg | mem), reg`

- Consider the following instructions in a high-level language:

```
int tmp = x;  
x       = y;  
y       = tmp;
```

- If `x` and `y` are held in registers, then a “clever enough” compiler can translate this code into a single `xchgl` instruction

55

The instruction pointer, `eip`

- The `eip` register holds the address of the next instruction to be executed
- As the processor reads each instruction, it increments the value in `eip` by the appropriate number of bytes to point to the following instruction
- This mechanism allows the processor to execute a sequence of instructions stored in contiguous locations in memory
- What would happen if we “move” a different value in to `eip`?

56

Jumping and labels

- We can transfer control and start executing instructions at address `addr` by using a jump instruction

`jmp addr`

- Labels can be attached to instructions in an assembly language program:

```
    jmp b  
a:   jmp c  
b:   jmp a  
c:   ...
```

- Modern, pipelined machines work well with sequences of instructions that appear in consecutive locations. Jumps can be expensive: one of the goals of an optimizing compiler is to avoid unnecessary jumps.

57

IA32 instructions: arithmetic and logic operations

58

Arithmetic instructions

- Combine a given `src` with a given `dst` value and leave the result in `dst`:

```
► addX  src, dst }
  subX  src, dst } integer arithmetic
  imulX src, dst } (signed)
  andX  src, dst }
  orX   src, dst } bitwise arithmetic
  xorX  src, dst }
```

- Similar to `dst += src`, `dst -= src`, etc.. in C/C++

59

Examples

- To compute $x^2 + y^2$ and store the result in `z`:

```
► movl x, %eax
  imull %eax, %eax
  movl y, %ebx
  imull %ebx, %ebx
  addl  %ebx, %eax
  movl  %eax, z
```

register	contents
eax	$x^2 + y^2$
ebx	y^2

```
    .data
x:  .long 4
y:  .long 3
z:  .long 0
```


60

IA32 instructions: conditional execution

61

Flags

- In addition to performing the required operation, arithmetic instructions also change bits in the `eflags` register



- The flags record details about the last operation, such as:
 - Was the result zero?
 - Was the result positive?
 - Did a carry occur?
 - etc...

62

Conditional jumps, `jcc`

We can test these flags in *conditional jump* instructions

<code>jz addr</code>	(jump to <code>addr</code> if the zero flag is set)	
<code>jnz addr</code>	(jump to <code>addr</code> if the zero flag is not set)	
<code>je addr</code>	(jump to <code>addr</code> if equal; same as <code>jz</code>)	
<code>jne addr</code>	(jump to <code>addr</code> if not equal; same as <code>jnz</code>)	
<code>jl addr</code>	(jump to <code>addr</code> if less than)	} (signed)
<code>jnl addr</code>	(jump to <code>addr</code> if not less than)	
<code>jg addr</code>	(jump to <code>addr</code> if greater than)	
<code>jng addr</code>	(jump to <code>addr</code> if not greater than)	
...		

63

Examples

```
subl    %eax,%ebx  
jz      addr
```

jump to addr
if ebx = eax

```
subl    %eax,%ebx  
jnz     addr
```

jump to addr
if ebx ≠ eax

```
subl    %eax,%ebx  
jl      addr
```

jump to addr
if ebx < eax

```
subl    %eax,%ebx  
jnl     addr
```

jump to addr
if ebx ≥ eax

If the specified condition does not apply, then execution just continues with the next instruction ...

64

The compare instruction

- The `cmpX` instruction behaves like `subX` except that the result is not saved; only the flags are changed
- For example:

```
cmpl    %eax,%ebx  
jl      addr
```

will jump to `addr` if the value in `ebx` is less than the value in `eax`, but it will not change the values in either register

65

Other conditional instructions

- There are some other instructions that perform an action based on the conditional flags without the cost of a jump
- `setCC reg8` sets the value in a specified 8-bit register to 0 or 1, based on the condition specified by CC:

```
cmpl    %ecx,%ebx    # set eax to 1 if  
setl    %al          # ebx < ecx, or  
movzbl  %al,%eax     # else to 0
```

- `cmoveCC src, dst` copies data from the specified `src` to `dst`, but only if the condition specified by CC holds:

```
cmpl    %ebx,%eax    # set eax to the max of  
cmovl   %ebx,%eax    # eax and ebx
```

condition code; no size suffix here!

66

IA32 instructions: more arithmetic

67

Unary operations

- The following arithmetic operations have only one argument (which serves as both source and destination)

► negX	(reg mem)	negate
notX	(reg mem)	complement
incX	(reg mem)	increment
decX	(reg mem)	decrement

- Like the binary operators, these instructions also set the flags for subsequent testing

68

Bitwise shift operations

- Shift operations are handled using instructions of the form:

op count, (reg | mem)

cf ← **shl/sal** ← 0 shift (logical/arithmetic) left

0 → **shr** → **cf** shift logical right

→ **sar** → **cf** shift arithmetic right

- count is either a constant or else the %cl register.
- In all cases, the count value will be masked to 31 bits.

69

Example


- Given two 32 bit input values:

• base:

• limit:

(Each box is one nibble (4 bits),
least significant bits on the right)

- Calculate a 64 bit descriptor:



- (Needed for the calculation of “GDT entries”)

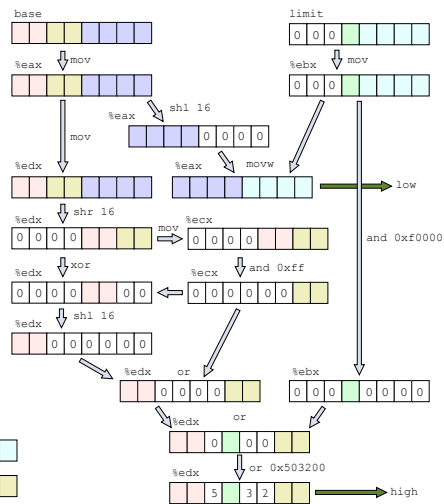

70

Example

```
movl    base, %eax
movl    limit, %ebx

mov     %eax, %edx
shl     $16, %eax
mov     %bx, %ax
movl    %eax, low

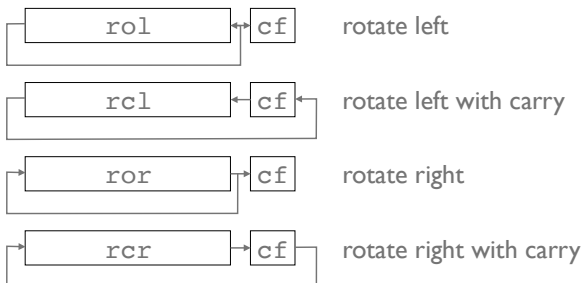
shr     $16, %edx
mov     %edx, %ecx
andl    $0xff, %ecx
xorl    %ecx, %edx
shl     $16, %edx
orl     %ecx, %edx
andl    $0xf0000, %ebx
orl     %ebx, %edx
orl     $0x503200, %edx
movl    %edx, high
```



71

Bitwise rotate operations

- Rotate operations use the same instruction format:



- [Aside: Curiously, “higher level” languages often include shift operators, but not rotates, even though the latter have more interesting/uniform behavior ...]


72

Division

- Divide implicit destination (`edx:eax`) (a 64-bit quantity) by a specified argument with result in `eax` and remainder in `edx`

`idivl (reg | mem)`

- Often used in conjunction with the `cld` instruction (“convert long to double”, a.k.a. `cdq`), which converts a signed 32-bit value in `eax` into the corresponding signed 64-bit value in `edx:eax`.



73

Example 1

Divide 4,660 (i.e., 0x1234) by 25:

```
► movl    $0x1234, %eax
   cld
   movl    $25, %ecx
   idivl   %ecx
```

Results: `eax` = 0xBA (186)
`edx` = 0xA (10)

Sure enough: $186 \cdot 25 + 10 = 4,660$

74

Example 2

Divide -1,105,577,712 (i.e., 0xBE1A3910) by 256

```
► movl    $0xBE1A3910, %eax
   cld
   movl    $256, %ecx
   idivl   %ecx
```

Results: `eax` = 0xFFBE1A3A (-4,318,662)
`edx` = 0xffffffff10 (-240)

Sure enough: $-4,318,662 \cdot 256 - 240 = -1,105,577,712$

75

Complications of division

- Division produces multiple results: a quotient and a remainder
- Division uses special registers: we'd better not store any other values in `eax` or `edx` if there's a chance that a division instruction might be executed
- Doesn't set flags: requires separate tests, for example, to determine whether quotient or remainder was zero
- Division can raise an exception if the `src` is zero (or `-1`)

76

IA32 instructions: using the stack

77

Stack

- The IA32 includes features that allow the programmer to use a region of memory as a simple stack:
 - the `esp` (stack pointer) register
 - special instructions like `push`, `pop`, `call`, `ret`, ...
- There is no obligation for the programmer to use these features, but it is often convenient to do so:
 - for temporary/scratch storage when a calculation needs more storage than the CPU registers can provide
 - to support calling and returning from functions

78

A typical memory layout

- A typical operating system reserves an area of scratch memory for each program, and sets the `esp` register to point to the end of this region when the program begins



- The stack pointer moves
 - down (decreases) as values are pushed on to the stack
 - up (increases) as values are popped off of the stack
- So long as they never overlap, the data and stack areas can grow or shrink as necessary as the program runs

79

Stack operations

- Push a value onto the stack

`pushl (reg | mem | imm)`

- Pop a value of the stack

`popl (reg | mem)`

- Roughly speaking:

```
pushl src  =  subl $4, %esp;  movl src, (%esp)
popl dst   =  movl (%esp), dst;  addl $4, %esp
```

80

Spilling temporaries on the stack

- The stack is often used for saving the contents of a register on the stack ("spilling") so that the register can be used, temporarily, for some other reason

- For example:

```
pushl %eax
pushl %edx
...code that changes eax and/or edx ...
popl  %edx
popl  %eax
```

pop values in reverse order
that was used to push them!

- Note that values on the stack can still be accessed, from memory, using `(%esp)`, `4(%esp)`, `8(%esp)`, `12(%esp)`, ...

81

Call and return

- There is a special instruction for calling a function

```
call addr    ≈    pushl $lab
                jmp  addr
lab: ...
```

- And a special instruction for returning from a function

```
ret          ≈    popl  %eax ← assuming
                    jmp  *%eax  eax isn't being
                                used for
                                something else
                                ...
```

- In practice, additional instructions are often needed to deal with parameter passing, etc. ...

special syntax: jump
to the address given
by the contents of
eax

82

Functions and the System V ABI

83

Implementing functions


- How do we pass arguments to a function?
- How does a function return a result?
- How do we handle local variables?
- In principle, especially in a bare metal setting, we can implement these features any way we like, using the basic tools that the IA32 instruction set provides
- But there are some existing standards we can follow, notably the “System V IA32 Application Binary Interface (ABI)”:

<http://www.sco.com/developers/devspecs/abi386-4.pdf>
particularly Section 3-9

84

Stack frames

The code for any given function/procedure call runs in the context of a stack frame of the form:



- Frame (base) pointer: ebp points to the stack frame; the caller's frame pointer is stored in old (i.e., (%ebp))
- Return address: retn is the return address
- Actual parameters: a_1, \dots, a_n are the function's arguments. We can access a_1 as 8 (%ebp), etc...
- Local variables: l_1, \dots, l_m are the function's local variables. We can access l_1 as -4 (%ebp), etc...

85

Building the stack frame ... in the caller

- | | | |
|-----|-----|--|
| | ... | |
| esp | ebp | |
- The caller starts by pushing the arguments:

	a_1	...	a_n	...	
esp					ebp
- Then it executes a `call` instruction, which pushes the return address:

	retn	a_1	...	a_n	...	
esp						ebp
- ... and jumps to the code for the callee ...

86

Building the stack frame ... in the callee

- | | | | | | | |
|-----|------|-------|-----|-------|-----|-----|
| | retn | a_1 | ... | a_n | ... | |
| esp | | | | | | ebp |
- The callee saves the old frame pointer, and sets a new value:

	old	retn	a_1	...	a_n	...	
ebp=esp	4	8	12	...			
- Then it decrements the stack pointer to reserve space for any local variables:

	l_m	...	l_1	old	retn	a_1	...	a_n	...	
esp	...	-4	ebp	4	8	12	...			
- ... and now the callee can start work ...

87

Function prologue

- The code that builds the stack frame at the start of a function body is called the prologue:

- At the beginning of a function body, the parameters and return address have already been pushed on to the stack.

We need to:

```
pushl %ebp      # save old frame pointer
movl  %esp, %ebp # and set new value
```

- If local variables taking M bytes of storage are required, then we need to reserve space for them:

```
subl  $M, %esp # allocate space for
                # locals (skip if  $M=0$ )
```

88

Function epilogue

- When a function completes, we must dismantle the stack frame and return the machine to the state it was in before the call. The code to do this is called the epilogue:

- Running the previous process in reverse:

```
movl  %ebp, %esp # discard locals/temps
popl  %ebp       # restore frame pointer
ret    # return to caller
```

- The first two instructions here can be replaced with the more efficient, but otherwise equivalent `leave` instruction

89

Removing the parameters

- Once we return to the caller, the result of the function is in `eax`, but the parameters are still on the stack:



- We restore the stack pointer to its original value by adding on the number of bytes that are used by the parameters:

```
addl  $N, %esp
```

- If no parameters were passed, then this step can be omitted

90

Example: a leaf function

```
int g(int u) {  
    return u*u;  
}
```



```
g: pushl %ebp  
   movl %esp, %ebp  
   movl 8(%ebp), %eax  
   imull %eax, %eax  
   movl %ebp, %esp  
   popl %ebp  
   ret
```

	...	old	retn	u	...	
		esp=ebp	ebp+4	ebp+8		

91

Example: multiple parameters + call

```
int f(int x,  
      int y,  
      int z) {  
    return g(x+y);  
}
```



```
f: pushl %ebp  
   movl %esp, %ebp  
   movl 8(%ebp), %eax  
   addl 12(%ebp), %eax  
   pushl %eax  
   call g  
   addl $4, %esp  
   movl %ebp, %esp  
   popl %ebp  
   ret
```

	...	old	retn	x	y	z	...	
		esp=ebp	ebp+4	ebp+8	ebp+12	ebp+16		

92

Example: spilling

```
int h(int x,  
      int y,  
      int z) {  
    return g(x)+g(y);  
}
```



```
h: pushl %ebp  
   movl %esp, %ebp  
   pushl 8(%ebp)  
   call g  
   addl $4, %esp  
   pushl %eax -- spill  
   pushl 12(%ebp)  
   call g  
   addl $4, %esp  
   popl %ecx -- unspill  
   addl %ecx, %eax  
   movl %ebp, %esp  
   popl %ebp  
   ret
```

	spill	old	retn	x	y	z	
	esp	ebp	ebp+4	ebp+8	ebp+12	ebp+16	

93

Observations

- There is a four instruction overhead for each function that uses the frame pointer
 - Increases execution time
 - Prevents use of `ebp` as a general purpose register
- For larger functions, the four instruction overhead is less of an issue
- For small functions, we would prefer to inline rather than copy
- Nevertheless, it is common to produce code that doesn't use `ebp` as a frame pointer (e.g., `-fomit-frame-pointer` in `gcc`)

94

Caller and callee saves

We (System V) can designate some registers as:

- **caller saves** (`eax`, `ecx`, and `edx`)
 - can be freely used by the callee
 - the caller is responsible for saving (and later restoring) the value of a caller save register before a call
- **callee saves** (`ebp`, `ebx`, `esi`, and `edi`)
 - can be freely used by the caller
 - the callee is responsible for saving (and later restoring) the value of a callee saves register before using it to store temporary values

95

Revisiting the previous example: h

```
int h(int x,  
      int y,  
      int z) {  
    return g(x)+g(y);  
}
```



```
h: pushl %ebp  
   movl %esp, %ebp  
  
   pushl 8(%ebp) -- x  
   call g  
   addl $4,%esp  
   pushl %eax -- spill  
   pushl 12(%ebp) -- y  
   call g  
   addl $4, %esp  
   popl %ecx -- unspill  
   addl %ecx, %eax  
  
   movl %ebp, %esp  
   popl %ebp  
   ret
```

instead of having the compiler save this value on the stack ...

96

Revisiting the previous example: h

```
int h(int x,      → h: pushl %ebp
    int y,      movl %esp, %ebp
    int z) {
    return g(x)+g(y);
}
```

... we can move it to a callee saves register, esi

g will preserve the value in esi, if necessary

so it will still contain the correct value here...

```
    pushl 8(%ebp) -- x
    call g
    addl $4,%esp
    movl %eax, %esi
    pushl 12(%ebp) -- y
    call g
    addl $4,%esp
    addl %esi, %eax
    movl %ebp, %esp
    popl %ebp
    ret
```

97

Revisiting the previous example: h

```
int h(int x,      → h: pushl %ebp
    int y,      movl %esp, %ebp
    int z) {      pushl %esi
    return g(x)+g(y);
}
```

... now h has to save the value in register, esi

one save in h is better than one saves in each of two calls to g

empirically, more than 50% of calls are to leaf functions

```
    pushl 8(%ebp) -- x
    call g
    addl $4,%esp
    movl %eax, %esi
    pushl 12(%ebp) -- y
    call g
    addl $4,%esp
    addl %esi, %eax
    popl %esi
    movl %ebp, %esp
    popl %ebp
    ret
```

98

Closing thoughts

99

Assembly “Language”?

- Highly imperative, primitive instructions, no expressions
- No high-level abstractions, but all the building blocks:
 - No arrays, records, variants, objects, closures, ...
 - No loops, switch statements, functions, local variables, ...
- Type System?
 - Values classified by size (e.g., 8 vs 32 bits) and storage class (e.g., memory, flag, integer register, floating point register, ...)
 - Limited protection against common programming mistakes
 - Programmer has full control over data representation

100

Macros

macro name parameters defaults

```
.macro    idtcalc handler, slot, dpl=0, type=IDT_INTR, seg=KERN_CS
mov      $seg, %eax          # eax = ? # seg
shl      $16, %eax           # eax = seg # 0
movl     $handler, %edx      # edx = hhi # hlo
mov      %dx, %ax            # eax = seg # hlo
mov      $(0x8e00 | (\dpl<<13) | \type), %dx
movl     %eax, idt + (      8*\slot)
movl     %edx, idt + (4 + 8*\slot)
.endm
```

loop constructs

```
initIDT: .irp    num, 0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,16,17,18,19
idtcalc  exc\num, slot=\num
.endr
```

macro invocation

```
idtcalc handler=syscall, slot=0x80, dpl=3
lidt     idtptr
ret
```

Macros support compile-time code generation:

- Macros are like “used-defined instructions”
- Also supported: conditional code generation, loops, ...
- One way to compensate for language weaknesses?

101

Summary

- IA32 provides a very basic programming language:
 - A fixed set of registers
 - Instructions for moving and operating on data
 - Instructions for testing and control transfer
- In programming language terms:
 - Low-level, primitive instructions, loosely typed
 - No high-level abstractions, but all the building blocks
 - Very close to the metal, low-level control, “predictable” performance
- Let’s write some programs!

102