



CS 410/510

Languages & Low-Level Programming

Mark P Jones
Portland State University

Spring 2016

Week 6: L4 Implementation

1

Introducing “pork”

- pork = the “Portland Oregon Research Kernel”
- An implementation of (a subset of) L4 X.2
- Similar API to Pistachio, but specific to IA32 platform
- Written around the start of 2007
- “I have almost all the pieces that I need to build an L4 kernel ... perhaps I should try putting them together?”
- Built using the techniques we have seen so far in this course ...
- ... let’s take a tour!

2

Boot

3

boot.S should look very familiar ...

```
entry: .global entry
       cll                                # Turn off interrupts

#-----
# Create initial page directory:
...
#-----
# Turn on paging/protected mode execution:
...
#-----
# Initialize GDT:
...
#-----
# Initialize IDT:
...
#-----
# Initialize PIC:
...
jmp     init          # Jump off into kernel, no return!

#-----
# Halt processor: Also used as code for the idle thread.
.global halt
halt:   hit
       jmp     halt

#-----
# Data areas:
.data
...
```

4

Exception handlers

```
# Descriptors and handlers for exceptions: -----
intr   0, divideError
intr   1, debug
intr   2, nmiInterrupt
intr   3, breakpoint
intr   4, overflow

intr   5, boundRangeExceeded
intr   6, invalidOpcode
intr   7, deviceNotAvailable
intr   8, doubleFault,      err=HWERR
intr   9, coprocessorSegmentOverrun

intr  10, invalidTSS,        err=HWERR
intr  11, segmentNotPresent, err=HWERR
intr  12, stackSegmentFault, err=HWERR
intr  13, generalProtection, err=HWERR
intr  14, pageFault,         err=HWERR

// Slot 15 is Intel Reserved
intr  16, floatingPointError
intr  17, alignmentCheck,    err=HWERR
intr  18, machineCheck
intr  19, simdFloatingPointException

// Slots 20-31 are Intel Reserved
```

5

Hardware interrupt handlers

```
# Add descriptors for hardware irqs: -----
.equ   IRQ_BASE, 0x20          # lowest hw irq number

.irp   num,          0x21,0x22,0x23, 0x24,0x25,0x26,0x27, \
                   0x28,0x29,0x2a,0x2b, 0x2c,0x2d,0x2e,0x2f
intr   \num, service=hardwareIRQ, err=(\num-IRQ_BASE)
.endr

intr   0x20, timerInterrupt
```


6

System call entry points

```
# Add descriptors for system calls: -----
# These are the only idt entries that we will allow to be called
# from user mode without generating a general protection fault,
# so they will be tagged with dpl=3.
intr    INT_THREADCONTROL, threadControl,    err=NOERR, dpl=3
intr    INT_SPACECONTROL,  spaceControl,      err=NOERR, dpl=3
intr    INT_IPC,           ipc,               err=NOERR, dpl=3
intr    INT_EXCHANGEREGS,  exchangeRegisters, err=NOERR, dpl=3
intr    INT_SCHEDULE,      schedule,          err=NOERR, dpl=3
intr    INT_THREADSWITCH,  threadSwitch,      err=NOERR, dpl=3
intr    INT_UNMAP,         unmap,             err=NOERR, dpl=3
intr    INT_PROCCONTROL,   processorControl,  err=NOERR, dpl=3
intr    INT_MEMCONTROL,    memoryControl,     err=NOERR, dpl=3
intr    INT_SYSTEMCLOCK,   systemClock,       err=NOERR, dpl=3
```

7

Overall kernel structure



8

An example exception handler

```
ENTRY invalidOpcode() {
    byte* eip = (byte*)current->context.iret.eip;
    if (eip[0]==0xf0 && eip[1]==0x90) { // Check for LOCK NOP instruction
        current->context.iret.eip += 2; // found => KernelInterface syscall
        KernelInterface_SetBaseAddress = kipStart(current->space);
        KernelInterface_SetAPIVersion  = API_VERSION;
        KernelInterface_SetAPIFlags    = API_FLAGS;
        KernelInterface_SetKernelId    = KERNEL_ID;
        resume();
    }
    handleException(6);
}
```

9

The KIP

10

What's in the KIP?

~	SCHEDULE_SC	THREADSWITCH_SC	Reserved	+F0 / +1E0
EXCHANGEREGISTERS_SC	UNMAP_SC	LIPC_SC	IPC_SC	+E0 / +1C0
MEMORYCONTROL_pSC	PROCESSORCONTROL_pSC	THREADCONTROL_pSC	SPACECONTROL_pSC	+D0 / +1A0
ProcessorInfo	PageInfo	ThreadInfo	ClockInfo	+C0 / +180
ProcDescPtr	BootInfo	~		+B0 / +160
KipAreaInfo	UtcInfo	VirtualRegInfo	~	+A0 / +140
~				+90 / +120
~				+80 / +100
~				+70 / +E0
~				+60 / +C0
~		MemoryInfo	~	+50 / +A0
~				+40 / +80
~				+30 / +60
~				+20 / +40
~				+10 / +20
KernDescPtr	APIFlags	APIVersion	0 _(0/32) K ₂₃₀ L ₄ L ₁	+0

11

kip.S

```
.data
.align (1<<PAGESIZE)
.global Kip, KipEnd
Kip:
.byte 'L', '4', 230, 'K'
.long API_VERSION, API_FLAGS, (KernelDesc - Kip)

.global Sigma0Server, Sigma1Server, RootServer
Kdebug: .long 0, 0, 0, 0 # Kernel debugger information
Sigma0Server: .long 0, 0, 0, 0 # Sigma0 information
Sigma1Server: .long 0, 0, 0, 0 # Sigma1 information
RootServer: .long 0, 0, 0, 0 # Rootserver information
.long RESERVED

.global MemoryInfo
macro memoryInfo offset, number
.long ((\offset<<16) | \number)
.endm
MemoryInfo: memoryInfo offset=(MemDesc-Kip), number=0

KdebugConfig: .long 0, 0

.long RESERVED, RESERVED, RESERVED, RESERVED
.long RESERVED, RESERVED, RESERVED, RESERVED
.long RESERVED, RESERVED, RESERVED, RESERVED
.long RESERVED, RESERVED, RESERVED, RESERVED

VirtRegInfo: .long NUMMRS-1 # virtual register information
...
```

12

Onetime macros

```
KernelDesc:      .long    KERNEL_ID          # Kernel Descriptor

                .macro    kernelGenDate day, month, year
                .long      (\year-2000)<<9 | (\month<<5) | \day
                .endm
kernelGenDate day=4, month=2, year=2007

                .macro    kernelVer ver, subver, subsubver
                .long      (((\ver<<8) | \subver)<<16) | \subsubver
                .endm
kernelVer ver=1, subver=2, subsubver=0
```

13

Kernel entry points

```
SystemCalls:     .long      (spaceControlEntry - Kip)
                 .long      (threadControlEntry - Kip)
                 .long      (ipcEntry - Kip)
                 ...
                 .long      (exchangeRegistersEntry - Kip)
                 .long      (threadSwitchEntry - Kip)
                 ...

                #-- Privileged system call entry points: -----
                .align 128

spaceControlEntry:
int               $INT_SPACECONTROL
ret
threadControlEntry:
int               $INT_THREADCONTROL
ret
...

                #-- System call entry points: -----
ipcEntry:         int         $INT_IPC
ret

threadSwitchEntry:
int               $INT_THREADSWITCH
ret
...
```


14

Thread Ids

15

Thread Ids

- User programs can reference other threads using thread ids



```

/*-----
 * Thread Ids:
 *-----*/
typedef unsigned ThreadId;           // Global thread id
#define nilthread 0
#define anythread (-1)
#define anylocalthread ((-1)<<6)
#define threadId(t,v) ((t<<VERSIONBITS)|v)
#define threadNo(tid) mask((tid)>>VERSIONBITS, THREADBITS)
#define isGlobal(tid) (mask(tid,6))

```


16

Flexpages


17

Flexpages (fpages)

- A generalized form of “page” that can vary in size:



- Includes both 4KB pages and 4MB superpages as special cases
- Also includes special cases to represent the full address space (complete) and the empty address space (nilpage):



- Can be represented, in practice, using collections of 4KB and 4MB pages
- If two flexpages overlap, then one includes the other

18

Flexpage implementation

```
/*-----
 * The Flexpage datatype:
 *-----*/
typedef unsigned Fpage;

static inline Fpage fpage(unsigned base, unsigned size) {
    return align(base, size) | (size<<4);
}

static inline Fpage completeFpage(void) { // {0::Bit 22 | 1::Bit 6 |0|r|w|x}
    return (1<<4);
}

extern unsigned fsize[];
// initialized to 0 -> 0, 1 -> 32, 2 -> 0, ..., 11 -> 0,
// 12 -> 12, 13 -> 13, ..., 32 -> 32, 33 -> 0, ...
extern unsigned fpmask[];
// initialized to 0 -> 0, 1 -> ~0, 2 -> 0, ..., 11 -> 0,
// 12 -> 0xffff, 13 -> 0xffff, ..., 32 -> 0xffffffff, 33 -> 0, ...

static inline unsigned fpageMask(Fpage fp) { return fpmask[(fp>>4)&0x3f]; }
static inline unsigned fpageSize(Fpage fp) { return fsize[(fp>>4)&0x3f]; }
static inline bool isComplete(Fpage fp) { return ~fpageMask(fp) == 0; }
static inline bool isNilpage(Fpage fp) { return fpageMask(fp) == 0; }
static inline unsigned fpageStart(Fpage fp) { return fp & ~fpageMask(fp); }
static inline unsigned fpageEnd(Fpage fp) { return fp | fpageMask(fp); }
```

19

Initialization of fsize and fpmask arrays

```
void initSpaces() {
    // Basic consistency checks:
    ASSERT(mask((unsigned)Kip,PAGESIZE) == 0, "KIP alignment error");
    ASSERT((KipEnd-Kip) <= (1<<KIPAREASIZE), "KIP size error");
    ASSERT(KIPAREASIZE <= PAGESIZE, "KIP area size error");
    ASSERT(UTCBSize <= PAGESIZE, "UTCB area size error");

    // Initialize fpage mask and size arrays.
    unsigned i;
    for (i=0; i<64; i++) {
        fsize[i] = fpmask[i] = 0;
    }
    unsigned k = 0xffff;
    for (i=12; i<=32; i++) {
        fsize[i] = i;
        fpmask[i] = k;
        k = (k<<1)|1;
    }
    fsize[1] = 32;
    fpmask[1] = ~0;

    ...
}
```

20

Memory Management

21

Kernel Memory Allocator

- `void initMemory(void);`
The kernel reserves a pool of 4K pages as part of the initialization process.
- `void* allocPage1(void);`
Allocates a single page from the kernel pool
- `void freePage(void* p);`
Returns a single page to the kernel pool
- `bool availPages(unsigned n);`
Checks to see if there are (at least) n free pages
- Around ~150 lines of code, most in `initMemory()`
- No automatic GC in pork ...

22

Why `alloc1()`?

- A function f that requires the allocation of up to N pages (but never more) has a name of the form fN
- A function that calls $fN()$ will either:
 - Call `availPages(N)` beforehand
 - Have a name of the form gM , where M is N plus the number of additional pages that gM might require ...
- Goal: minimize number of checks for free pages
 - Reduce code size
 - Improve performance
 - Fewer places to write error handling code

23

Alas, this could fail:

- Consider the following function:

```
void g1() {    // 1 suffix because this function
              // allocates a page
    f();
    void* p = allocPage1();
    ...
}
```
- But now suppose `f()` takes the form:

```
void f() {
    if (availPages(1)) { ... allocPage1(); ... }
}
```
- Pork still uses this naming convention, but relies on “disciplined use”
- Maybe a type system could do better ... ?

24

Thread Control Blocks

25

Thread control blocks (TCBs)

```
struct TCB {
    ThreadId    tid;           // this thread's id and version number
    byte        status;       // thread status
    byte        prio;         // thread priority
    byte        padding;
    byte        count;        // for gc of TCBs in kernel memory
    struct UTCB* utcb;        // pointer to this thread's utcb
    unsigned    vutcb;        // virtual address of utcb


    struct TCB* sendqueue;    // list of threads waiting to send
    struct TCB* receiver;    // pointer to owner of sendqueue
    struct TCB* prev;
    struct TCB* next;

    struct Space* space;      // pointer to this thread's addr space
    unsigned    faultCode;    // exception number or page fault addr
    struct Context context;   // context of user level process

    ThreadId    scheduler;    // scheduling parameters
    unsigned    timeslice;
    unsigned    timeleft;
    unsigned    quantleft;
};
```

26

Thread control blocks (TCBs)




```
struct TCB* existsTCB(unsigned threadNo) {
    TCBTable* tab = tcbDir[threadNo >> TCBDIRBITS];
    if (tab) {
        struct TCB* tcb = ((struct TCB*)tab) + mask(threadNo, TCBDIRBITS);
        if (tcb->space) {
            return tcb;
        }
    }
    return 0;
}

struct TCB* findTCB(ThreadId tid) {
    struct TCB* tcb = existsTCB(threadNo(tid));
    return (tcb && tcb->tid==tid) ? tcb : 0;
}
```

27

Thread control blocks (TCBs)



28

Allocating and initializing TCBs

```
struct TCB* allocTCB1(ThreadId tid, struct Space* space, ThreadId scheduler) {
    unsigned threadNo = threadNo(tid);
    TCBTable* tab = tcbDir[threadNo >> TCBDIRBITS];
    if (!tab) {
        tab = tcbDir[threadNo >> TCBDIRBITS] = (TCBTable*)allocPage1();
    }
    ++tab[0]->count; // Count an additional TCB in this page
    struct TCB* tcb = ((struct TCB*)tab) + mask(threadNo, TCBDIRBITS);
    tcb->tid = tid;
    tcb->status = Halted;
    tcb->space = space;
    tcb->utcb = 0;
    tcb->vutcb = 0xffffffff;
    tcb->sendqueue = 0;
    tcb->next = tcb;
    tcb->prev = tcb;
    tcb->prio = 128; // Default is unspecified
    tcb->scheduler = scheduler;
    tcb->timeslice = 0;
    tcb->timeleft = 10000; // Default timeslice is 10ms
    tcb->quantleft = 0; // Default quantum is infinite
    initUserContext(&(tcb->context));
    enterSpace(space); // Register the thread in this space
    return tcb;
}
```

29


Allocating/Freeing TCBs

- Primitives: `allocTCB1()` and `destroyTCB()`
- Every TCB contains a count field
- Reference count TCBtables using the count field for the first TCB in the table
- New TCBTables obtained via `allocPage1()`
- Empty TCBTables recycled using `freePage()`
- Lazy initialization (page already zeroed)

30

Type Safety?


- To maintain type safety if we were doing this in Habit:



31

Type Safety?

- To maintain type safety if we were doing this in Habit:




- must be fully initialized at allocation time ...

32

Type Safety?


- To maintain type safety if we were doing this in Habit:



33

Type Safety?

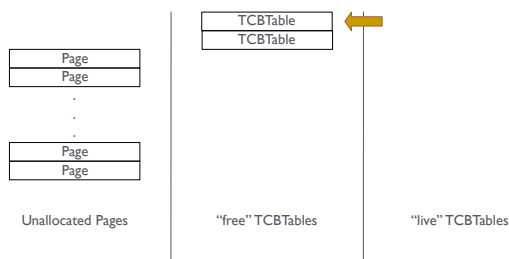
- To maintain type safety if we were doing this in Habit:



34

Type Safety?


- To maintain type safety if we were doing this in Habit:



35

Type Safety?

- To maintain type safety if we were doing this in Habit:




- Pages are never truly freed ...

36

Type Safety?

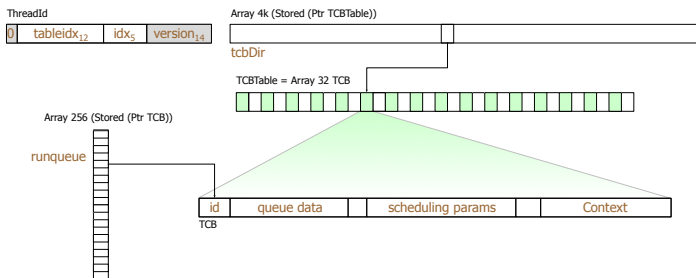
- To maintain type safety if we were doing this in Habit:



- ... unless we have garbage collection ... ?


37

Thread Control Blocks (TCBs)




38

Scheduling data structures: runqueue



39

Scheduling data structures: prioiset



```
/*-----  
 * Select a new thread to execute. We pick the next runnable thread with  
 * the highest priority.  
 */  
void reschedule() {  
    switchTo(holder = priosetSize ? runqueue[prioiset[0]] : idleTCB);  
}
```

40

Switching to a new thread

```
static void inline switchTo(struct TCB* tcb) {  
    struct Context* ctxt = &(tcb->context);  
    DEBUG(printf("Switching to thread %x (tcb=%x)\n", tcb->tid, tcb));  
    current = tcb; // Change current thread  
    *utcbptr = tcb->vutcb; // Change UTCB address  
    + (unsigned)&(((struct UTCB*)0)->mr[0]);  
    DEBUG(printf("set utcbptr to %x\n", *utcbptr));  
    esp0 = (unsigned)(ctxt + 1); // Change esp0  
    DEBUG(extern void showSpace(struct Space* space);  
    DEBUG(showSpace(tcb->space);)  
    switchSpace(tcb->space); // Change address space  
    DEBUG(printf("switched space, context is at %x\n\n", ctxt));  
    returnToContext(ctxt);  
}
```

41

Switching to a new thread (w/o debugging)

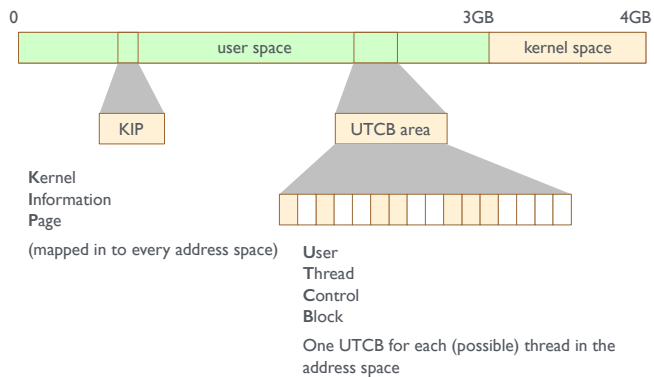
```
static void inline switchTo(struct TCB* tcb) {  
    struct Context* ctxt = &(tcb->context);  
    current = tcb; // Change current thread  
    *utcbptr = tcb->vutcb; // Change UTCB address  
    + (unsigned)&(((struct UTCB*)0)->mr[0]);  
    esp0 = (unsigned)(ctxt + 1); // Change esp0  
    switchSpace(tcb->space); // Change address space  
    returnToContext(ctxt);  
}  
  
...  
  
void switchSpace(struct Space* space) {  
    if (space->pdir) { // No switch for kernel/inactive threads  
        if (currentSpace != space) {  
            currentSpace = space;  
            setPdir(currentSpace->pdir);  
            currentSpace->loaded = 1;  
        } else {  
            refreshSpace();  
        }  
    }  
}
```

42

Address Spaces

43

Address space layout



44

Representing address spaces

```
struct Space {  
    unsigned pdir;           // Physical address of page directory  
    struct Mapping* mem;      // Memory map  
    Fpage kipArea;           // Location of kernel interface page  
    Fpage utcbArea;          // Location of UTCBs  
    unsigned count;          // Count of threads in this space  
    unsigned active;         // Count of active threads in this space  
    unsigned loaded;         // 1 => already loaded in cr3  
};  
  
...  
  
void enterSpace(struct Space* space) {  
    space->count++; // increment reference count;  
}  
  
...  
  
void configureSpace(struct Space* space, Fpage kipArea, Fpage utcbArea) {  
    ASSERT(!activeSpace(space), "configuring active space");  
    space->kipArea = kipArea;  
    space->utcbArea = utcbArea;  
}
```

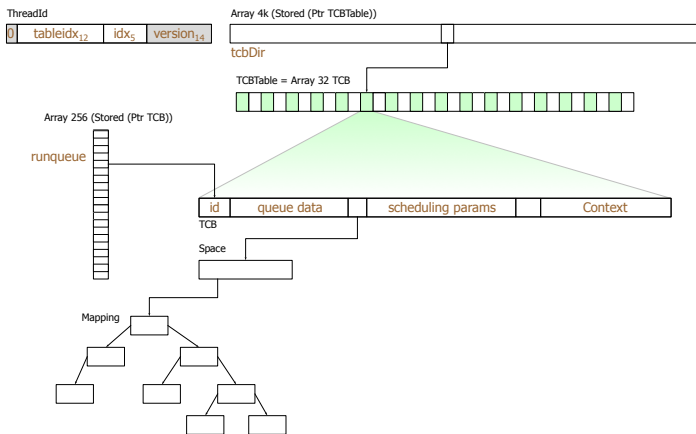
45

A typical system call

```
ENTRY spaceControl() {
    if (!privileged(current->space)) { /* check for privileged thread */
        retError(SpaceControl_Result, NO_PRIVILEGE);
    } else {
        struct TCB* dest = findTCB(SpaceControl_SpaceSpecifier);
        if (!dest) {
            retError(SpaceControl_Result, INVALID_SPACE);
        } else if (!activeSpace(dest->space)) { /* ignore if active threads */
            Fpage kipArea = SpaceControl_KipArea;
            Fpage utcbArea = SpaceControl_UtcbArea;
            unsigned kipEnd, utcbEnd;
            if (!isNilpage(utcbArea)) /* validate utcb area */
            || fpageSize(utcbArea) < MIN_UTCBAREASIZE
            || (utcbEnd=fpageEnd(utcbArea)) >= KERNEL_SPACE) {
                retError(SpaceControl_Result, INVALID_UTCB);
            } else if (!isNilpage(kipArea)) /* validate KIP area */
            || fpageSize(kipArea) != KIPAREASIZE
            || (kipEnd=fpageEnd(kipArea)) >= KERNEL_SPACE
            || (kipEnd=fpageStart(utcbArea) && utcbEnd=fpageStart(kipArea)) {
                retError(SpaceControl_Result, INVALID_KIPAREA);
            } else {
                configureSpace(dest->space, kipArea, utcbArea);
            }
        }
        SpaceControl_Result = 1;
        SpaceControl_Control = 0; /* control parameter is not used */
        resume();
    }
}
```

46

Spaces and mappings



47

Representing mappings

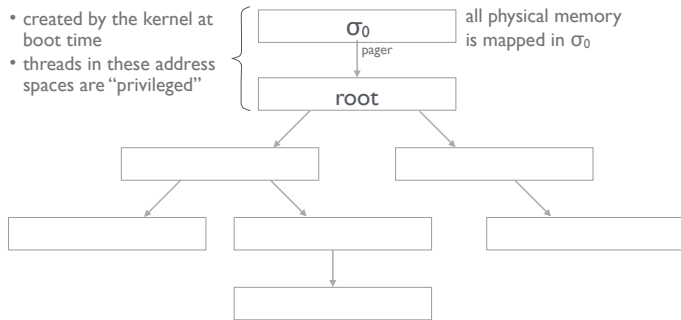
```
struct Mapping {
    struct Space* space; /* Which address space is this in? */
    struct Mapping* next;
    struct Mapping* prev;
    unsigned level;
    Fpage vfp; /* Virtual fpage */
    unsigned phys; /* Physical address */
    struct Mapping* left;
    struct Mapping* right;
};
```

- A binary search tree of memory regions within a single address space
- A mapping data base that documents the way that memory regions have been mapped between address spaces

48

The “recursive address space model”

- created by the kernel at boot time
- threads in these address spaces are “privileged”




- In a dynamic system, we need the ability to revoke a previous mappings ... this will get interesting ...

49

Small Objects:

- Pork uses only two “small” object types (≤ 32 bytes):
 - Address space descriptors (Space)
 - Mapping descriptors (Mapping)
- Kernel allocates/frees pages to store small objects (each page can store up to 127 objects)
- Pages with free slots are linked together



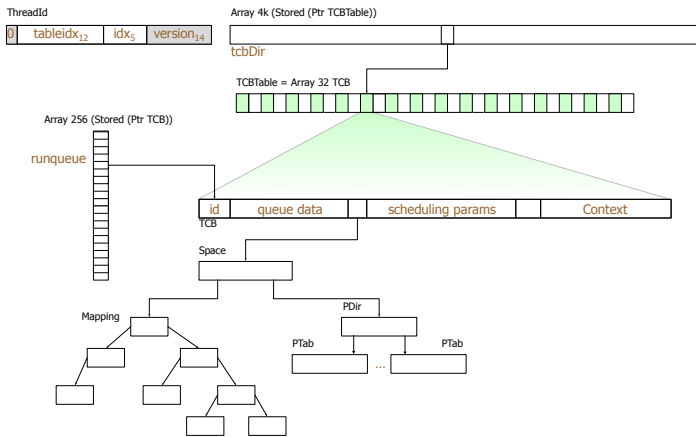
50

Small Objects, continued:

- Issues:
 - Fragmentation: Compacting GC is probably feasible but possibly expensive
 - Denial of service/interference between spaces: Could have a separate pool of object pages for each address space ...
- Costs?
 - Current implementation for allocating/freeing small objects takes ~70 lines, including header structure declaration
- For memory safety: separate pools of Space and Mapping objects ...

51

Page Dirs and Page Tables:



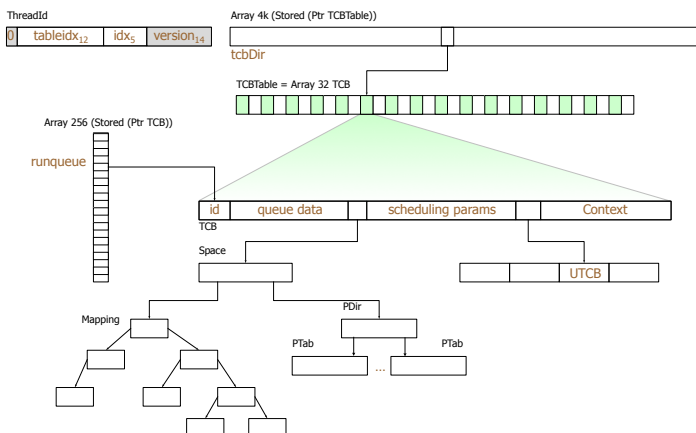
52

Page Dirs and Page Tables:

- Page directories and page tables use full pages
- Aligned pointers take fewer than 32 bits; how can this be tracked in the type system?
- Physical rather than virtual addresses
- Allocation and deallocation controlled via reference counts in Space structures

53

User TCBs (UTCBs):



54

User TCBs (UTCBS):

- Mapped into kernel space, also fully accessible in user space
- Just bits: No pointers, references, indexes
- Could be recycled into the pool of kernel pages
- Currently freed only when all active threads in a space have terminated

55

IPC

56

Thread status

```
/*-----  
 * Thread status:  
 * A byte field in each TCB specifies the current status of that thread:  
 * +-----+  
 * | b6 | b5 | b4 | ipctype |  
 * +-----+  
 * b3-b0: ipctype (4 bits)  
 * b4: 1=>halted, or halt requested (i.e., will halt after IPC)  
 * b5: 1=>blocked waiting to send an ipc of the specified type  
 * b6: 1=>blocked waiting to receive an ipc of the specified type  
 * A zero status byte indicates that the thread is Runnable.  
 *-----*/  
#define Runnable      0  
#define Halted        0x10  
#define Sending(type) (0x20|(type))  
#define Receiving(type) (0x40|(type))  
  
typedef enum {  
    MRS, PageFault, Exception, Interrupt, Preempt, Startup  
} IPCType;  
  
static inline IPCType ipctype(struct TCB* tcb) {  
    return (IPCType)(tcb->status & 0xf);  
}
```

57

The ipc system call

```
/*-----  
 * The "IPC" System Call:  
 *-----*/  
ENTRY ipc() {  
  
    ThreadId to = IPC_GetTo;                // Send Phase  
    if (to!=nilthread) {  
        if (!sendPhase(MRs, current, to)) {  
            reschedule();  
        }  
    }  
  
    ThreadId fromSpec = IPC_GetFromSpec(current); // Receive Phase  
    if (fromSpec!=nilthread) {  
        current->utcb->mr[0] = 0;  
        recvPhase(MRs, current, fromSpec);  
    }  
  
    reschedule();  
}
```

58

The send phase (Part 1)

```
static bool sendPhase(IPCType sendtype, struct TCB* send, ThreadId recvId) {  
    // Find the receiver TCB: -----  
    struct TCB* recv;  
    if (recvId==anythread) ||  
        recvId==anylocalthread ||  
        ! (recv=findTCB(recvId)) {  
        sendError(sendtype, send, NonExistingPartner);  
        return 0;  
    }  
  
    // Determine whether we can send the message immediately: -----  
    if (isReceiving(recv)) {  
        IPCType recvtype = iptype(recv);  
        ThreadId srcId = recvFromSpec(recvtype, recv);  
        if ((srcId==send->tid) ||  
            (srcId==anythread) ||  
            (srcId==anylocalthread && send->space==recv->space)) {  
            // Destination is blocked and ready to receive from send:  
            IPCErr err = transferMessage(sendtype, send, recvtype, recv);  
            if (err==NoError) {  
                resumeThread(recv);  
                return 1;  
            } else {  
                sendError(sendtype, send, err);  
                recvError(recvtype, recv, err);  
                return 0;  
            }  
        }  
    }  
    ...  
}
```

59

The send phase (Part 2)

```
...  
// Destination is not ready to receive a message, so try to block: -----  
if (sendCanBlock(sendtype, send)) {  
    if (send->status==Runnable) {  
        removeRunnable(send);  
    }  
    send->status = Sending(sendtype) | (Halted & send->status);  
    send->receiver = recv;  
    recv->sendqueue = insertTCB(recv->sendqueue, send);  
} else {  
    sendError(sendtype, send, NoPartner);  
}  
return 0;  
}
```

60

Transferring messages

```
static IPCErr transferMessage(IPCType sendtype, struct TCB* send,
                             IPCType recvtype, struct TCB* recv) {
    if (recvtype==MRS) {          // Send to MRS (Destination is user ipc)
        ...
        switch (sendtype) {
            case MRS : ... // Send between sets of message registers
            case PageFault : ... // Send pagefault message to pager
            case Exception : ... // Send message to an exception handler
            case Interrupt : ... // Send message to an interrupt handler
        }
    } else if (sendtype==MRS) { // Receive from MRS (Source is user ipc)
        ...
        switch (recvtype) {
            case PageFault : ... // Receive a response from a pager
            case Exception : ... // Receive a response from an exception handler
            case Interrupt : ... // Receive a response from an interrupt handler
            case Startup : ... // Receive startup message from thread's pager
        }
    }
    return Protocol; // Protocol error: incompatible types/format
}
```

61

Regular IPC:

MRS \Rightarrow MRS

```
struct UTCB* rutcb = recv->utcb;
struct UTCB* sutcb = send->utcb;
unsigned u = mask(sutcb->mr[0], 6); // untyped items
unsigned t = mask(sutcb->mr[0]>>6, 6); // typed items
if ((u+t>=NUMMRS) || (t&1)) {
    return MessageOverflow;
} else {
    unsigned i;
    rutcb->mr[0] = MsgTag(sutcb->mr[0]>>16, 0, t, u);
    for (i=1; i<=u; i++) {
        rutcb->mr[i] = sutcb->mr[i];
    }
    if (t>0) {
        Fpage acc = rutcb->acceptor;
        do {
            IPCErr err = transferTyped(send, recv, acc,
                                       rutcb->mr[i] = sutcb->mr[i],
                                       rutcb->mr[i+1] = sutcb->mr[i+1]);
            if (err!=NoError) {
                return err;
            }
            i += 2;
        } while ((t-=2)>0);
    }
    return NoError;
}
```

62

Example: IPCs from hardware interrupts

```
ENTRY hardwareIRQ() {
    unsigned n = current->context.iret.error;
    maskAckIRQ(n); // Mask and acknowledge the interrupt with the PIC
    struct TCB* irqTCB = existsTCB(n);

    // An irq thread may be Halted, Sending (i.e., waiting for the user level
    // handler to receive notice of the interrupt), or Receiving (i.e., waiting
    // for the user level handler to finish processing the interrupt and send
    // an acknowledgement). In theory, an interrupt can only occur if it is
    // unmasked at the PIC, and that, in turn, should only be possible when
    // (1) the corresponding irq thread is Halted; and (2) the "pager" for
    // the irq thread (stored in the vutcb field) is set to a non-nilthread id.

    if (irqTCB->status==Halted && irqTCB->vutcb!=nilthread) {
        if (sendPhase(Interrupt, irqTCB, irqTCB->vutcb)) {
            irqTCB->status = Receiving(Interrupt) | Halted;
        }
    }
    reschedule(); // allow the user level handler to begin ...
}
```

63

Interrupt handler protocol

- When a hardware interrupt occurs, the kernel sends an IPC message from the interrupt thread to its pager with the tag:

From Interrupt Thread

-1 _(12/44)	0 ₍₄₎	0 ₍₄₎	t = 0 ₍₆₎	u = 0 ₍₆₎	MR ₀
-----------------------	------------------	------------------	----------------------	----------------------	-----------------

- When the pager has finished handling the error, it sends an IPC message back to the interrupt thread to reenale the corresponding interrupt

To Interrupt Thread

0 _(16/48)	0 ₍₄₎	t = 0 ₍₆₎	u = 0 ₍₆₎	MR ₀
----------------------	------------------	----------------------	----------------------	-----------------

64

Interrupt handler protocol

Interrupt \Rightarrow MRs

- When a hardware interrupt occurs, the kernel sends an IPC message from the interrupt thread to its pager with the tag:

From Interrupt Thread

-1 _(12/44)	0 ₍₄₎	0 ₍₄₎	t = 0 ₍₆₎	u = 0 ₍₆₎	MR ₀
-----------------------	------------------	------------------	----------------------	----------------------	-----------------

```
case Interrupt :    // Send message to an interrupt handler
    rutcb->mr[0] = MsgTag((-1)<<4, 0, 0, 0);
    return NoError;
```

65

Interrupt handler protocol

MRs \Rightarrow Interrupt

```
case Interrupt :    // Receive a response from an interrupt handler
    if (mask(sutcb->mr[0],12)==0) {
        ASSERT(mask(recv->tid, VERSIONBITS)==1, "Wrong irq version");
        ASSERT(threadNo(recv->tid) < NUMIRQs, "IRQ out of range");
        enableIRQ(threadNo(recv->tid));    // Reenable interrupt
        return NoError;
    }
    break;
```

- When the pager has finished handling the error, it sends an IPC message back to the interrupt thread to reenale the corresponding interrupt

To Interrupt Thread

0 _(16/48)	0 ₍₄₎	t = 0 ₍₆₎	u = 0 ₍₆₎	MR ₀
----------------------	------------------	----------------------	----------------------	-----------------

66

Example: IPCs from page faults

```
ENTRY pageFault() {
    asm("    movl %%cr2, %0\n" : "=r"(current->faultCode));

    if (current->space==sigma0Space && sigma0map(current->faultCode)) {
        printf("sigma0 case succeeded!\n");
    } else {
        ThreadId pagerId = current->utcb->pager;
        if (pagerId==nilthread) {
            haltThread(current);
        } else if (sendPhase(PageFault, current, pagerId)) {
            removeRunnable(current);    // Block current if message already delivered
            current->status = Receiving(PageFault);
        }
    }
    refreshSpace();
    reschedule();
}
```

67

Page fault protocol

- When a thread triggers a page fault, the kernel translates that event into an IPC to the thread's pager:

To Pager

faulting user-level IP (32/64)					MR ₂
fault address (32/64)					MR ₁
-2 (12/44)	0 r w x	0 (4)	t = 0 (6)	u = 2 (6)	MR ₀

- The pager can respond by sending back a reply with a new mapping ... that also restarts the faulting thread:

From Pager

MapItem / GrantItem					MR _{1,2}
0 (16/48)	0 (4)	t = 2 (6)	u = 0 (6)		MR ₀

68

Page fault protocol

PageFault ⇒ MRs

- When a thread triggers a page fault, the kernel translates that event into an IPC to the thread's pager:

To Pager

faulting user-level IP (32/64)					MR ₂
fault address (32/64)					MR ₁
-2 (12/44)	0 r w x	0 (4)	t = 0 (6)	u = 2 (6)	MR ₀

```
case PageFault : { // Send pagefault message to pager
    unsigned rwx = (send->context.iret.error & 2) ? 2 : 4;
    rutcb->mr[0] = MsgTag((-2)<<4|rwx, 0, 0, 2);
    rutcb->mr[1] = send->faultCode;
    rutcb->mr[2] = send->context.iret.eip;
}
return NoError;
```

69

Page fault protocol

MRS \Rightarrow PageFault

```
case PageFault : // Receive a response from a pager
  if (mask(sutcb->mr[0],12)==MsgTag(0, 0, 2, 0)) {
    return transferTyped(send, recv,
      completeFpage(), sutcb->mr[1], sutcb->mr[2]);
  }
  break;
```

- The pager can respond by sending back a reply with a new mapping ... that also restarts the faulting thread:

From Pager

MapItem / GrantItem				MR _{1,2}
0 (16/48)	0 (4)	$t = 2$ (6)	$u = 0$ (6)	MR ₀

70

Let's poke around ... !

71


72

Performance Benchmarking: Pingpong, Pistachio, and Pork

The pingpong benchmark

- A small L4 benchmark from the Karlsruhe pistachio distribution, written in C++
- A single ipc call transfers contents of n message registers (MRs) between threads
- create two threads, “ping” & “pong”:
for $n = 0, 4, 8, \dots, 60$:
for 128K times:
 send n MRs from “ping” to “pong”
 send n MRs from “pong” to “ping”
 measure cycles & time per ipc call
- Cycles measured using rdtsc, time measured using interrupts

Expected Performance Model



Test Platform



- Dell Mini 9 netbook (1.6GHz Atom N270 CPU)
- Booting via grub from a flashdrive

Pistachio “Output”

```
5: INTER-BS (small)
s: print SQLite table
Benchmarking Inter-BS IPC...
IPC ( 0 MRs): 1240.67 cycles, 0.77us, 0.00 instrs
IPC ( 4 MRs): 1293.50 cycles, 0.81us, 0.00 instrs
IPC ( 8 MRs): 1301.64 cycles, 0.81us, 0.00 instrs
IPC (12 MRs): 1306.29 cycles, 0.81us, 0.00 instrs
IPC (16 MRs): 1317.96 cycles, 0.82us, 0.00 instrs
IPC (20 MRs): 1325.16 cycles, 0.83us, 0.00 instrs
IPC (24 MRs): 1333.26 cycles, 0.83us, 0.00 instrs
IPC (28 MRs): 1342.20 cycles, 0.84us, 0.00 instrs
IPC (32 MRs): 1350.34 cycles, 0.84us, 0.00 instrs
IPC (36 MRs): 1358.46 cycles, 0.85us, 0.00 instrs
IPC (40 MRs): 1362.08 cycles, 0.85us, 0.00 instrs
IPC (44 MRs): 1374.64 cycles, 0.86us, 0.00 instrs
IPC (48 MRs): 1382.00 cycles, 0.86us, 0.00 instrs
IPC (52 MRs): 1390.00 cycles, 0.87us, 0.00 instrs
IPC (56 MRs): 1398.02 cycles, 0.87us, 0.00 instrs
IPC (60 MRs): 1406.13 cycles, 0.88us, 0.00 instrs

What now?
g - Continue
q - Quit/Show measurement
ESC - Enter KIB
```

Pork “Output”

```
The Portland IA Manual (pork), February 2007

IPC ( 0 MRs): 1519.50 cycles, 0.95us, 0.00 instrs
IPC ( 4 MRs): 1530.14 cycles, 0.95us, 0.00 instrs
IPC ( 8 MRs): 1556.71 cycles, 0.99us, 0.00 instrs
IPC (12 MRs): 1579.67 cycles, 0.99us, 0.00 instrs
IPC (16 MRs): 1607.34 cycles, 1.02us, 0.00 instrs
IPC (20 MRs): 1634.90 cycles, 1.02us, 0.00 instrs
IPC (24 MRs): 1664.64 cycles, 1.02us, 0.00 instrs
IPC (28 MRs): 1687.47 cycles, 1.02us, 0.00 instrs
IPC (32 MRs): 1702.00 cycles, 1.06us, 0.00 instrs
IPC (36 MRs): 1721.46 cycles, 1.06us, 0.00 instrs
IPC (40 MRs): 1745.56 cycles, 1.10us, 0.00 instrs
IPC (44 MRs): 1707.86 cycles, 1.14us, 0.00 instrs
IPC (48 MRs): 1804.40 cycles, 1.14us, 0.00 instrs
IPC (52 MRs): 1818.70 cycles, 1.14us, 0.00 instrs
IPC (56 MRs): 1842.79 cycles, 1.14us, 0.00 instrs
IPC (60 MRs): 1875.66 cycles, 1.18us, 0.00 instrs


What now?
g - Continue
q - Quit/Show measurement
```

Transcribed Data (Inter-AS)

ping pong #MRs	pistachio Inter-AS IPC		pork Inter-AS IPC		Ratio, pork/pistachio	
	cycles	microseconds	cycles	microseconds	cycles	microseconds
0	1240.67	0.77	1519.59	0.95	1.22	1.23
4	1293.58	0.81	1530.14	0.95	1.18	1.17
8	1301.64	0.81	1556.71	0.99	1.20	1.22
12	1306.29	0.81	1579.67	0.99	1.21	1.22
16	1317.96	0.82	1607.34	1.02	1.22	1.24
20	1325.16	0.83	1634.98	1.02	1.23	1.23
24	1333.26	0.83	1664.64	1.02	1.25	1.23
28	1342.28	0.84	1687.47	1.02	1.26	1.21
32	1350.34	0.84	1702.89	1.06	1.26	1.26
36	1358.46	0.85	1721.46	1.06	1.27	1.25
40	1362.08	0.85	1745.56	1.10	1.28	1.29
44	1374.64	0.86	1787.86	1.14	1.30	1.33
48	1382.80	0.86	1804.40	1.14	1.30	1.33
52	1390.88	0.87	1818.78	1.14	1.31	1.31
56	1398.02	0.87	1842.79	1.14	1.32	1.31
60	1406.13	0.88	1875.66	1.18	1.33	1.34


Inter-AS = "ping" and "pong" in different address spaces

Cycles (Inter-AS)




pistachio = 1274.66 + 2.27n (least squares)
pork = 1512.57 + 6n

Microseconds (Inter-AS)



Pork : Pistachio (Inter-AS)




Transcribed Data (Intra-AS)

ping pong #MRs	pistachio Intra-AS IPC		pork Intra-AS IPC		Ratio, pork/pistachio	
	cycles	microseconds	cycles	microseconds	cycles	microseconds
0	729.19	0.45	1078.71	0.68	1.48	1.51
4	774.74	0.48	1097.90	0.68	1.42	1.42
8	778.49	0.48	1115.55	0.72	1.43	1.50
12	790.04	0.49	1143.99	0.72	1.45	1.47
16	795.65	0.49	1171.99	0.72	1.47	1.47
20	806.12	0.50	1193.23	0.76	1.48	1.52
24	811.85	0.50	1219.75	0.76	1.50	1.52
28	822.54	0.51	1247.19	0.76	1.52	1.49
32	827.20	0.51	1271.19	0.80	1.54	1.57
36	838.69	0.52	1295.20	0.80	1.54	1.54
40	843.37	0.52	1319.39	0.83	1.56	1.60
44	855.89	0.53	1343.43	0.83	1.57	1.57
48	859.57	0.53	1363.04	0.87	1.59	1.64
52	871.08	0.54	1391.45	0.87	1.60	1.61
56	875.72	0.54	1415.61	0.91	1.62	1.69
60	887.38	0.55	1439.58	0.91	1.62	1.65


Intra-AS = "ping" and "pong" in same address space

Cycles (Intra-AS)




pistachio = 756.54 + 2.21n (least squares)
pork = 1073.54 + 6.11n

Microseconds (Intra-AS)



Pork : Pistachio (Intra-AS)



Estimating Clock Frequency

cycles/microsecond	
pistachio	pork
1611.26	1599.57
1597.01	1610.67
1606.96	1572.43
1612.70	1595.63
1607.27	1575.82
1596.58	1602.92
1606.34	1632.00
1597.95	1654.38
1607.55	1606.50
1598.19	1624.02
1602.45	1586.87
1598.42	1568.30
1607.91	1582.81
1598.71	1595.42
1606.92	1616.48
1597.88	1589.54

Inter-AS

cycles/microsecond	
pistachio	pork
1620.42	1586.34
1614.04	1614.56
1621.85	1549.38
1612.33	1588.88
1623.78	1627.76
1612.24	1570.04
1623.70	1604.93
1612.82	1641.04
1621.96	1588.99
1612.87	1619.00
1621.87	1589.63
1614.89	1618.59
1621.83	1566.71
1613.11	1599.37
1621.70	1555.62
1613.42	1581.96

Intra-AS

Pretty consistent with 1.6GHz processor frequency, but estimates from pork are typically a little lower than those for pistachio

Summary

Comparison	Range
Pork/Pistachio (Inter-AS)	1.17 – 1.35
Pork/Pistachio (Intra-AS)	1.42 – 1.65
Inter-AS/Intra-AS (Pork)	1.58 – 1.70
Inter-AS/Intra-AS (Pistachio)	1.30 – 1.40

- IPC in Pork is slower than Pistachio (17-65%)
- Overhead for crossing address spaces is higher in pork than Pistachio (65% vs 35%)

Performance Tuning Opportunities?

- Are there opportunities for performance-tuning pork to reduce the gap?
- Inter-AS: $\text{pistachio} = 1274.66 + 2.27n$ (least squares)
 $\text{pork} = 1512.57 + 6n$
- Intra-AS: $\text{pistachio} = 756.54 + 2.21n$ (least squares)
 $\text{pork} = 1073.54 + 6.11n$
- Example: pork takes ~6 cycles to transfer a machine word, where pistachio uses around ~2

Transfer Message in pork

- Source:

```
for (i=1; i<=u; i++) {
    rutcb->mr[i] = sutcb->mr[i];
}
```

- Machine Code:

209: ba 01 00 00 00

mov \$0x1,%edx initialization

20e: 8b 84 97 00 01 00 00

mov 0x100(%edi,%edx,4),%eax

215: 89 84 91 00 01 00 00

mov %eax,0x100(%ecx,%edx,4)

21c: 83 c2 01

add \$0x1,%edx

21f: 39 d3

cmp %edx,%ebx

221: 73 eb

jae 20e

loop

Transfer Message in pistachio

- Source:

```

inline void tcb_t::copy_msr(tcb_t * dest, word_t start, word_t count)
{
    ASSERT(start + count <= IPC_NUM_MR);
    ASSERT(count > 0);
    word_t dummy;

    #if defined(CONFIG_X86_SMALL_SPACES)
    asm volatile ("mov %0, %%es" : : "r" (X86_KDS));
    #endif

    /* use optimized IA32 copy loop -- uses complete cacheline
    transfer */
    __asm__ volatile__ (
        "cld\n"
        "rep movsl (%0), (%1)\n"
        : /* output */
        : %S(dummy), ~D(dummy), ~C(dummy)
        : /* input */
        : "c"(count), "D"(&get_tcb(T->mr[start]),
        "D"(&dest->get_tcb(T->mr[start]));

    #if defined(CONFIG_X86_SMALL_SPACES)
    asm volatile ("mov %0, %%es" : : "r" (X86_UDS));
    #endif
}

```

Transfer Message in pistachio

- Machine Code:

```

b15: 31 c9
b17: 8b 73 0c
b1a: 8b 7d 0c
b1d: 88 d1
b1f: 81 c6 04 01 00 00
b25: 81 c7 04 01 00 00
b2b: fc
b2c: f3 a5

```

```

xor    %ecx,%ecx      initialization
mov    0xc(%ebx),%esi
mov    0xc(%ebp),%edi
mov    %dl,%cl
add    $0x104,%esi
add    $0x104,%edi
cld

```

```

rep movsl %ds:(%esi), %es:(%edi)
loop

```

Reflections

- In this case, the performance differences between pork and pistachio can be understood and addressed
 - Could be handled by a compiler intrinsic (looks like a function, but treated specially by the compiler)
 - Familiar in C (memcpy)
- How easily can other performance gaps be closed?
 - Other opportunities for intrinsics? Special handling for fast paths? Algorithmic tweaks? Refined choice of data structures? etc.