



## CS 410/510


### Languages & Low-Level Programming

Mark P Jones  
Portland State University

Spring 2016

Week 10: Habit and HaL4

1



High-level  
languages

Low-level  
programming


2

## The CEMLaBS Project

- “Using a Capability-Enhanced Microkernel as a Testbed for Language-Based Security”
- Started October 2014, Funded by The National Science Foundation
- Three main questions:
  - **Feasibility:** Is it possible to build an inherently “unsafe” system like seL4 in a “safe” language like Habit?
  - **Benefit:** What benefits might this have, for example, in reducing verification costs?
  - **Performance:** Is it possible to meet reasonable performance goals for this kind of system?


3

Chipping away ...



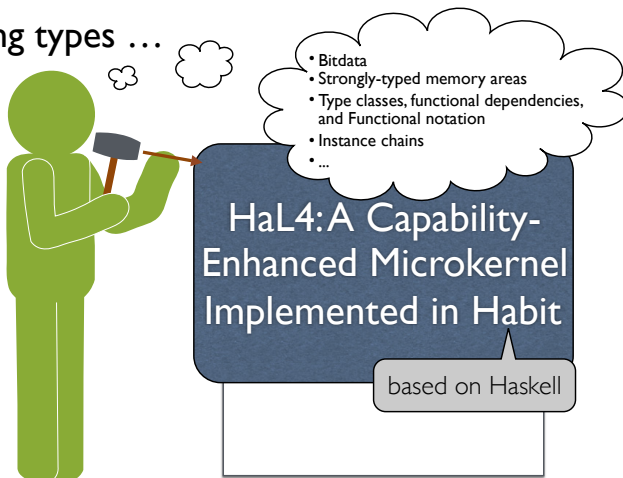
4

Chipping away ...



5

Using types ...



6

## Example: IA32 Paging Structures

Remember this?

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0						
Address of page directory <sup>1</sup>																Ignored				P C D	Pw T	Ignored				CR3											
Bits 31:22 of address of 4MB page frame								Reserved (must be 0)				Bits 39:32 of address <sup>2</sup>				P A T	Ignored				G	1	D	A	P C D	Pw T	U / S	R / W	1	PDE: 4MB page							
Address of page table																Ignored				0	I g n	A	P C D	Pw T	U / S	R / W	1	PDE: page table									
Ignored																										0	PDE: not present										
Address of 4KB page frame																Ignored				G	P A T	D	A	P C D	Pw T	U / S	R / W	1	PTE: 4KB page								
Ignored																										0	PTE: not present										

Figure 4-4. Formats of CR3 and Paging-Structure Entries with 32-Bit Paging

7

## Example: IA32 Paging Structures

Here is how we describe page directory entries in Habit:

```
bitdata PDE /WordSize -- Page Directory Entries
= UnmappedPDE { unused=0 :: Bit 31 | B0 } -- Unused entry (present bit reset)
| PageTablePDE { ptab :: Phys PageTable -- physical address of page table
| unused=0 :: Bit 4
| B0 -- signals PageTablePDE
| attrs=readWrite :: PagingAttrs -- paging attributes
| B1 } -- present bit set
| SuperPagePDE { super :: Phys SuperPage -- physical address of superpage
| unused=0 :: Bit 13
| global=0 :: Bit 1 -- 1 => global translation (if cr4.pge=1)
| B1 -- signals SuperPagePDE
| attrs :: PagingAttrs -- paging attributes
| B1 } -- present bit set
bitdata PagingAttrs /6
= PagingAttrs { dirty = 0 :: Bit 1 -- Dirty; 1 => data written to page
| accessed = 0 :: Bit 1 -- Accessed; 1 => page accessed
| caching = Caching[] :: Caching
| us :: Bit 1 -- User/supervisor; 1 => user access allowed
| rw :: Bit 1 } -- Read/write; 1 => write access allowed
```

8

## Example: IA32 Address Space Layout

Remember this?

user space

kernel space

0GB
1GB
2GB
3GB
4GB

```
struct PageDir /4K { pdes :: Array UserSuperPages (Stored PDE)
| kpdes :: Array KernelSuperPages (Stored KPDE) }

type UserSuperPages = 768 -- Number of superpages in the user space address range
type KernelSuperPages = 256 -- Number of superpages in the kernel space address range

bitdata KPDE /WordSize -- Kernel Page Directory Entries, a variant of SuperPagePDE
= KPDE { B00 -- leading zeros
| ix :: Ix KernelSuperPages -- physical superpage number
| unused=0 :: Bit 13
| global=1 :: Bit 1 -- 1 => global translation (if cr4.pge=1)
| B1 -- indicates a form of SuperPagePDE
| attrs = kernelOnly :: PagingAttrs -- paging attributes
| B1 }

instance Initable PageDir where
initialize = PageDir { pdes <- initArray (\ix -> initStored UnmappedPDE[])
| kpdes <- initArray (\ix -> initStored KPDE[ix]) }
```

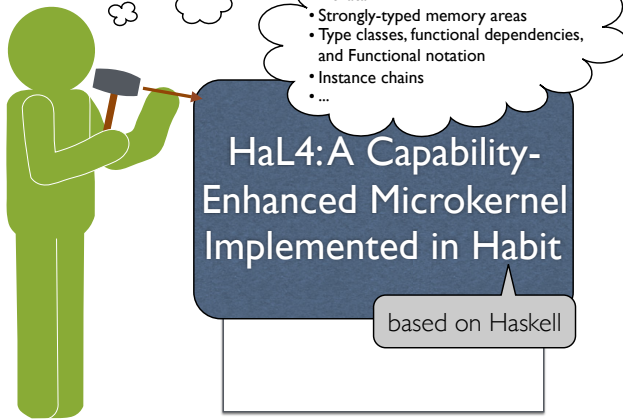
9

## Summary

- The code is not simple ...
- ... but this reflects the underlying data structures
- Correct usage is enforced by the type system:
  - The separation between user space and kernel space
  - The requirement for physical rather than virtual addresses in page directory entries
  - Invariants on data structure size
  - Initialization


10

## Using types ...




11

## Using lambda ...




12

Using lambda ...




13

Using lambda ...



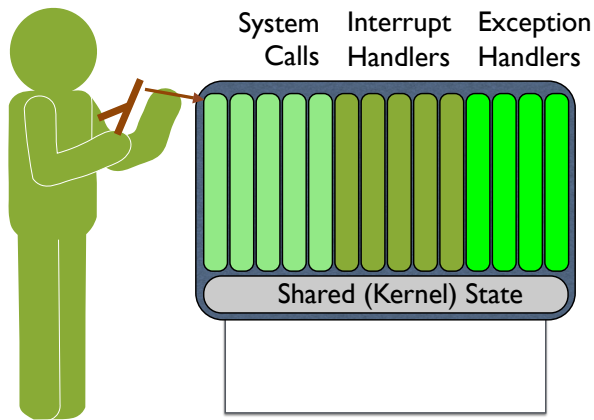
14

Using lambda ...



15

## Using lambda ...



16

Validate  
Parameters

Return with  
error code

Perform  
Action

Return with  
success code

17

```
syscallMapPageDir :: (KE k, KW k) => k a
syscallMapPageDir
= do curr <- getCurrent
    asidIdx <- getReg asidCapReg curr
    case<- lookupCapAll curr.cspace asidIdx of
      Ref asidCap ->
        case<- get asidCap.objptr of
          ASIDTableObj[] ->
            range <- getCapdata asidCap
            offset <- getReg offsetReg
            case offset `inRange` range of
              Just asid ->
                let slot = asidTable @@ asid
                count <- get slot.count
                if count==0 then
                  pdirIdx <- getReg pdirCapReg curr
                  case<- lookupCapAll curr.cspace pdirIdx curr of
                    Ref pdirCap ->
                      case<- get pdirCap.objptr of
                        PageDirObj[pdir] ->
                          case<- getCapdata pdirCap of
                            UnmappedPD[] ->
                              set slot.pdir (Ref pdir)
                              set slot.count 1
                              setCapdata pdirCap MappedPD[asid]
                              success curr
                              -> mappedErr curr
                            -> invalidCapabilityErr curr
                            Null -> invalidCapabilityErr curr
                            else mappedErr curr
                            Nothing -> rangeErrorErr curr
                            _ -> invalidCapabilityErr curr
                            Null -> invalidCapabilityErr curr
```

Parameter  
Validation

Action

Error  
Reporting

18

## Programming with continuations

- Traditional control flow structure

```
do f <- openFile "file.txt"
  l1 <- readLine f
  l2 <- readLine f
  out (l1, l2)
  closeFile f
```

- How to deal with errors?
  - Make functions return error codes (and hope that callers will check those codes)?
  - Add the ability to throw and catch exceptions?
  - Use continuations ...

19

## Programming with continuations

- Instead of

```
openFile :: String -> IO FileHandle
```

- Try:

```
openFile :: String
          -> (ErrorCode -> IO a)
          -> (FileHandle -> IO a)
          -> IO a
```

higher-order, or  
first-class functions

- It's as if we've given openFile two return addresses: one to use when an error occurs, and one to use when the call is successful.

20

## Programming with continuations

- Our original program using continuations:

```
openFile "file.txt"
  (\error -> ...)
  (\f -> do l1 <- readLine f
           l2 <- readLine f
           out (l1, l2)
           closeFile f)
```

lambda  
expressions:  
\var -> expr

- Could we do the same for readLine?

21

## Programming with continuations

- Our original program using continuations:

```
openFile "file.txt"
  (\error -> ...)
  (\f -> readLine f
    (\error -> ...)
    (\l1 -> readLine f
      (\error -> ...)
      (\l2 <- do out (l1, l2)
        closeFile f)))
```

- Hmm, not so pretty ...

22

## Programming with continuations

- Name the error handlers:

```
openFile "file.txt"
  err1
  (\f -> readLine f
    err2
    (\l1 -> readLine f
      err3
      (\l2 <- do out (l1, l2)
        closeFile f)))
```

23

## Programming with continuations

- Reformat:

```
openFile "file.txt" err1 (\f ->
  readLine f          err2 (\l1 ->
  readLine f          err3 (\l2 ->
    do out (l1, l2)
      closeFile f)))
```

- Looking better ...

24



## Programming with continuations

- Add an infix operator: `f $ x = f x`

```
openFile "file.txt" err1 $ \f ->
readLine f      err2 $ \l1 ->
readLine f      err3 $ \l2 ->
do out (l1, l2)
  closeFile f
```

- Fewer parentheses ...
- Easier to add or remove individual lines ...
- ... still a little cluttered by error handling behavior

25

## Programming with continuations

- Integrate the error handlers in to the main operations:

```
openFile "file.txt" $ \f ->
readLine f      $ \l1 ->
readLine f      $ \l2 ->
do out (l1, l2)
  closeFile f
```

- Not always applicable ...
- ... but a good choice for HaL4 where the response to a particular type of invalid parameter is always the same (typically, returning an error code to the caller)
- ... and this also ensures a consistent API

26

## “Validators”

The implementation of HaL4 includes a small library of validator functions:

```
getCurrent      :: KR k => (TCBRef -> k a) -> k a
getRegCap       :: KE k => #r -> TCBRef
                -> (CapRef -> k a) -> k a
emptyCapability :: KE k => TCBRef -> CapRef -> k a -> k a
cdtLeaf        :: KE k => TCBRef -> CapRef -> k a -> k a
notMaxDepth    :: KE k => TCBRef -> CapRef -> k a -> k a
untypedCapability :: KE k => TCBRef -> CapRef
                -> (UntypedRef -> k a) -> k a
pageDirCapability :: KE k => TCBRef -> CapRef
                -> (PageDirRef -> PMapData -> k a) -> k a
pageTableCapability :: KE k => TCBRef -> CapRef
                -> (PageTableRef -> MapData -> k a) -> k a
```

27

## “Validators”

- In effect, we have built an embedded domain specific language, just for validating parameters in HaL4
- Benefits include:
  - Ease of reuse
  - Consistency
  - Clarity
  - Ability to pass multiple results on to continuation

28

```
syscallMapPageDir :: (KE k, KW k) => k a
syscallMapPageDir
```

<pre>= getCurrent   getMapPageDirASIDTab curr   asidTableCapability curr asidcap    getMapPageDirOffset curr   asidInRange curr offset range   asidNotUsed curr asid    getMapPageDirPDir curr   pageDirCapability curr pdcap   unmappedPD curr pdmd</pre>	<pre>\$ \curr    -&gt; \$ \asidcap -&gt; \$ \range   -&gt;  \$ \offset  -&gt; \$ \asid    -&gt; \$ \slot    -&gt;  \$ \pdcap   -&gt; \$ \pdir pdmd -&gt; \$</pre>
--	---

Validators

```
do set slot.pdir (Ref pdir)
  set slot.count 1
  setCapdata pdcap MappedPD[asid]
  success curr
```

Action

29

```
syscallMapPageDir :: (KE k, KW k) => k a
syscallMapPageDir
```

```
= getCurrent      $ \curr    ->
  getMapPageDirASIDTab curr  $ \asidcap ->
  asidTableCapability curr asidcap $ \range   ->

  getMapPageDirOffset curr    $ \offset  ->
  asidInRange curr offset range $ \asid    ->
  asidNotUsed curr asid       $ \slot    ->

  getMapPageDirPDir curr      $ \pdcap   ->
  pageDirCapability curr pdcap $ \pdir pdmd ->
  unmappedPD curr pdmd       $
```

```
do set slot.pdir (Ref pdir)
  set slot.count 1
  setCapdata pdcap MappedPD[asid]
  success curr
```

clear and  
concise

30

```
syscallMapPageDir :: (KE k, KW k) => k a
syscallMapPageDir
```

```
=
  getCurrent          $ \curr      ->
  getMapPageDirASIDTab curr $ \asidcap ->
  asidTableCapability curr asidcap $ \range ->

  getMapPageDirOffset curr $ \offset ->
  asidInRange curr offset range $ \asid ->
  asidNotUsed curr asid $ \slot ->

  getMapPageDirPDir curr $ \pdcap ->
  pageDirCapability curr pdcap $ \pdir pdmd ->
  unmappedPD curr pdmd $
```

```
do set slot.pdir (Ref pdir)
    set slot.count 1
    setCapdata pdcap MappedPD[asid]
    success curr
```

reusable

31

```
syscallMapPageDir :: (KE k, KW k) => k a
syscallMapPageDir
```

```
=
  getCurrent          $ \curr      ->
  getMapPageDirASIDTab curr $ \asidcap ->
  asidTableCapability curr asidcap $ \range ->

  getMapPageDirOffset curr $ \offset ->
  asidInRange curr offset range $ \asid ->
  asidNotUsed curr asid $ \slot ->

  getMapPageDirPDir curr $ \pdcap ->
  pageDirCapability curr pdcap $ \pdir pdmd ->
  unmappedPD curr pdmd $
```

```
do set slot.pdir (Ref pdir)
    set slot.count 1
    setCapdata pdcap MappedPD[asid]
    success curr
```

high-level

32

```
syscallMapPageDir :: (KE k, KW k) => k a
syscallMapPageDir
```

```
=
  getCurrent          $ \curr      ->
  getMapPageDirASIDTab curr $ \asidcap ->
  asidTableCapability curr asidcap $ \range ->

  getMapPageDirOffset curr $ \offset ->
  asidInRange curr offset range $ \asid ->
  asidNotUsed curr asid $ \slot ->

  getMapPageDirPDir curr $ \pdcap ->
  pageDirCapability curr pdcap $ \pdir pdmd ->
  unmappedPD curr pdmd $
```

```
do set slot.pdir (Ref pdir)
    set slot.count 1
    setCapdata pdcap MappedPD[asid]
    success curr
```

performance  
concerns?

33

## Other examples: reading/writing registers

```
syscallRtype :: (KE k, RW k) => k a
syscallRtype
  = getCurrent          $ \curr ->
    getRootCap curr     $ \root ->
      untypedCapability curr root $ \ut ->
        cdtLeaf curr root    $
          notMaxDepth curr root $
            getCNodeCap curr   $ \cncap ->
              cnodeCapability curr cncap $ \cnode ->
                getCapArray curr cnode  $ \caps ->
                  allEmptyCapabilities curr caps $
                    validType curr ut caps    $ \objs ->
                      addChildren curr root caps objs
```

34

## Other examples: reading/writing registers

```
syscallReadRegs :: (KE k, RW k) => k a
syscallReadRegs
  = getCurrent          $ \curr -> -- get current thread
    getDstCap curr      $ \cap -> -- get specified capability
      tcbCapability curr cap $ \tcb ps -> -- ... to a tcb object
        permitsRead curr ps.perms $ -- ... that we can read
          transferRegs tcb (>->) curr

syscallWriteRegs :: (KE k, RW k) => k a
syscallWriteRegs
  = getCurrent          $ \curr -> -- get current thread
    getDstCap curr      $ \cap -> -- get specified capability
      tcbCapability curr cap $ \tcb ps -> -- ... to a tcb object
        permitsWrite curr ps.perms $ -- ... that we can write
          transferRegs tcb (<-<) curr
```

$r_1 \leftarrow r_2 = r_2 \rightarrow r_1$

$r_1 \rightarrow r_2$   
 = do x <- readRef r1  
 writeRef r2 x

$r_1 \rightarrow r_2$   
 = readRef r1 >>= writeRef r2

35

## Other examples: transferring registers

```
transferRegs :: KE k => TCBRef
-> (RegRef -> RegRef -> k ())
-> TCBRef -> k a

transferRegs tcb xfer curr
  = do let tcbIframe = tcb.context.iframe
        currRegs = curr.context.regs
        tcbIframe.eip `xfer` currRegs.esi
        tcbIframe.esp `xfer` currRegs.ecx
        case<- getIPCBuffer curr of
          Ref buf -> let tcbRegs = tcb.context.regs
                      tcbRegs.eax `xfer` (buf.mrs @@ 3)
                      tcbRegs.ebx `xfer` (buf.mrs @@ 4)
                      tcbRegs.ecx `xfer` (buf.mrs @@ 5)
                      tcbRegs.edx `xfer` (buf.mrs @@ 6)
                      tcbRegs.esi `xfer` (buf.mrs @@ 7)
                      tcbRegs.edi `xfer` (buf.mrs @@ 8)
                      tcbRegs.ebp `xfer` (buf.mrs @@ 9)
                      success curr
```

<-< or >->  
 would work here!

36

## Other examples: reading/writing registers

```
syscallReadRegs :: (KE k, RW k) => k a
syscallReadRegs
  = getThread          $ \curr -> -- get current thread
    getDstCap curr    $ \cap -> -- get specified capability
    tcbCapability curr cap $ \tcb ps -> -- ... to a tcb object
    permitsRead curr ps.perms $ -- ... that we can read
    transferRegs tcb (>->) curr

syscallWriteRegs :: (KE k, RW k) => k a
syscallWriteRegs
  = getThread          $ \curr -> -- get current thread
    getDstCap curr    $ \cap -> -- get specified capability
    tcbCapability curr cap $ \tcb ps -> -- ... to a tcb object
    permitsWrite curr ps.perms $ -- ... that we can write
    transferRegs tcb (<-<) curr
```

37

## Other examples: reading/writing registers

```
syscallReadRegs :: (KE k, RW k) => k a
syscallReadRegs
  = regAccess          $ \curr tcb ps -> -- find thread
    permitsRead curr ps $ -- check for read perm
    transferRegs tcb (>->) curr

syscallWriteRegs :: (KE k, RW k) => k a
syscallWriteRegs
  = regAccess          $ \curr tcb ps -> -- find thread
    permitsWrite curr ps $ -- check for write perm
    transferRegs tcb (<-<) curr

regAccess :: (KE k) => (TCBRef -> TCBRef -> Perms -> k a) -> k a
regAccess k
  = getThread          $ \curr -> -- get current thread
    getDstCap curr    $ \cap -> -- get specified capability
    tcbCapability curr cap $ \tcb ps -> -- ... to a tcb object
    k curr tcb ps.perms
```

38

## Representing function values

- How should we represent values of type `Int -> Int`?
  - There are many different values, including: `(\z -> x+z)`, `(\x -> x+1)`, `(\x -> x*2)`, `(\x -> f(g(x)))`, ...
  - ... any of which could be passed as arguments to other functions ...
  - ... so we need a uniform, but flexible way to represent them
- A common answer is to represent functions like these by a pointer to a “closure”, a heap allocated object that contains:
  - a code pointer (i.e., the code for the function)
  - the values of its free variables

39

## Closures

- Every function of type `Int -> Int` will be represented using the same basic structure:

codeptr	...
---------	-----

- The code pointer and list of variables vary from one function value to the next:

(\z -> x+z)	codeptr <sub>1</sub>	x	
(\x -> x+1)	codeptr <sub>2</sub>		
(\x -> x*2)	codeptr <sub>3</sub>		
(\x -> f(g(x)))	codeptr <sub>4</sub>	f	g

- To make a closure, allocate a suitably sized block of memory and save the required code pointer and variable values

40

## Calling unknown functions

- If `f` is a known function, then we call it by pushing its arguments on the stack and jumping directly to its code
- What if `f` is an unknown function, represented by a variable that points to a closure structure instead?
- The System V ABI doesn't cover this case, but we can make up our own convention:
  - push the arguments
  - push a pointer to the closure (for access to free variables)
  - call the code that is pointed to at the start of the closure
- This process is known as “entering a closure”

41

## Thunks

- A closure that can be entered without any arguments is sometimes referred to as a “thunk”
  - Thunks represent suspended/delayed computations that can be executed (or “entered” or “invoked”) repeatedly at a later stage.
  - Examples:
    - procedure values / monadic computations
- ```
r1 >-> r2 = do x <- readRef r1
              writeRef r2 x
```
- interrupt/exception/system call handlers
  - ...

42

## Pros/cons of using higher order functions

- Pros:
  - high-level
  - clear and concise
  - facilitates reuse (less code, greater consistency)
- Cons:
  - what impact do all these closures have on performance?
    - it's hard to explore this question without a way to talk about these things more precisely ...
  - in other words, we need a *language* that we can use to express these ideas...

43

## A notation for closures and thunks

- Closures:  $k\{x_1, \dots, x_n\}$ 
  - code pointer:  $k$
  - stored fields:  $x_1, \dots, x_n$
  - If  $f$  is a closure, then we write  $f @ x$  for the result of entering  $f$  with argument  $x$
- Thunks:  $m[x_1, \dots, x_n]$ 
  - code pointer:  $m$
  - stored fields:  $x_1, \dots, x_n$
  - if  $t$  is a thunk, then we write `invoke t` for the result that is obtained by invoking  $t$

44

## Do notation and three address code

- Some uses of do notation look a lot like three address code:

|                               |                         |
|-------------------------------|-------------------------|
| <code>do x &lt;- f y z</code> | <code>x := y + z</code> |
| <code>  t &lt;- g x y</code>  | <code>t := x * y</code> |
| <code>  u &lt;- h t</code>    | <code>u := -t</code>    |
| <code>  p x t</code>          | <code>goto p</code>     |

- Some optimization techniques are like algebraic properties:

`(do x <- return y; c)`    =    `[y/x] c`

- Is this a “monad law” or is it “copy propagation”?

45

## Comprehending Monads

| Types                               |                                                                       |
|-------------------------------------|-----------------------------------------------------------------------|
| $K^*$                               | $= K$                                                                 |
| $(U \rightarrow V)^*$               | $= (U^* \rightarrow M V^*)$                                           |
| $(U, V)^*$                          | $= (U^*, V^*)$                                                        |
| Terms                               |                                                                       |
| $x^*$                               | $= [x]^M$                                                             |
| $(\lambda x \rightarrow v)^*$       | $= [(\lambda x \rightarrow v^*)]^M$                                   |
| $(t \ u)^*$                         | $= [y \mid f \leftarrow t^*, x \leftarrow u^*, y \leftarrow (f x)]^M$ |
| $(u, v)^*$                          | $= [(x, y) \mid x \leftarrow u^*, y \leftarrow v^*]^M$                |
| $(fst \ t)^*$                       | $= [(fst \ z) \mid z \leftarrow t^*]^M$                               |
| Assumptions                         |                                                                       |
| $(x_1 :: T_1, \dots, x_n :: T_n)^*$ | $= x_1 :: T_1^*, \dots, x_n :: T_n^*$                                 |
| Typings                             |                                                                       |
| $(A \vdash t :: T)^*$               | $= A^* \vdash t^* :: M T^*$                                           |

Moggi,  
Wadler,  
Kennedy,  
Benton,  
...  
Bailey

Figure 7: Call-by-value translation.

46

## From source terms ...

```
id      = \x -> x

compose = \f g x -> f (g x)

map     = \f xs ->
          case xs of
            Nil      -> Nil
            Cons y ys -> Cons (f y) (map f ys)
```

47

## ... to MIL programs

```
id      ← k0{ }
k0{ } x = b0(x)
b0(x)   = return x

map      ← k4{ }
k4{ } f  = k5{f}
k5{f} xs = b2(f, xs)
b2(f, xs) = case xs of
              Nil() → b3()
              Cons(y, ys) → b4(f, y, ys)

b3()     = Nil()
b4(f, y, ys) = z ← f @ y
              m ← map @ f
              zs ← m @ ys
              Cons(z, zs)

compose  ← k1{ }
k1{ } f  = k2{f}
k2{f} g  = k3{f, g}
k3{f, g} x = b1(f, g, x)
b1(f, g, x) = u ← g @ x
             f @ u
```

48



## MIL, a monadic intermediate language

```
a  ::= v          -- variable
    | n          -- integer literal

t  ::= return a   -- simple value
    | p((a1, ..., an)) -- primitive call
    | b(a1, ..., an) -- block call
    | C(a1, ..., an) -- data constructor
    | K{a1, ..., an} -- closure constructor
    | f @ a        -- enter closure
    | m[a1, ..., an] -- thunk constructor
    | invoke a     -- invoke closure

c  ::= v <- t; c   -- monadic bind
    | t           -- (generalized) tail call
    | case v of alt1; ...; altn -- conditional branch

alt ::= C(v1, ..., vn) -> b(a1, ..., an) -- constructor match
    | _ -> b(a1, ..., an) -- default match

def ::= b(v1, ..., vn) = c -- block entry point
    | K{v1, ..., vn} v = c -- closure definition
    | v <- t -- top-level constant
```

49

## MIL, a monadic intermediate language

```
a  ::= v          -- variable
    | n          -- integer literal

t  ::= return a   -- simple value
    | p((a1, ..., an)) -- primitive call
    | b(a1, ..., an) -- block call
    | C(a1, ..., an) -- data constructor
    | K{a1, ..., an} -- closure constructor
    | f @ a        -- enter closure
    | m[a1, ..., an] -- thunk constructor
    | invoke a     -- invoke closure

c  ::= v <- t; c   -- monadic bind
    | t           -- (generalized) tail call
    | case v of alt1; ...; altn -- conditional branch

alt ::= C(v1, ..., vn) -> b(a1, ..., an) -- constructor match
    | _ -> b(a1, ..., an) -- default match

def ::= b(v1, ..., vn) = c -- block entry point
    | K{v1, ..., vn} v = c -- closure definition
    | v <- t -- top-level constant
```

Definitions

50

## MIL, a monadic intermediate language

```
a  ::= v          -- variable
    | n          -- integer literal

t  ::= return a   -- simple value
    | p((a1, ..., an)) -- primitive call
    | b(a1, ..., an) -- block call
    | C(a1, ..., an) -- data constructor
    | K{a1, ..., an} -- closure constructor
    | f @ a        -- enter closure
    | m[a1, ..., an] -- thunk constructor
    | invoke a     -- invoke closure

c  ::= v <- t; c -- monadic bind
    | t -- (generalized) tail call
    | case v of alt1; ...; altn -- conditional branch

alt ::= C(v1, ..., vn) -> b(a1, ..., an) -- constructor match
    | _ -> b(a1, ..., an) -- default match

def ::= b(v1, ..., vn) = c -- block entry point
    | K{v1, ..., vn} v = c -- closure definition
    | v <- t -- top-level constant
```

Code

51

## MIL, a monadic intermediate language

```
a  ::= v          -- variable
    | n          -- integer literal

t  ::= return a   -- simple value
    | p((a1, ..., an)) -- primitive call
    | b(a1, ..., an) -- block call
    | C(a1, ..., an) -- data constructor
    | k{a1, ..., an} -- closure constructor
    | f @ a        -- enter closure
    | m[a1, ..., an] -- thunk constructor
    | invoke a     -- invoke closure

c  ::= v <- t; c   -- monadic bind
    | t           -- (generalized) tail call
    | case v of alt1; ...; altn -- conditional branch

alt ::= C(v1, ..., vn) -> b(a1, ..., an) -- constructor match
    | _ -> b(a1, ..., an) -- default match

def ::= b(v1, ..., vn) = c -- block entry point
    | k{v1, ..., vn} v = c -- closure definition
    | v <- t -- top-level constant
```

52

## MIL, a monadic intermediate language

```
a  ::= v          -- variable
    | n          -- integer literal

t  ::= return a   -- simple value
    | p((a1, ..., an)) -- primitive call
    | b(a1, ..., an) -- block call
    | C(a1, ..., an) -- data constructor
    | k{a1, ..., an} -- closure constructor
    | f @ a        -- enter closure
    | m[a1, ..., an] -- thunk constructor
    | invoke a     -- invoke closure

c  ::= v <- t; c   -- monadic bind
    | t           -- (generalized) tail call
    | case v of alt1; ...; altn -- conditional branch

alt ::= C(v1, ..., vn) -> b(a1, ..., an) -- constructor match
    | _ -> b(a1, ..., an) -- default match

def ::= b(v1, ..., vn) = c -- block entry point
    | k{v1, ..., vn} v = c -- closure definition
    | v <- t -- top-level constant
```

53

## ... to MIL programs

```
id      ← k0{ }
k0{ } x = b0(x)
b0(x) = return x

map      ← k4{ }
k4{ } f  = k5{ f }
k5{ f } xs = b2(f, xs)
b2(f, xs) = case xs of
    Nil() → b3()
    Cons(y, ys) → b4(f, y, ys)

b3() = Nil()
b4(f, y, ys) = z ← f @ y
               m ← map @ f
               zs ← m @ ys
               Cons(z, zs)
```

compose ← k<sub>1</sub>{ }  
k<sub>1</sub>{ } f = k<sub>2</sub>{ f }  
k<sub>2</sub>{ f } g = k<sub>3</sub>{ f, g }  
k<sub>3</sub>{ f, g } x = b<sub>1</sub>(f, g, x)  
b<sub>1</sub>(f, g, x) = u ← g @ x  
f @ u

unknown function call

known function call

54

## ... to MIL programs

```
id      ← k0{}  
k0{x}  = b0(x)  
b0(x)  = return x  
  
map      ← k4{}  
k4{f}  = k5{f}  
k5{f} xs = b2(f,xs)  
b2(f,xs) = case xs of  
    Nil() → b3()  
    Cons(y,ys) → b4(f,y,ys)  
  
b3()    = Nil()  
b4(f,y,ys) = z ← f @ y  
            m ← map @ f  
            zs ← m @ ys  
            Cons(z,zs)
```

unknown function call

55

## ... to optimized MIL programs

```
id      ← k0{}  
k0{x}  = b0(x)  
b0(x)  = return x  
  
map      ← k4{}  
k4{f}  = k5{f}  
k5{f} xs = b2(f,xs)  
b2(f,xs) = case xs of  
    Nil() → b3()  
    Cons(y,ys) → b4(f,y,ys)  
  
b3()    = Nil()  
b4(f,y,ys) = z ← f @ y  
            m ← k5{f}  
            zs ← m @ ys  
            Cons(z,zs)
```

unknown function call

known function call

56

## ... to optimized MIL programs

```
id      ← k0{}  
k0{x}  = b0(x)  
b0(x)  = return x  
  
map      ← k4{}  
k4{f}  = k5{f}  
k5{f} xs = b2(f,xs)  
b2(f,xs) = case xs of  
    Nil() → b3()  
    Cons(y,ys) → b4(f,y,ys)  
  
b3()    = Nil()  
b4(f,y,ys) = z ← f @ y  
            m ← k5{f}  
            zs ← m @ ys  
            Cons(z,zs)
```

unknown function call

57

## ... to optimized MIL programs

```
id      ← k0{}  
k0{ } x = b0(x)  
b0(x)  = return x  
  
map      ← k4{}  
k4{ } f  = k5{f}  
k5{f} xs = b2(f,xs)  
b2(f,xs) = case xs of  
    Nil() → b3()  
    Cons(y,ys) → b4(f,y,ys)  
  
b3()    = Nil()  
b4(f,y,ys) = z ← f @ y  
            m ← k5{f} pure, dead code  
            zs ← b2(f,ys)  
            Cons(z,zs)
```

58

## ... to optimized MIL programs

```
id      ← k0{}  
k0{ } x = b0(x)  
b0(x)  = return x  
  
map      ← k4{}  
k4{ } f  = k5{f}  
k5{f} xs = b2(f,xs)  
b2(f,xs) = case xs of  
    Nil() → b3()  
    Cons(y,ys) → b4(f,y,ys)  
  
b3()    = Nil()  
b4(f,y,ys) = z ← f @ y  
            zs ← b2(f,ys)  
            Cons(z,zs)
```

59

## Constant/copy propagation / left monad law:

$x \leftarrow \text{return } a; c \longrightarrow [a/x] c$

## Tail call introduction / right monad law:

$x \leftarrow t; \text{return } x \longrightarrow t$

60

**Prefix inlining:**  $b(x) = v_1 \leftarrow t_1; \dots; v_n \leftarrow t_n; t$   
(monad associativity law)

$v \leftarrow b(x); c \longrightarrow v_1 \leftarrow t_1; \dots; v_n \leftarrow t_n; v \leftarrow t; c$

**Suffix inlining:**

$v \leftarrow t_0; b(x) \longrightarrow v \leftarrow t_0; v_1 \leftarrow t_1; \dots; v_n \leftarrow t_n; t$

61

**Wildcard introduction:**

$v \leftarrow t; c \longrightarrow \_ \leftarrow t; c \quad (v \text{ not free in } c)$

**Dead tail elimination:**

$\_ \leftarrow t; c \longrightarrow c \quad (t \text{ pure})$

62

**Identity laws:**

$\text{or}((x, 0)) \longrightarrow \text{return } x$

**Idempotence:**

$\text{or}((x, x)) \longrightarrow \text{return } x$

**Common subexpression elimination:**

$\{y = t\} \longrightarrow v \leftarrow \text{return } y$   
 $v \leftarrow t$

63

## Constant folding:

(M, N constants)

$\{x=M, y=N\}$   
 $v \leftarrow \text{and}((x, y)) \longrightarrow v \leftarrow \text{return}(M \& N)$

## Argument ordering:

(M, N constants)

$v \leftarrow \text{and}((M, x)) \longrightarrow v \leftarrow \text{and}((x, M))$

## Associative folding:

$(u \& M) \& N = u \& (M \& N)$

$\{v = \text{and}((w, M))\}$   
 $x \leftarrow \text{and}((v, N)) \longrightarrow x \leftarrow \text{and}((w, M \& N))$

64

## Distributive folding:

(M, N, P constants)

$\{v = \text{or}((u, M))\}$   
 $x \leftarrow \text{and}((v, N)) \longrightarrow \begin{array}{l} t \leftarrow \text{and}((u, N)) \\ x \leftarrow \text{or}((t, M \& N)) \end{array} \quad (t \text{ new})$

$(u \mid M) \& N = (u \& N) \mid (M \& N)$

$\{v = \text{or}((u, M)), w = \text{and}((v, N))\}$   
 $x \leftarrow \text{or}((w, P)) \longrightarrow \begin{array}{l} t \leftarrow \text{and}((u, N)) \\ x \leftarrow \text{or}((w, (M \& N) \mid P)) \end{array}$

$((u \mid M) \& N) \mid P = (u \& N) \mid ((M \& N) \mid P)$

65

## Known closure:

$\{f = k\{x\}\}$   
 $v \leftarrow f @ y \longrightarrow v \leftarrow t \quad \text{if } k\{x\} \ y = t$

## Known thunk:

$\{t = m\{x\}\}$   
 $v \leftarrow \text{invoke } t \longrightarrow v \leftarrow m(x)$

## Known constructor:

$\{v = C(x)\}$   
 $\text{case } v \text{ of}$   
     $C(x) \rightarrow b(x, y)$   
     $\dots$   
 $\longrightarrow b(x, y)$

66

## Derived blocks - trailing enter:

$v \leftarrow b(x); v @ a \longrightarrow b1(x, a)$

$b(x) = \begin{matrix} v_1 \leftarrow t_1 \\ \dots \\ v_n \leftarrow t_n \\ t \end{matrix} \longrightarrow b1(x, a) = \begin{matrix} v_1 \leftarrow t_1 \\ \dots \\ v_n \leftarrow t_n \\ v \leftarrow t \\ v @ a \end{matrix}$

67

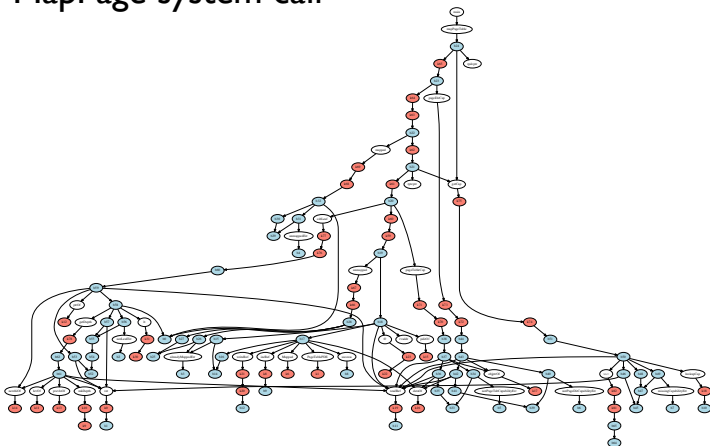
## Derived blocks - known constructors:

$\{u=C(p), f=k\{x\}\} \longrightarrow b1(p, x, r)$   
 $b(u, f, r)$



$b(u, f, r) = c \longrightarrow b1(p, x, r) = \begin{matrix} u \leftarrow C(p) \\ f \leftarrow k\{x\} \\ c \end{matrix}$

68

## MapPage system call



69



# Prioset

```

prioSet = \l prio -> do writeRef (at prioSet l) prio
              writeRef (at prioSet prio) l

insertPriority = \prio -> do s <- readRef prioSetSize
              writeRef prioSetSize (add s 1)
              heapHeapPairUp (mod s 4) prio

heapHeapPairUp = \l prio ->
  case dest l of
    Nothing -> prioSet 0 prio
    Just l' -> do parent <- ref (shift l 1)
              prio <- readRef (at prioSet parent)
              if lt prio prio then
                prioSet l prio
                heapHeapPairUp parent prio
              else
                prioSet l prio

removePriority = \prio ->
  do s <- readRef prioSetSize
  writeRef prioSetSize (sub s 1)
  prio <- readRef (at prioSet (mod s 4) 1)
  if not prio prio then
    i <- readRef (at prioSet 0)
    heapHeapPairDown i prio (mod s 4) 2)
    prio <- readRef (at prioSet i)
    heapHeapPairUp i prio

heapHeapPairUp = \l prio last ->
  do let u = undefined // <- ref (unassigned i)
  case lq (add (mul 2 u) 1) last of
    Nothing -> prioSet 1 prio
    Just l' ->
      do prio <- readRef (at prioSet l)
      case lq (add (mul 2 u) 2) last of
        Nothing ->
          if gt prio prio then
            prioSet l prio
          else
            prioSet l prio
        Just r' ->
          do prio <- readRef (at prioSet r)
          if gt prio prio && gt prio prio then
            prioSet l prio
          else if gt lt prio prio then
            prioSet l prio // left is higher
            heapHeapPairDown l prio last
          else
            prioSet l prio // right is higher
            heapHeapPairDown r prio last

```

72



## Prioret


```

prioret = \l prio -> do writeRef (at prioret l) prio
                    writeRef (at prioret prio) l
insertPriority = \prio -> do s <- readRef prioretSize
                    writeRef prioretSize (add s 1)
                    heapheapairup (modix s) prio

heapheapairup = \l prio ->
  case dec l of
    Nothing -> prioret 0 prio
    Just j -> do parent <- ret (shiftR j 1)
                rprio <- readRef (at prioret parent)
                if lt rprio prio then
                  prioret l rprio
                  heapheapairup parent prio
                else
                  prioret l prio

removePriority = \prio ->
  do s <- readRef prioretSize
  writeRef prioretSize (sub s 1)
  rprio <- readRef (at prioret (modix (sub s 1)))
  if neg prio rprio then
    i <- readRef (at prioret prio)
    heapheapairdown i rprio (modix (sub s 2))
    rprio <- readRef (at prioret i)
    heapheapairup i rprio

heapheapairdown = \l prio last ->
  do let u = unassigned l // <- set (unassigned i)
     case lq (add (mul 2 u) 1) last of
       Nothing -> prioret l prio // i has no children
       Just l -> do lprio <- readRef (at prioret l)
                  case lq (add (mul 2 u) 2) last of
                    Nothing -> // i has no right child
                    Just r -> if lt lprio prio then
                              prioret l lprio
                              prioret l prio
                              else
                              prioret l prio
                              Just r -> // i has two children
                              do rprio <- readRef (at prioret r)
                                 if lt prio lprio & lt prio rprio then
                                   prioret l prio
                                   else if lt prio lprio rprio then
                                   prioret l lprio // left is higher
                                   heapheapairdown l prio last
                                   else
                                   prioret l rprio // right is higher
                                   heapheapairdown r prio last
  
```



73


## Fibonacci

```

fib = \n -> if eq n 0 then 0
           else if eq n 1 then 1
           else add (fib (sub n 1))
                  (fib (sub n 2))

main = fib 12
  
```

initial compilation



74


## Fibonacci

```

fib = \n -> if eq n 0 then 0
           else if eq n 1 then 1
           else add (fib (sub n 1))
                  (fib (sub n 2))

main = fib 12
  
```

47 transformation steps later ...



```

b10(n) =
  t39 <- eq((n, 0))
  case t39 of
    True() -> b5()
    False() -> b9(n)

b9(n) =
  t43 <- eq((n, 1))
  case t43 of
    True() -> b7()
    False() -> b8(n)
  
```

```

b8(n) =
  t47 <- add((n, -1))
  t48 <- b10(t47)
  t52 <- add((n, -2))
  t53 <- b10(t52)
  add((t48, t53))

b7() = return 1
b5() = return 0

main <- b8(12)
  
```

75


Church

```
succ = \n f x -> n f (f x)
plus = \n m -> \f x -> n f (m f x)
mult = \n m -> \f x -> n (m f) x
expn = \n m -> m n

zero = \f x -> x
one = \f x -> f x
two = \f x -> f (f x)
three = succ two
```

```
toInt = \n -> n (add 1) 0

main
= do print (toInt (plus one
                    (mult two three)))
  print (toInt (mult one two))
  print (toInt (expn two three))
  print (toInt (expn three two))
```



initial compilation

Church

```
succ = \n f x -> n f (f x)
plus = \n m -> \f x -> n f (m f x)
mult = \n m -> \f x -> n (m f) x
expn = \n m -> m n


zero = \f x -> x
one = \f x -> f x
two = \f x -> f (f x)
three = succ two
```

```
toInt = \n -> n (add 1) 0

main
= do print (toInt (plus one
                    (mult two three)))
  print (toInt (mult one two))
  print (toInt (expn two three))
  print (toInt (expn three two))
```

```
main <- b4[]

b4() =
  - <- print((7))
  - <- print((2))
  - <- print((8))
  print((9))
```



570 transformation steps later ...

PREPRINT: To be presented at ICFP '14, September 1-4, 2014, Gothenburg, Sweden.

Practical and Effective Higher-Order Optimizations

Lars Bergström  
larsberg@cs.cmu.edu

Matthew Fluet  
Matthew Le  
Research Institute of Technology  
(mit-rit)@cs.cmu.edu

John Reppy  
Nora Smedley  
University of Chicago  
(john.reppy@chicago.edu)

**Abstract**  
Inlining is an optimization that replaces a call to a function with that function's body. This optimization not only reduces the overhead of a function call, but can expose additional optimization opportunities to the compiler, such as removing redundant operations or unrolling conditional branches. Another optimization, copy propagation, replaces a redundant copy of a still live variable with the original. Copy propagation can reduce the total number of flow variables, allowing register pressure and memory usage, and thereby eliminating redundant memory-to-memory copies. In practice, both of these optimizations are implemented as nearly every modern compiler. These two optimizations are practical to implement and effective in free-order languages, but in languages with statically-typed function types, such as Haskell, these optimizations are not available in code generated in higher-order style. With higher-order functions, the analysis challenge lies here: how to transform the code into the same as if the function were inlined, up to the flow variables, or the meaning of the program may change. This paper's 1997 Ph.D. dissertation called this family of optimizations "flow-J" and he proposed one analysis technique, called *flow-J*, to support these optimizations. Unfortunately, *flow-J* has proven too expensive to implement in practice. Because these higher-order optimizations are not available in functional languages, compiler programmers routinely avoid use of higher-order code when they cannot be optimized (particularly in compiler benchmarks). This paper presents the first practical and effective technique for Super-J (higher-order) inlining and copy propagation, which we call *reduced-variables analysis*. We show that this technique is practical by implementing it in the context of a real compiler for an ML-family language and showing that the reported analysis has run times below 7% of the total compilation time. This technique's effectiveness is shown through use of benchmarks and an analysis program, where this analysis exposes additional potential optimization sites.

\*Portions of this work were performed while the author was at the University of Chicago.

**Categories and Subject Descriptors** D.3.1 [Programming Languages]: General; D.3.2 [Programming Languages]: Language Classifications—Application (functional) languages; D.3.4 [Programming Languages]: Processors—Optimization

**Keywords** control flow analysis, inlining, optimization

**1. Introduction**  
All high-level programming languages rely on complex optimizations to transform a language that is convenient for software developers into one that runs efficiently on target hardware. Two such common compiler optimizations are copy propagation and function inlining. Copy propagation is a language like ML is a simple substitution. Given a program of the form:

```
let val x = 7
in x + 2
end
```

We want to propagate the definition of *x* to its uses, resulting in

```
let val x = 7
in x + 2
end
```

At this point, we can eliminate the now unused *x*, resulting in

```
let
in 7 + 2
end
```

This optimization can reduce the memory requirements (i.e., register pressure) of a program, and it may open the possibility for further optimizations.

Inlining replaces a function call with the body of the function with the help of the function by performing straightforward  $\beta$ -substitution. For example, given the program

```
let fun f x = x + 2
in f 7
end
```

Inlining *f* and inlining the inlined definition results in

```
7 + 2
```

This optimization removes the cost of the function call and opens the possibility of further optimizations, such as constant folding. Inlining does require some care, however, since it can increase the program size, which can negatively affect the instruction cache performance, negating the benefits of eliminating and overhead. The importance of inlining for functional languages and techniques for providing practically performance are well-covered in the context of CFC.

Both copy propagation and function inlining have been well-studied and are widely implemented in modern compilers for both first-order and higher-order programming languages. In this paper,

(Copyright notice will appear here once "preprint" option is removed.)

## Practical and Effective Higher-Order Optimizations

Lars Bergstrom  
Mozilla Research \*  
larsberg@mozilla.com

Matthew Fluet  
Matthew Le  
Rochester Institute of Technology  
{mtf,ml9951}@cs.rit.edu

John Reppy  
Nora Sandler  
University of Chicago  
{jhr,nlsandler}@cs.uchicago.edu

### Abstract

Inlining is an optimization that replaces a call to a function with that function's body. This optimization not only reduces the overhead of a function call, but can expose additional optimization opportunities to the compiler, such as removing redundant operations or unused conditional branches. Another optimization, copy propagation, replaces a redundant copy of a still-live variable with the original. Copy propagation can reduce the total number of live variables, reducing register pressure and memory usage, and possibly eliminating redundant memory-to-memory copies. In practice, both of these optimizations are implemented in nearly every modern compiler.

These two optimizations are practical to implement and effective in first-order languages, but in languages with lexically-scoped first-class functions (aka, closures), these optimizations are not available to code programmed in a higher-order style. With higher-order functions, the analysis challenge has been that the environment at the call site must be the same as at the closure capture location, up to the free variables, or the meaning of the program may change. Olin Shivers' 1991 dissertation called this family of

*Categories and Subject Descriptors* D.3.0 [Programming Languages]: General; D.3.2 [Programming Languages]: Language Classifications—Applicative (functional) languages; D.3.4 [Programming Languages]: Processors—Optimization

*Keywords* control-flow analysis, inlining, optimization

### 1. Introduction

All high level programming languages rely on compiler optimizations to transform a language that is convenient for software developers into one that runs efficiently on target hardware. Two such common compiler optimizations are copy propagation and function inlining. Copy propagation in a language like ML is a simple substitution. Given a program of the form:

```
let val x = y
in
  x+2*y
end
```

We want to propagate the definition of  $x$  to its uses, resulting in:

```
let val x = y
in
  x+2*x
end
```

## A "challenging example"

```
fun mk i =
  let
    fun g j = j + i
    fun f (h : int -> int, k) =
      (h (k * i))
  in
    (f, g)
  end
val (f1, g1) = mk 1
val (f2, g2) = mk 2
val res = f1 (g2, 3)
```

```
fun mk i =
  let
    fun g j = j + i
    fun f (h : int -> int, k) =
      ((k * i) + i)
  in
    (f, g)
  end
res <- return 4
val (f1, g1) = mk 1
val (f2, g2) = mk 2
val res = f1 (g2, 3)
```

"While this example is obviously contrived, this situation occurs regularly in idiomatic higher-order programs and the inability to handle the environment problem in general is a limit in most compilers, leading developers to avoid higher-order language features in performance-critical code."

80

## An earlier example in the paper

```
let
  fun emit x = print (Int.toString x)
  fun fact i m k =
    if i=0 then k m
    else fact (i-1) (m*i) k
in
  fact 6 1 emit
end
```



```
let
  fun emit x = print (Int.toString x)
  fun fact i m k =
    if i=0 then emit m
    else fact (i-1) (m*i) emit
in
  fact 6 1 emit
end
```

```
b2() =
  t46 <- intToStr((720))
  print((t46))
```

```
main <-
  b2()
```

81

## ... generalized

```
emit x = print (intToStr x)
fact i m k
  = if i==0
    then k m
    else fact (i-1) (m*i) k
main m n = fact m n emit
```

```
b7(x) =
  t14 <- intToStr((x))
  print((t14))
b10(m, i) =
  t9 <- add((i, -1))
  t12 <- mul((m, i))
  t15 <- eq((t9, 0))
  case t15 of
    True -> b7(t12)
    False -> b10(t12, t9)
b8(i, m) =
  t5 <- eq((i, 0))
  case t5 of
    True -> b7(m)
    False -> b10(m, i)
k36{m} n = b8(m, n)
k37{} m = k36{m}
main <-
  k37{}
```

+ Third arg for fact has gone!

+ Direct call to (inlined) emit

- Some duplication of code ...

## Very good optimization, for low effort?

- 50% chance that I just got lucky with my choice of examples ...
- 50% chance that this is a result of aggressive inlining and specialization
- 1% chance that there is something fundamentally new about this approach to optimization ...
- 101% chance that I don't understand percentages :-)

## Summary

- Back to the three main questions for CEMLaBS:
  - **Feasibility:** Still chipping away ... but getting closer!
  - **Benefit:** Good evidence that we will benefit from the use of functional language features
    - + Types
    - + Higher-order functions
  - **Performance:** acceptable performance may be within reach
    - + We can generate good quality code, even when lambdas are used in fundamental ways
    - + Some code duplication (but, so far, this is entirely tolerable for our specific use case ...)