



## CS 410/510

### Languages & Low-Level Programming

Mark P Jones  
Portland State University

Spring 2016

Week 1: Introduction, Assembly Language


1

## Introduction and Goals

2

### Origins

- For a long time, a group of us at PSU have been looking at the role that high-level programming languages can play in the construction of (very) low-level software.
- By using **high-level languages**, we can hope to increase programmer productivity, and improve software quality
- By focussing on very **low-level software**, we hope to provide strong foundations for the complete software stack



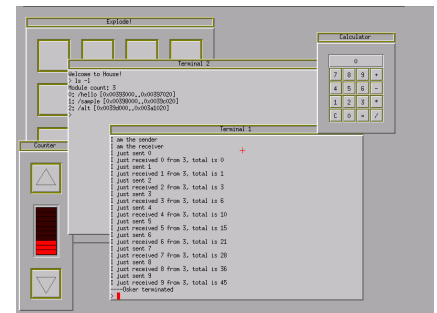
3

### House (2005)

Kernel, GUI, drivers, network stack, and apps

Boots and runs in a bare metal environment

... all written in Haskell, a “purely functional” programming language



4

### Why “House”?

“The Haskell User’s Operating System Environment”

You are more secure in a house ...



than if you only have Windows ...

5


### Performance concerns

- By design, higher-level languages abstract away from the details of how the underlying machine works
- Can we obtain the levels of performance and predictability that are typically required/expected in the systems programming domain?
- Can we write good systems software in a language that intentionally distances users from details of memory layout, representation, instruction selection, alignment, caching, etc.?
- Traditional approaches to building system software resort to using old, low-level languages like assembly and C
- Do “modern” languages have anything to offer in this area?

6


## The Habit programming language

- “a dialect of Haskell that is designed to meet the needs of high assurance systems programming”
- How do you design a programming language for a specific domain?
- Experiment with existing languages
- Understand the domain ...



## The seL4 experience

- In 2009, a group from NICTA, UNSW, and OK Labs in Australia announced seL4, as “the world's first operating system kernel with an end-to-end proof of implementation correctness and security enforcement.”



- A landmark achievement for formal verification, and a strong foundation for building trustworthy systems

8

## seL4 and capabilities

- Even without the verification result, the design of seL4 is interesting in its own right:
  - seL4 is a “capability enhanced” version of an earlier microkernel design called L4
  - The “capability” abstraction in seL4 provides facilities for implementing “least privilege” security policies and novel mechanisms for controlling resource usage

9

## Proof vs performance?

- Security properties established in the seL4 verification include:
  - Absence of code injection attacks
  - Absence of buffer overflows
  - Absence of null pointer dereferences
  - ...
- Many of these properties could be established “for free” if the implementation had been written in a “safer” language
- How might things be different if we built something like seL4 in Habit?

10

## The CEMLaBS project

- “Using a Capability-Enhanced Microkernel as a Testbed for Language-Based Security”
- Started October 2014, funded by The National Science Foundation
- Three main questions:
  - **Feasibility:** Is it possible to build an inherently “unsafe” system like seL4 in a “safe” language like Habit?
  - **Benefit:** What benefits might this have, for example, in reducing verification costs?
  - **Performance:** Is it possible to meet reasonable performance goals for this kind of system?

11

## CEMLaBS workplan

- The original plan was to use Habit to design and build HAL4, a “capability-enhanced” microkernel in the style of seL4 ...
- At the time, neither the specification or the implementation of seL4 was publicly available ... but some key aspects of its design had been revealed in published papers, dissertations, and talks
- In July 2014, seL4 was released as open source software!
- Now we are free to build HAL4 as an implementation of seL4, which will allow use to make more direct comparisons between the two implementations ...

12

## CEMLaBS status

- The original HAL4 implementation is incomplete (and will need to be reworked to match the seL4 specification)
- The Habit compiler is not complete (a full compiler pipeline has been implemented, but it is not feature complete, and it is not robust)
- We'll be working on addressing both of these shortcomings this coming summer
- In the meantime, there is still much to study ...

13

## Course description

- **An overview of conventional low-level programming techniques:**
  - Bare metal programming
  - Fundamental programmable hardware components
- **Case studies of practical microkernel implementations:**
  - OS abstractions (address spaces, threads, capabilities, ...)
  - The L4 and seL4 microkernels
- **Reflections on the design of programming languages for use in this application domain:**
  - Assembly, C, Rust, Habit, domain specific languages, ...

14

## Course learning objectives

Upon the successful completion of this course, students will be able to:

1. Write simple programs that can run in a bare-metal environment using low-level programming languages.
2. Discuss common challenges in low-level systems software development, including debugging in a bare-metal environment.
3. Explain how conventional operating system features (multiple address spaces, context switching, protection, etc.) motivate the desire for (and benefit from) hardware support.

15

## Course learning objectives, continued

4. Develop code to configure and use programmable hardware components such as a memory management unit (MMU), interrupt controller (PIC), and interval timer (PIT).
5. Describe the key steps in a typical boot process, including the role of a bootloader.
6. Describe the motivation, implementation, and application of microkernel abstractions for managing address spaces, threads, and interprocess communication (IPC).
7. Explain the use and implementation of capabilities in access control and resource management.
8. Develop programs using the capability abstraction provided by the seL4 microkernel.

16

## Course learning objectives, continued

9. Illustrate the use of a range of domain specific languages in the development of systems software.
10. Use practical case studies to evaluate and compare language design proposals.
11. Describe features of modern, high-level programming languages—including abstract datatypes and higher-order functions—and show how they can be leveraged in the construction of low-level software.
12. Explain how the requirements of low-level systems programming motivate the desire for (and benefit from) language-based support.

17

## The “programming languages” perspective

- We will survey and evaluate a range of programming languages during this course:
  - Low-level machine and assembly languages
  - Systems programming languages (e.g., C, Rust, ...)
  - Object-oriented languages (e.g., the seL4 API)
  - Domain specific languages
  - Functional languages (e.g., Habit, Haskell, ...)
- What are the driving needs of the systems domain?
- How can a programming language design best meet those needs?

18

## Context

- Basic Platform: Generic “IBM PC” compatible
  - 32 bits ... not 64
  - IA32 ... not x86\_64 or ARM
  - BIOS ... not EFI or UEFI
  - `int` and `iret` ... not `sysenter`/`sysexit`
  - PIC ... not APIC
  - No PAE, PCI, ACPI, MMX, SSE, SMM, SMP, VTx, ...
  - etc., ...
- Already complicated enough for our purposes!
- Well supported by current hardware, emulators, and tools
- Underlying concepts still very broadly applicable

19

## Development environment

- Ubuntu Linux
  - Week 1: using the LinuxLab machines (Mac OS also an option)
  - Weeks 2+: using a VirtualBox virtual machine, preconfigured with appropriate development tools (can be used on Linux, Mac OS, or Windows)
- Bare metal emulation using the QEMU emulator

20

## Rough schedule

Week	Topic
1	Assembly language programming
2	Bare metal programming
3	Hardware support for OS abstractions
4	Memory management & protection
5	Case Study: L4 use & implementation
6	
7	Case Study 2: seL4 use & implementation
8	
9	Language design for low-level programming
10	

21

## Lectures and labs

- Monday Lectures (UTS 210), Wednesday Labs (FAB 88-09)
  - Students are expected to attend all lectures and labs
  - Sets of guided practical exercises will be provided for each lab
  - Students will work on labs during lab sessions with instructor supervision
  - Students are expected to complete each set of exercises before the following week's lab session
  - We will use D2L forums for discussion of exercises outside lab sessions
  - Lab sessions may also be used for direct instruction, as necessary

22

## Exams

- No midterm or final exams
  - Students are still required to attend the “final” on 6/9
  - This will be used for student presentations or demos, or else for other instruction

23

## Portfolio assessment

- Final grades will be assigned on the basis of a “Portfolio” comprising the results of lab work and independent projects
- Based on previously listed collection of 12 course objectives
- General process:
  - Student uploads work item (typically a tar ball or zip file), accompanied by a write up explaining how the submitted work addresses the expectations of one or more objectives
  - Grading assigns a score of N/A, Below, Meets or Exceeds for each objective (with +/- versions)
  - Different final letter grades awarded depending on the combination of individual scores

24

## Portfolio assessment, continued

- Students may request consideration of at most two objectives in any given week (Saturday-Friday)
  - discourages procrastination, balances grading load
- I will attempt to provide feedback on submitted items within a week
- Students may resubmit in an attempt to improve score for any given objective
  - (Each resubmission counts as a new submission)
- Completion of lab exercises for a given week can typically be used to satisfy the “Meets” condition for specific objectives
- Process starts in Week 2, runs until end of finals week

25

## General approach and caveats

- This course will be experimental, open-ended, hands-on, and interactive
- Your flexibility, patience, and tolerance of informal or underspecified aspects of the course will be very much appreciated
- Questions, suggestions, and contributions are welcome at any point

26

## An introduction to IA32 assembly language programming

27

## What is IA32?

- We'll be using the IA32 (x86) architecture as our main target:
  - A “32-bit” instruction set
  - Broadly adopted by:
    - processors from Intel, AMD, Via, ...
    - laptops, desktops, servers, gaming consoles, ...
    - Linux, Mac OS X, Windows, ...
- Arguably, a bit dated ... but still very relevant, and a good platform for learning and exploration
- (... and one of the two architectures supported by seL4)

28

## Other architectures:

- Not to be confused with:
  - x86-64/AMD64: a 64 bit architecture supported (in addition to IA32) by more recent AMD/Intel designs
  - IA64: a completely different 64-bit Intel architecture (Itanium)
  - ARM: widely used in phones, tablets, and more
  - IBM Power: used in Xbox 360, PS3, Wii, servers, and more
  - SPARC: used by some of the college's Unix servers
- Except for x86-64, you can't run IA32 code directly on a machine that uses one of these alternatives instruction sets!

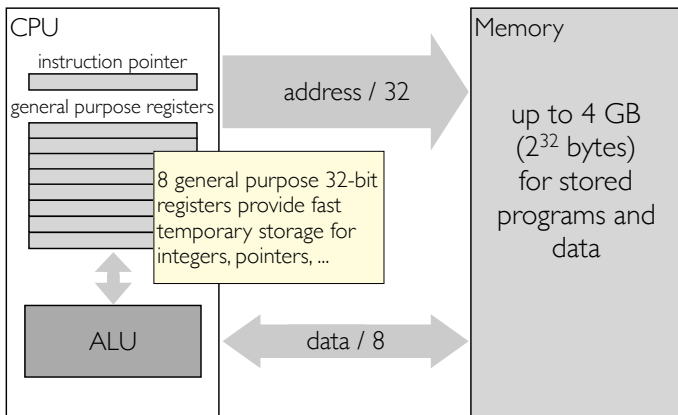
29

## Notes

- No prior or in-depth knowledge of IA32 programming will be assumed
- We will only use a small subset of the full instruction set
- If you're looking to become an expert on IA32 programming, you'll want to look for another class!
- We'll be using the *AT&T syntax* for IA32 assembly language rather than the *Intel syntax*. This is the default syntax used by the free GNU tools in Linux, MacOS, and DJGPP or Cygwin on Windows, and others
- For simplicity, I recommend: [ssh yourid@linuxlab.cs.pdx.edu](ssh://yourid@linuxlab.cs.pdx.edu) (or, on Windows, the equivalent using PuTTY)

30

## A (greatly) simplified view of the IA32



31

## Programming for IA32

- In concrete terms, an IA32 program is just a collection of byte values (*machine code*)
- Once it has been loaded in to memory, the processor can execute a program by interpreting the byte values as *instructions* for the processor to act on
- For practical purposes, we will usually write IA32 programs in a textual format called *assembly language* that is easier to read than the raw byte values
- The program that translates assembly language programs in to machine code is called an *assembler*

32

## The GNU assembler, as

- Assembly code goes in files with a .s suffix
- We will typically use gcc to invoke the assembler
 

```
gcc -m32 -o output assemblyCode.s extras.c
```
- You can also invoke the assembler directly: detailed documentation is available from:
 

```
http://sourceware.org/binutils/docs/as/
```

 For IA32 programming, look in particular at the section on "80386 Dependent Features"

33

## An assembly code listing

```

.global f
f:
    pushl   %ebp
    movl    %esp,%ebp
    pushl   %ebx
    movl    8(%ebp), %ebx

    movl    $0, %eax    # initialize length count in eax

loop:
    jmp     test
    incl    %eax        # increment count
    addl    $4, %ebx     # and move to next array element

test:
    movl    (%ebx), %ecx # load array element
    cmpl    $0, %ecx    # test for end of array
    jne     loop        # repeat if we're not done ...

    popl    %ebx
    movl    %ebp,%esp
    popl    %ebp
    ret
  
```

Assembly code

34

## An assembly code listing

Machine code	Assembly code
0000 55	f: .globl f
0001 89E5	pushl %ebp
0003 53	movl %esp,%ebp
0004 8B5D08	pushl %ebx
	movl 8(%ebp), %ebx
0007 B8000000	movl \$0, %eax # initialize length count in eax
0008 00	
000c EB04	jmp test
000e 40	loop: incl %eax # increment count
000f 83C304	addl \$4, %ebx # and move to next array element
0012 8B0B	test: movl (%ebx), %ecx # load array element
0014 83F900	cmpl \$0, %ecx # test for end of array
0017 75F5	jne loop # repeat if we're not done ...
0019 5B	popl %ebx
001a 89EC	movl %ebp,%esp
001c 5D	popl %ebp
001d C3	ret

35

addresses  
/offsets

labels


addresses /offsets	labels	directive	comments
0000 55	f:	.globl f	
0001 89E5		pushl %ebp	
0003 53		movl %esp,%ebp	
0004 8B5D08		pushl %ebx	
		movl 8(%ebp), %ebx	
0007 B8000000		movl \$0, %eax	# initialize length count in eax
0008 00			
000c EB04		jmp test	
000e 40	loop:	incl %eax	# increment count
000f 83C304		addl \$4, %ebx	# and move to next array element
0012 8B0B	test:	movl (%ebx), %ecx	# load array element
0014 83F900		cmpl \$0, %ecx	# test for end of array
0017 75F5		jne loop	# repeat if we're not done ...
0019 5B		popl %ebx	
001a 89EC		movl %ebp,%esp	
001c 5D		popl %ebp	
001d C3		ret	
machine code		instructions	

36

## IA32 registers

37


## 8-bit registers (holding a single byte, 0-255)



Introduced in 1978 as part of the 8086 architecture

38


## 16-bit registers (“word”)



Introduced in 1978 as part of the 8086 architecture

39

## 32-bit registers (“double word”)



“e” for extended

sometimes referred to as “long word”s

Introduced in 1985 as part of the 80386 architecture

40

## Special vs. general purpose registers

- `eip`: the instruction pointer register
- `esp`: the stack pointer register
- `eflags`: the flags register, stores information about the results of the most recent arithmetic or logic instruction
- Other registers can typically be used for any purpose (although some instructions—division, for example—work only with specific registers)

41

## IA32 instructions

42

## Instruction format

- A typical IA32 instruction has the form:



- A suffix on the opcode indicates the size of the data that is being operated on:
  - 32-bit values use the suffix **l**(ong)
  - 16-bit values use the suffix **w**(ord)
  - 8-bit values use the suffix **b**(yte)

43

## Addressing modes

- **Register access**, reg:
  - `%eax`: the value in register `eax`
  - Can typically use any registers except `eip` and `eflags`
- **Memory access**, mem:
  - `var`: the value in the memory location at address `var`
  - `( %eax )`: the value in memory at the address in `eax`
  - `8 ( %eax )`: the value in memory at the address given by adding 8 to the value in `eax`
- **Immediate**, immed:
  - `$42`: the constant value 42 (decimal; use `$0x2A` for hex)
  - `$var`: the address of memory location `var`

44

## Directives for “declaring” variables

```
.data                # put variables in the "data" section
                    # (code usually goes in .text)

myvar: .align 4      # make sure address is multiple of 4
      .long 42       # Simple variable, initialized to 42

.global days         # A globally accessible array of ints
days: .long 31, 28, 31, 30, 30, 30
      .long 31, 31, 30, 31, 30, 31


scratch: .space 4*100 # reserve uninitialized space

medium: .long 123    # a 32-bit integer (takes 4 bytes)
regular: .short 123  # a 16-bit integer (takes 2 bytes)
small: .byte 123     # an 8-bit integer (takes 1 byte)
```

45

## How values are stored in memory

- A double word holds 32 binary digits (“bits”) (i.e., 4 bytes)



- `0xBE1A3910` can be interpreted as -1,105,577,712 (signed) or 3,189,389,584 (unsigned)

- Stored in memory with the least significant byte at the lowest address (“little endian”):

stored byte	0x10	0x39	0x1A	0xBE
address	400	401	402	403

46

## IA32 instructions: data movement

47

## Move instructions

- Copy data from a source to a destination (where X is one of the size suffixes: b,w,l):

`movX src, dst`

- Any of the following combinations of arguments is allowed:

`movX reg, (reg | mem)`

`movX mem, reg`

`movX immed, (reg | mem)`

- Note that you can’t move mem to mem in one instruction

48



## Examples

Suppose that the memory (starting at address 0) contains the following (four byte) values:

8	6	2	8	0	2	4	1	7	3	4	5	6
0	4	8	12	16	20	24	28	32	36	40	44	48


Then

instruction	contents of <code>eax</code>
<code>movl \$12, %eax</code>	12
<code>movl (%eax), %eax</code>	8
<code>movl 8(%eax), %eax</code>	0

49


## Zero and sign-extension

- Suppose we want to copy a value from a 16-bit register in to a 32-bit register



- Two common strategies:

- Zero extension: for unsigned values



- Sign extension: for signed values



50

## Move with sign, move with zero extension

- Copy from source to larger destination with sign extension:

`movsFT src, dst`

- Copy from source to larger destination with zero extension:

`movzFT src, dst`

- F and T are the “from” and “to” sizes (either b, w, or l)

- Valid combinations: bw, bl, or wl

- Examples:

```
movsbw %al, %dx    # byte to word
movzwl %ax, %edx   # word to long
```

51

## Scaled indexed addressing

- `[base]([reg1], reg2[, index])`

a memory operand whose address is the value in `reg1`, plus the specified base constant, plus the value of `reg2` times the index (which must be 1, 2, 4, or 8)

- Any of the parts in [...] can be omitted

- Examples:

```
(eax, ebx, 4)  the ebxth element in the array of 32-bit
                words starting at the address in eax
days(, ebx, 4) the ebxth element in the array of 32-bit
                words starting at the address days
```

52

## More examples

Suppose that the memory (starting at address 0) contains the following (four byte) values:

8	6	2	8	0	2	4	1	7	3	4	5	6
0	4	8	12	16	20	24	28	32	36	40	44	48

Then

instruction	<code>eax</code>	<code>ebx</code>
<code>movl \$12, %eax</code>	12	
<code>movl 8(%eax), %ebx</code>	12	2
<code>movl 12(%eax, %ebx, 4), %eax</code>	7	2

53

## The lea (load effective address) instruction

- Load the address of the source operand (must be memory) to a destination (where X is one of the size suffixes: b, w, l):

`leax src, dst`

- Can also be used to co-opt the addressing mode circuitry into performing arithmetic operations:

```
leal 4(%eax), %eax    # eax += 4
leal 1(%eax, %eax, 2), %eax # eax = 3*eax + 1
leal 1(%eax, %eax), %eax # eax = 2*eax + 1
leal 4(, %eax, 8), %eax # eax = 8*eax + 4
```

- These instructions just do an address calculation and do not attempt to read the data at that address.

54

## The exchange instruction

- Exchange data between two locations

`xchgX (reg | mem), reg`

- Consider the following instructions in a high-level language:

```
int tmp = x;
x       = y;
y       = tmp;
```

- If `x` and `y` are held in registers, then a “clever enough” compiler can translate this code into a single `xchgl` instruction

55

## The instruction pointer, `eip`

- The `eip` register holds the address of the next instruction to be executed
- As the processor reads each instruction, it increments the value in `eip` by the appropriate number of bytes to point to the following instruction
- This mechanism allows the processor to execute a sequence of instructions stored in contiguous locations in memory
- What would happen if we “move” a different value in to `eip`?

56

## Jumping and labels

- We can transfer control and start executing instructions at address `addr` by using a jump instruction

`jmp addr`

- Labels can be attached to instructions in an assembly language program:

```
    jmp b
a:   jmp c
b:   jmp a
c:   ...
```

- Modern, pipelined machines work well with sequences of instructions that appear in consecutive locations. Jumps can be expensive: one of the goals of an optimizing compiler is to avoid unnecessary jumps.

57

## IA32 instructions: arithmetic and logic operations

58

## Arithmetic instructions

- Combine a given `src` with a given `dst` value and leave the result in `dst`:

```
    addX  src, dst } integer arithmetic
    subX  src, dst } (signed)
    imulX src, dst }
    andX  src, dst } bitwise arithmetic
    orX   src, dst }
    xorX  src, dst }
```

- Similar to `dst += src`, `dst -= src`, etc.. in C/C++

59

## Examples

- To compute  $x^2 + y^2$  and store the result in `z`:

```
    movl  x, %eax
    imull %eax, %eax
    movl  y, %ebx
    imull %ebx, %ebx
    addl  %ebx, %eax
    movl  %eax, z
```

register	contents
<code>eax</code>	$x^2 + y^2$
<code>ebx</code>	$y^2$

```
    .data
x:   .long 4
y:   .long 3
z:   .long 0
```


60

## IA32 instructions: conditional execution

61

## Flags

- In addition to performing the required operation, arithmetic instructions also change bits in the `eflags` register



- The flags record details about the last operation, such as:
  - Was the result zero?
  - Was the result positive?
  - Did a carry occur?
  - etc...

62

## Conditional jumps, `jCC`

We can test these flags in *conditional jump* instructions

```

jz  addr    (jump to addr if the zero flag is set)
jnz addr    (jump to addr if the zero flag is not set)
je  addr    (jump to addr if equal; same as jz)
jne addr    (jump to addr if not equal; same as jnz)
jl  addr    (jump to addr if less than)
jnl addr    (jump to addr if not less than)
jg  addr    (jump to addr if greater than)
jng addr    (jump to addr if not greater than)
...
    
```

(signed)

63

## Examples

```

subl  %eax,%ebx    jump to addr if ebx = eax
jz     addr
subl  %eax,%ebx    jump to addr if ebx ≠ eax
jnz     addr
subl  %eax,%ebx    jump to addr if ebx < eax
jl     addr
subl  %eax,%ebx    jump to addr if ebx ≥ eax
jnl     addr
    
```

If the specified condition does not apply, then execution just continues with the next instruction ...

64

## The compare instruction

- The `cmpX` instruction behaves like `subX` except that the result is not saved; only the flags are changed

For example: `cmpl %eax,%ebx`  
`j1 addr`

will jump to `addr` if the value in `ebx` is less than the value in `eax`, but it will not change the values in either register

65

## Other conditional instructions

- There are some other instructions that perform an action based on the conditional flags without the cost of a jump
- `setCC reg8` sets the value in a specified 8-bit register to 0 or 1, based on the condition specified by CC:

```

cmpl  %ecx,%ebx    # set eax to 1 if
setl  %al           # ebx < ecx, or
movzbl %al,%eax    # else to 0
    
```

- `cmovCC src, dst` copies data from the specified `src` to `dst`, but only if the condition specified by CC holds:

```

cmpl  %ebx,%eax    # set eax to the max of
cmovl %ebx,%eax    # eax and ebx
    
```

condition code; no size suffix here!

66

## IA32 instructions: more arithmetic

67

## Unary operations

- The following arithmetic operations have only one argument (which serves as both source and destination)

► <b>negX</b>	(reg   mem)	negate
<b>notX</b>	(reg   mem)	complement
<b>incX</b>	(reg   mem)	increment
<b>decX</b>	(reg   mem)	decrement

- Like the binary operators, these instructions also set the flags for subsequent testing

68

## Bitwise shift operations

- Shift operations are handled using instructions of the form:

**op** count, (reg | mem)

**shl/sal** ← 0 shift (logical/arithmetic) left

0 → **shr** → **cf** shift logical right

→ **sar** → **cf** shift arithmetic right

- count is either a constant or else the `%cl` register.
- In all cases, the count value will be masked to 31 bits.

69

## Example

- Given two 32 bit input values:

• base:

• limit:

(Each box is one nibble (4 bits),  
least significant bits on the right)

- Calculate a 64 bit descriptor:

- (Needed for the calculation of “GDT entries”)

70

## Example

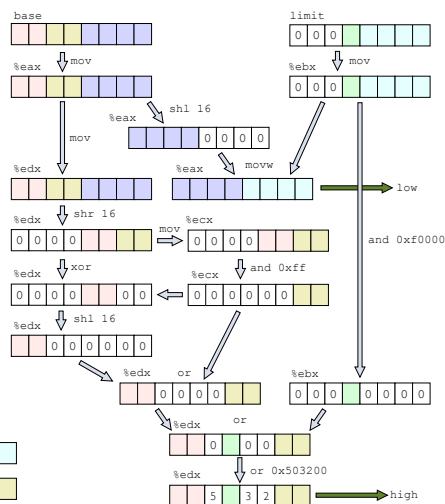
```
movl base, %eax
movl limit, %ebx
```

```
mov %eax, %edx
shl $16, %eax
mov %bx, %ax
movl %eax, low

shr $16, %edx
mov %edx, %ecx
andl $0xff, %ecx
xorl %ecx, %edx
shl $16, %edx
orl %ecx, %edx
andl $0xf0000, %ebx
orl %ebx, %edx
orl $0x503200, %edx
movl %edx, high
```

low


high



71

## Bitwise rotate operations

- Rotate operations use the same instruction format:



- [Aside: Curiously, “higher level” languages often include shift operators, but not rotates, even though the latter have more interesting/uniform behavior ...]


72

## Division

- Divide implicit destination (`edx:eax`) (a 64-bit quantity) by a specified argument with result in `eax` and remainder in `edx`

`idivl (reg | mem)`

- Often used in conjunction with the `cld` instruction (“convert long to double”, a.k.a. `cdq`), which converts a signed 32-bit value in `eax` into the corresponding signed 64-bit value in `edx:eax`.



73

## Example 1

Divide 4,660 (i.e., 0x1234) by 25:

```
movl    $0x1234, %eax
cld
movl    $25, %ecx
idivl   %ecx
```

Results: `eax` = 0xBA (186)  
`edx` = 0xA (10)

Sure enough:  $186 * 25 + 10 = 4,660$

74

## Example 2

Divide -1,105,577,712 (i.e., 0xBE1A3910) by 256

```
movl    $0xBE1A3910, %eax
cld
movl    $256, %ecx
idivl   %ecx
```

Results: `eax` = 0xFFBE1A3A (-4,318,662)  
`edx` = 0xffffffff10 (-240)

Sure enough:  $-4,318,662 * 256 - 240 = -1,105,577,712$

75

## Complications of division

- Division produces multiple results: a quotient and a remainder
- Division uses special registers: we'd better not store any other values in `eax` or `edx` if there's a chance that a division instruction might be executed
- Doesn't set flags: requires separate tests, for example, to determine whether quotient or remainder was zero
- Division can raise an exception if the src is zero (or -1)

76

## IA32 instructions: using the stack

77

## Stack

- The IA32 includes features that allow the programmer to use a region of memory as a simple stack:
  - the `esp` (stack pointer) register
  - special instructions like `push`, `pop`, `call`, `ret`, ...
- There is no obligation for the programmer to use these features, but it is often convenient to do so:
  - for temporary/scratch storage when a calculation needs more storage than the CPU registers can provide
  - to support calling and returning from functions

78

## A typical memory layout

- A typical operating system reserves an area of scratch memory for each program, and sets the `esp` register to point to the end of this region when the program begins



- The stack pointer moves
  - down (decreases) as values are pushed on to the stack
  - up (increases) as values are popped off of the stack
- So long as they never overlap, the data and stack areas can grow or shrink as necessary as the program runs

79

## Stack operations

- Push a value onto the stack

`pushl (reg | mem | immed)`

- Pop a value of the stack

`popl (reg | mem)`

- Roughly speaking:

```
pushl src  =  subl $4, %esp;  movl src, (%esp)
popl  dst   =  movl (%esp), dst; addl $4, %esp
```

80

## Spilling temporaries on the stack

- The stack is often used for saving the contents of a register on the stack (“spilling”) so that the register can be used, temporarily, for some other reason

- For example:

```
pushl %eax
pushl %edx
... code that changes eax and/or edx ...
popl  %edx
popl  %eax
```

pop values in reverse order  
that was used to push them!

- Note that values on the stack can still be accessed, from memory, using `(%esp)`, `4(%esp)`, `8(%esp)`, `12(%esp)`, ...

81

## Call and return

- There is a special instruction for calling a function

```
call addr      =      pushl $lab
                  jmp      addr
lab: ...
```

- And a special instruction for returning from a function

```
ret            =      popl  %eax ←
                  jmp    *%eax
```

assuming  
eax isn't being  
used for  
something else  
...

- In practice, additional instructions are often needed to deal with parameter passing, etc. ...

special syntax: jump  
to the address given  
by the contents of  
eax

82

## Functions and the System V ABI

83

## Implementing functions


- How do we pass arguments to a function?
- How does a function return a result?
- How do we handle local variables?
- In principle, especially in a bare metal setting, we can implement these features any way we like, using the basic tools that the IA32 instruction set provides
- But there are some existing standards we can follow, notably the “System V IA32 Application Binary Interface (ABI)”:

<http://www.sco.com/developers/devspecs/abi386-4.pdf>  
particularly Section 3-9

84

## Stack frames

The code for any given function/procedure call runs in the context of a stack frame of the form:



- **Frame (base) pointer:** ebp points to the stack frame; the caller's frame pointer is stored in old (i.e., ( %ebp ))
- **Return address:** retn is the return address
- **Actual parameters:**  $a_1, \dots, a_n$  are the function's arguments. We can access  $a_1$  as 8 ( %ebp ), etc...
- **Local variables:**  $l_1, \dots, l_m$  are the function's local variables. We can access  $l_1$  as -4 ( %ebp ), etc...

85


## Building the stack frame ... in the caller

- |     |     |     |
|-----|-----|-----|
|     | ... |     |
| esp |     | ebp |

- The caller starts by pushing the arguments:



- Then it executes a `call` instruction, which pushes the return address:




- ... and jumps to the code for the callee ...

86


## Building the stack frame ... in the callee

- |     |      |       |     |       |     |     |
|-----|------|-------|-----|-------|-----|-----|
|     | retn | $a_1$ | ... | $a_n$ | ... |     |
| esp |      |       |     |       |     | ebp |

- The callee saves the old frame pointer, and sets a new value:



- Then it decrements the stack pointer to reserve space for any local variables:



- ... and now the callee can start work ...

87

## Function prologue

- The code that builds the stack frame at the start of a function body is called the prologue:

- At the beginning of a function body, the parameters and return address have already been pushed on to the stack.

We need to:

```
pushl %ebp      # save old frame pointer
movl  %esp, %ebp # and set new value
```

- If local variables taking M bytes of storage are required, then we need to reserve space for them:

```
subl  $M, %esp # allocate space for
              # locals (skip if M=0)
```

88

## Function epilogue

- When a function completes, we must dismantle the stack frame and return the machine to the state it was in before the call. The code to do this is called the epilogue:

- Running the previous process in reverse:

```
movl  %ebp, %esp # discard locals/temps
popl  %ebp      # restore frame pointer
ret    # return to caller
```

- The first two instructions here can be replaced with the more efficient, but otherwise equivalent `leave` instruction

89

## Removing the parameters

- Once we return to the caller, the result of the function is in `eax`, but the parameters are still on the stack:



- We restore the stack pointer to its original value by adding on the number of bytes that are used by the parameters:

```
addl  $N, %esp
```

- If no parameters were passed, then this step can be omitted

90

## Example: a leaf function

```
int g(int u) {
    return u*u;
}
```

→

```
g: pushl %ebp
   movl %esp, %ebp
   movl 8(%ebp), %eax
   imull %eax, %eax
   movl %ebp, %esp
   popl %ebp
   ret
```

...	old	retn	u	...
esp=ebp	ebp+4	ebp+8		

91

## Example: multiple parameters + call

```
int f(int x,
      int y,
      int z) {
    return g(x+y);
}
```

→

```
f: pushl %ebp
   movl %esp, %ebp
   movl 8(%ebp), %eax
   addl 12(%ebp), %eax
   pushl %eax
   call g
   addl $4, %esp
   movl %ebp, %esp
   popl %ebp
   ret
```

...	old	retn	x	y	z	...
esp=ebp	ebp+4	ebp+8	ebp+12	ebp+16		

92

## Example: spilling

```
int h(int x,
      int y,
      int z) {
    return g(x)+g(y);
}
```

→

```
h: pushl %ebp
   movl %esp, %ebp
   pushl 8(%ebp)
   call g
   addl $4, %esp
   pushl %eax -- spill
   pushl 12(%ebp)
   call g
   addl $4, %esp
   popl %ecx -- unspill
   addl %ecx, %eax
   movl %ebp, %esp
   popl %ebp
   ret
```

spill	old	retn	x	y	z
esp	ebp	ebp+4	ebp+8	ebp+12	ebp+16

93

## Observations

- There is a four instruction overhead for each function that uses the frame pointer
  - Increases execution time
  - Prevents use of ebp as a general purpose register
- For larger functions, the four instruction overhead is less of an issue
- For small functions, we would prefer to inline rather than copy
- Nevertheless, it is common to produce code that doesn't use ebp as a frame pointer (e.g., `-fomit-frame-pointer` in gcc)

94

## Caller and callee saves

We (System V) can designate some registers as:

- **caller saves** (eax, ecx, and edx)
  - can be freely used by the callee
  - the caller is responsible for saving (and later restoring) the value of a caller save register before a call
- **callee saves** (ebp, ebx, esi, and edi)
  - can be freely used by the caller
  - the callee is responsible for saving (and later restoring) the value of a callee saves register before using it to store temporary values

95

## Revisiting the previous example: h

```
int h(int x,
      int y,
      int z) {
    return g(x)+g(y);
}
```

→

```
h: pushl %ebp
   movl %esp, %ebp
   pushl 8(%ebp) -- x
   call g
   addl $4, %esp
   pushl %eax -- spill
   pushl 12(%ebp) -- y
   call g
   addl $4, %esp
   popl %ecx -- unspill
   addl %ecx, %eax
   movl %ebp, %esp
   popl %ebp
   ret
```

instead of having the compiler save this value on the stack ...

96



## Revisiting the previous example: h

```

int h(int x,
      int y,
      int z) {
    return g(x)+g(y);
}

```

→

```

h: pushl %ebp
   movl %esp, %ebp

   pushl 8(%ebp) -- x
   call g
   addl $4, %esp
   movl %eax, %esi
   pushl 12(%ebp) -- y
   call g
   addl $4, %esp
   addl %esi, %eax
   movl %ebp, %esp
   popl %ebp
   ret

```

... we can move it to a callee saves register, esi

g will preserve the value in esi, if necessary

so it will still contain the correct value here...

97

## Revisiting the previous example: h

```

int h(int x,
      int y,
      int z) {
    return g(x)+g(y);
}

```

→

```

h: pushl %ebp
   movl %esp, %ebp
   pushl %esi
   pushl 8(%ebp) -- x
   call g
   addl $4, %esp
   movl %eax, %esi
   pushl 12(%ebp) -- y
   call g
   addl $4, %esp
   addl %esi, %eax
   popl %esi
   movl %ebp, %esp
   popl %ebp
   ret

```

... now h has to save the value in register, esi

one save in h is better than one saves in each of two calls to g

empirically, more than 50% of calls are to leaf functions

98

## Closing thoughts

99

## Assembly “Language”?

- Highly imperative, primitive instructions, no expressions
- No high-level abstractions, but all the building blocks:
  - No arrays, records, variants, objects, closures, ...
  - No loops, switch statements, functions, local variables, ...
- Type System?
  - Values classified by size (e.g., 8 vs 32 bits) and storage class (e.g., memory, flag, integer register, floating point register, ...)
  - Limited protection against common programming mistakes
  - Programmer has full control over data representation

100

## Macros

```

.macro idtcalc handler, slot, dpl=0, type=IDT_INTR, seg=KERN_CS
mov    $seg, %eax           # eax = ? # seg
shl    $16, %eax            # eax = seg # 0
movl   $handler, %edx       # edx = hhi # hlo
mov    %dx, %ax            # eax = seg # hlo
mov    $(0x8e00 | (\dpl<<13) | \type), %dx
movl   %eax, idt + ( 8*\slot)
movl   %edx, idt + (4 + 8*\slot)
.endm

```

macro name

parameters

defaults

loop constructs

```

initIDT: .irp num, 0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,16,17,18,19
idtcalc exc\num, slot=\num
.endr

idtcalc handler=syscall, slot=0x80, dpl=3
lidt   idtptr
ret

```

macro invocation

Macros support compile-time code generation:

- Macros are like “used-defined instructions”
- Also supported: conditional code generation, loops, ...
- One way to compensate for language weaknesses?

101

## Summary

- IA32 provides a very basic programming language:
  - A fixed set of registers
  - Instructions for moving and operating on data
  - Instructions for testing and control transfer
- In programming language terms:
  - Low-level, primitive instructions, loosely typed
  - No high-level abstractions, but all the building blocks
  - Very close to the metal, low-level control, “predictable” performance
- Let’s write some programs!

102