



# CS 410/510

## Languages & Low-Level Programming

Mark P Jones  
Portland State University

Spring 2016

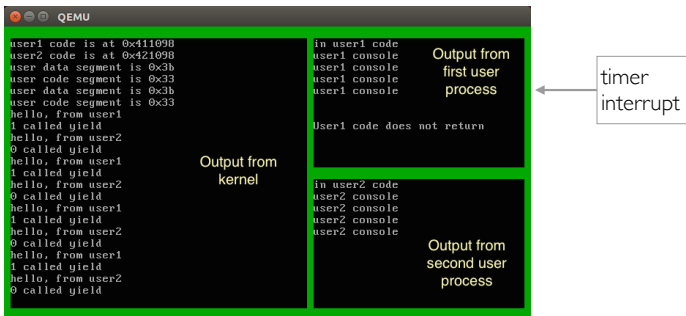
Week 4: Memory Management

1

## Loose Ends

2

## The Week 3 Lab: Context Switching

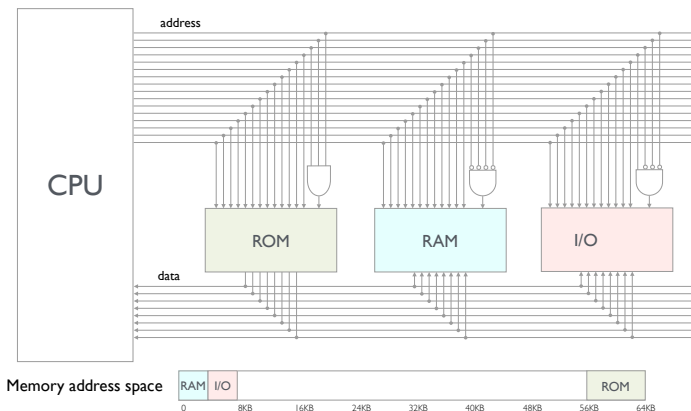


3


## Port I/O

4


## Memory mapped I/O



## Memory mapped I/O



## Port I/O



## Port I/O in the IA32 instruction set

- The IA32 has a 16 bit I/O Port address space
- The hardware actually uses the same address bus and data bus with a signal to indicate whether memory or port access is being used
- You can write a byte/short/word to an I/O port using:  

```
out[b|w|l] [%al,%ax,%eax], [imm8|%dx]
```

 (use imm8 for 8 bit port numbers, otherwise use %dx)
- You can read a byte/short/word from an I/O port using:  

```
in[b|w|l] [imm8|%dx], [%al,%ax,%eax]
```

8

## Port I/O using gcc inline assembly

```
static inline void outb(short port, byte b) {
    asm volatile("outb %1, %0\n" : : "dN"(port), "a"(b));
}
```

```
static inline byte inb(short port) {
    unsigned char b;
    asm volatile("inb %1, %0\n" : "=a"(b) : "dN"(port));
    return b;
}
```

- Arcane syntax, general form:  

```
asm ( template : output operands : input operands : clobbered registers );
```
- Operand constraints include:
  - "d" (use %edx), "a" (use %eax), "N" (imm8 constant),
  - "=" (write only), "r" (register), ...

9

## The role of inline assembly

- We can already call assembly code from C and vice versa by following calling conventions like the System V ABI
- Inline assembly allows for even tighter integration between C and assembly code: code can be inlined, can have an impact on register allocation, etc...
- But there is essentially no checking of the arguments: it's up to the programmer to specify the correct list of clobbered registers to ensure correct semantics
- Programmers might want to check the generated code ...
- How can a general language provide access to essential machine specific instructions and registers?

10

## Standard port numbers on the PC platform

Port Range	Device
0x00-0x1f	First DMA controller (8237)
0x20-0x3f	Master Programmable Interrupt Controller (PIC) (8259A)
0x40-0x5f	Programmable Interval Timer (PIT) (8253/8254)
0x60-0x6f	Keyboard (8042)
0x70-0x7f	Real Time Clock (RTC)
0x80-0x9f	DMA ports, Refresh
0xa0-0xbf	Second Programmable Interrupt Controller (PIC) (8259A)
0xc0-0xdf	Second DMA controller (8237)
...	...
0x3f0-0x3f7	Primary floppy disk drive controller
0x3f8-0x3ff	Serial Port I
...	...

11

## Serial port output in assembly

```
.set PORTCOM1, 0x3f8
serial_putc:
    pushl %eax
    pushl %edx

1:    movw $(PORTCOM1+5), %dx
    inb %dx, %al          # Wait for port to be ready
    andb $0x60, %al
    jz 1b
    movw PORTCOM1, %dx    # Output the character
    movb 12(%esp), %al
    outb %al, %dx

    cmpb $0xa, %al        # Was it a newline?
    jnz 2f

    movw $(PORTCOM1+5), %dx
    inb %dx, %al          # Wait again for port to be ready
    andb $0x60, %al
    jz 1b
    movw PORTCOM1, %dx    # Send a carriage return
    movb $0xd, %al
    outb %al, %dx

2:    popl %edx
    popl %eax
    ret
```

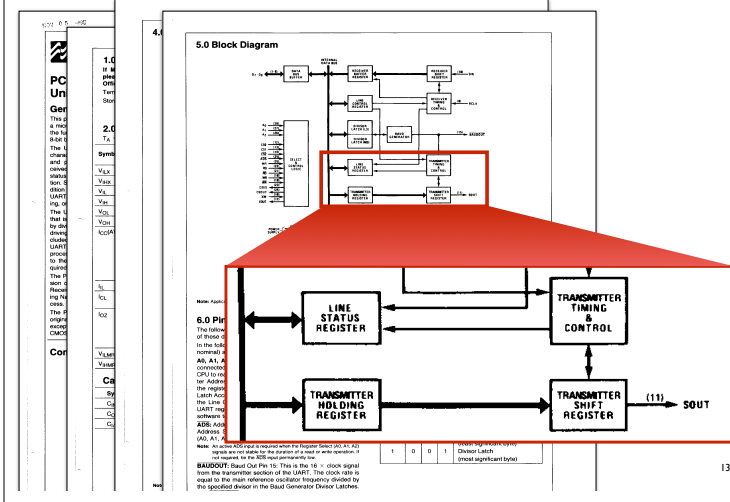
PC platform

why +5?

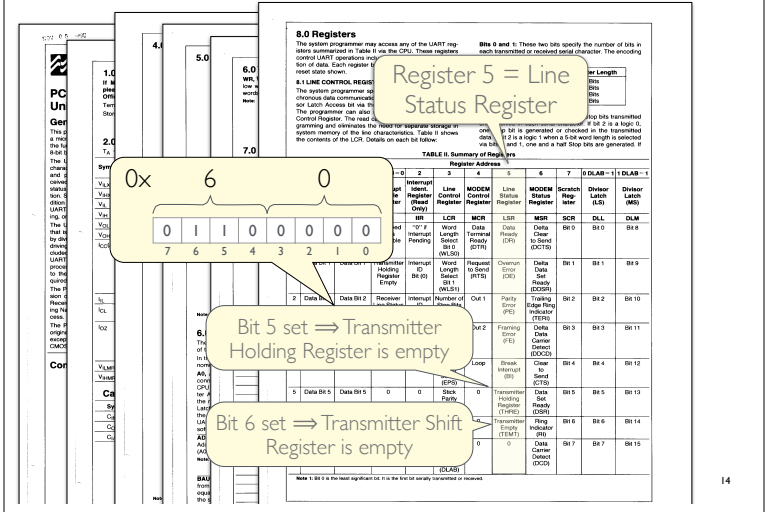
why 0x60?

12

## To the datasheet!



## To the datasheet!



## Serial port output in assembly

```
.set    PORTCOM1, 0x3f8

serial_putc:
    pushl    %eax
    pushl    %edx

1:     movw    $(PORTCOM1+5), %dx
    inb      %dx, %al    # Wait for port to be ready
    andb     $0x60, %al
    jz       1b          # check for available transmitter register
    movw     $PORTCOM1, %dx
    movb     12(%esp), %al
    outb     %al, %dx

    cmpb     $0xa, %al    # Was it a newline?
    jnz     2f

    movw     $(PORTCOM1+5), %dx
    inb      %dx, %al    # Wait again for port to be ready
    andb     $0x60, %al
    jz       1b
    movw     $PORTCOM1, %dx
    movb     $0xd, %al    # Send a carriage return
    outb     %al, %dx

2:     popl    %edx
    popl    %eax
    ret
```

15


## Reading datasheets

- Datasheets present detailed technical information in a very terse format
- Unless you are already familiar with the details, and just looking for a reference, it can be hard to find the information you need
- But persevere, and practice; this can be a useful skill
- One thing you'll often see is that computer systems typically only use a fraction of the available functionality(/transistors)
- Sample code, from the manufacturers, or on the web, can also be very useful!

16

## Interrupts

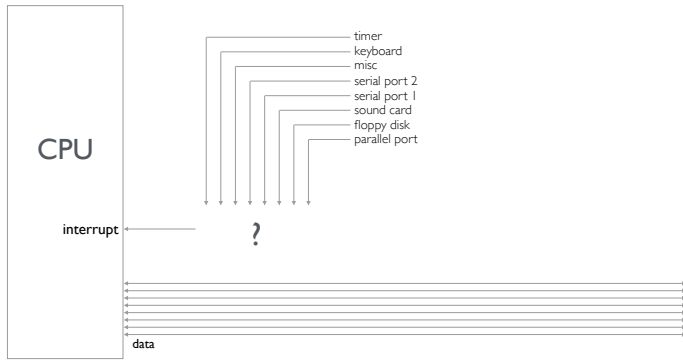
## Hardware interrupts



- The CPU has an interrupt pin
- Connect it to a timer to generate regular timer interrupts!

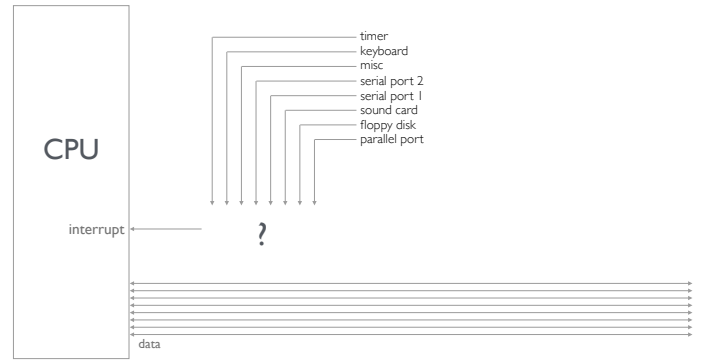
17

## How to handle multiple interrupt sources?



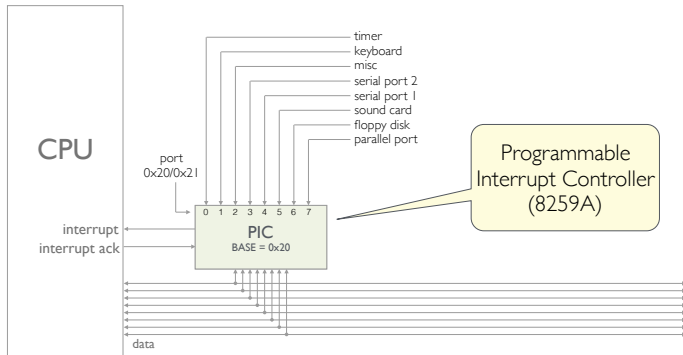
- How do we combine multiple interrupt signals?
- How do we identify and prioritize interrupt sources?

## How to handle multiple interrupt sources?



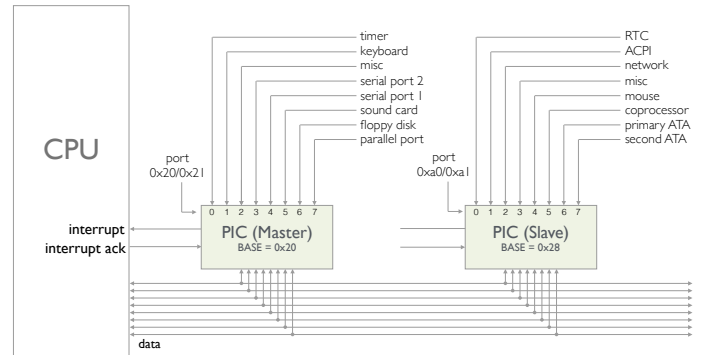
- One option: use an “or” to combine the interrupt signals
- Use the CPU to “poll” to determine which interrupt fired ...

## Adding an interrupt controller



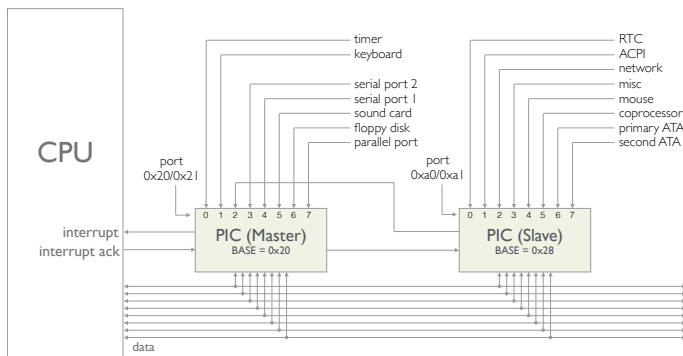
- The PIC allows individual interrupts to be masked/unmasked
- Responds to ack with programmed BASE + IRQ (interrupt request number) on data bus

## Adding multiple interrupt controllers




- Two PICs in master/slave configuration

## Adding multiple interrupt controllers



- Two PICs in master/slave configuration
- Any interrupt on the Slave triggers interrupt 2 on the Master

## IDT structure



## Initializing the PIC

```
.equ    IRQ_BASE,    0x20        # lowest hw irq number

.equ    PIC_MASTER, 0x20
.equ    PIC_SLAVE,  0xa0

# Send ICWs (initialization control words) to initialize PIC.
.macro  initpic port, base, info, init
    movb    $0x11, %al
    outb    %al, %port          # ICW1: Initialize + will be sending ICW4

    movb    %base, %al          # ICW2: Interrupt vector offset
    outb    %al, %port+1

    movb    %init, %al          # OCW1: set initial mask
    outb    %al, %port+1
.endm

initPIC: initpic PIC_MASTER, IRQ_BASE, 0x04, 0xfb # all but IRQ2 masked out
        initpic PIC_SLAVE,  IRQ_BASE+8, 0x02, 0xff
        ret
```

Interrupts on Master map  
to IDT entries 0x20-0x27

Interrupts on Slave map  
to IDT entries 0x28-0x2f

25

## Initializing the PIC

```
.equ    IRQ_BASE,    0x20        # lowest hw irq number



.equ    PIC_MASTER, 0x20
.equ    PIC_SLAVE,  0xa0

# Send ICWs (initialization control words) to initialize PIC.
.macro  initpic port, base, info, init
    movb    $0x11, %al
    outb    %al, %port          # ICW1: Initialize + will be sending ICW4

    movb    %base, %al          # ICW2: Interrupt vector offset
    outb    %al, %port+1

    movb    %init, %al          # OCW1: set initial mask
    outb    %al, %port+1
.endm

initPIC: initpic PIC_MASTER, IRQ_BASE, 0x04, 0xfb # all but IRQ2 masked out
        initpic PIC_SLAVE,  IRQ_BASE+8, 0x02, 0xff
        ret
```



26

## Enabling and disabling individual IRQs

- Individual IRQs are enabled by clearing the mask bit in the corresponding PIC:

```
static inline void enableIRQ(byte irq) {
    if (irq < 8) {
        outb(0xa1, ~(1 << (irq & 7)) & inb(0xa1));
    } else {
        outb(0x21, ~(1 << (irq & 7)) & inb(0x21));
    }
}
```

- IRQs are disabled by setting the mask bit in the corresponding PIC:

```
static inline void disableIRQ(byte irq) {
    if (irq < 8) {
        outb(0xa1, (1 << (irq & 7)) | inb(0xa1));
    } else {
        outb(0x21, (1 << (irq & 7)) | inb(0x21));
    }
}
```

27

## IRQ handling lifecycle

- Install handler for IRQ in IDT
- Use the PIC to enable that specific IRQ (the CPU will still ignore the interrupt if the IF flag is clear)
- If the interrupt is triggered, disable the IRQ and send an EOI (end of interrupt) to reenale the PIC for other IRQs:

```
static inline void maskAckIRQ(byte irq) {
    if (irq < 8) {
        outb(0xa1, (1 << (irq & 7)) | inb(0xa1));
        outb(0xa0, 0x60 | (irq & 7)); // EOI to slave
        outb(0x20, 0x62); // EOI for IRQ2 on master
    } else {
        outb(0x21, (1 << (irq & 7)) | inb(0x21));
        outb(0x20, 0x60 | (irq & 7)); // EOI to master
    }
}
```

- When the interrupt has been handled, reenale the IRQ

28

## Timers

29

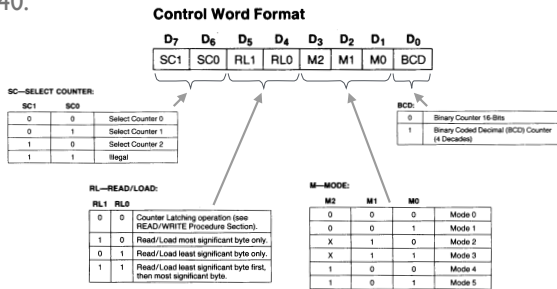
## The programmable interval timer (PIT)

- The IBM PC included an Intel 8253/54 programmable interval timer (PIT) chip
- The PIT was clocked at 1,193,181.8181Hz, for compatibility with the NTSC TV standard
- The PIT provides three counter/timers. On the PC, these were used to handle:
  - Counter 0: Timer interrupts
  - Counter 1: DRAM refresh
  - Counter 2: Playing tones via the PC's speaker

30

## ... continued

- The PIT is programmed by sending a control word to port 0x43 followed by a two byte counter value (lsb first) to port 0x40.



- Each timer/counter runs in one of six modes.


31

## Example: Programming the PIT

To configure for timer interrupts:

```
#define HZ 100 // Frequency of timer interrupts
#define PIT_INTERVAL ((1193182 + (HZ/2)) / HZ)
#define TIMERIRQ 0

static inline void startTimer() {
    outb(0x43, 0x34); // PIT control (0x43), counter 0, 2 bytes, mode 2, binary
    outb(0x40, PIT_INTERVAL & 0xff); // counter 0, lsb
    outb(0x40, (PIT_INTERVAL >> 8) & 0xff); // counter 0, msb
    enableIRQ(TIMERIRQ);
}
```



32

## Time stamp counter

- Modern Intel CPUs include a 64 bit time stamp counter that tracks the number of cycles since reset
- The current TSC value can be read in `edx:eax` using the `rdtsc` instruction
- `rdtsc` is privileged, but the CPU can be configured to allow access to `rdtsc` in user level code
- Can use differences in TSC value before and after an event to measure elapsed time
- But beware of complications related to multiprocessor systems; power management (e.g., variable clock speed); ...
- ... and virtualization .... (e.g., QEMU, VirtualBox, ...)

33

## Volatile Memory

34

## The first user program

```
unsigned flag = 0;
for (i=0; i<600; i++) {
    ...
}
printf("My flag is at 0x%x\n", &flag);
while (flag==0) {
    /* do nothing */
}
printf("Somebody set my flag to %d\n", flag);
...
```

user

- According to the semantics of C, there is no way for the value of the variable `flag` to change during the while loop ...
- ... so there is no way that the "Somebody set my flag ..." message could appear
- ... the compiler could delete the code after the while loop ...

35

## The second user program

```
unsigned flag = 0;
for (i=0; i<600; i++) {
    ...
}
printf("My flag is at 0x%x\n", &flag);
while (flag==0) {
    /* do nothing */
}
printf("Somebody set my flag to %d\n", flag);
...
```

user

```
for (i=0; i<1200; i++) {
    ...
}
unsigned* flagAddr = (unsigned*)0x4025b0;
printf("flagAddr = 0x%x\n", flagAddr);
*flagAddr = 1234;
printf("\n\nUser2 code does not return\n");
for (;;) { /* Don't return! */
}
```

user2

36

## Marking the flag as volatile

```
volatile unsigned flag = 0;
for (i=0; i<600; i++) {
    ...
}
printf("My flag is at 0x%x\n", &flag);
while (flag==0) {
    /* do nothing */
}
printf("Somebody set my flag to %d!\n", flag);
...
```

user

"My flag is at 0x4025b0"

"Somebody set my flag to 1234!"

```
for (i=0; i<1200; i++) {
    ...
}
unsigned* flagAddr = (unsigned*)0x4025b0;
printf("flagAddr = 0x%x\n", flagAddr);
*flagAddr = 1234;
printf("\n\nUser2 code does not return\n");
for (;;) { /* Don't return! */
}
```

user2

37

## The volatile modifier

- Under normal circumstances, a C compiler can treat an expression like  $x+x$  as being equivalent to  $2*x$ :
  - There is no way for the value in  $x$  to change from one side of the  $+$  to the other (no intervening assignments)
  - The compiler can replace two attempts to read  $x$  with a single read, without changing the behavior of the code
- Marking a variable as `volatile` indicates that the compiler should allow for the possibility that the stored value might change from one read to the next
- The `volatile` modifier is often necessary when working with memory mapped I/O

38

## Unresolved issues

39

## Issues with the Week 3 lab example

- Although we are running in protected mode, we are using segments that span the full address space, so there is **no true protection** between the different programs
- Address space layout is ad hoc: different programs load and run at different addresses; there is no consistency
- We had to choose different (but essentially arbitrary) start addresses for user and user2, even when they were just two copies of the same program
- Why should worries about low level memory layout & size propagate in to the design of higher-level applications?
- Our user programs included duplicate code (e.g., each one has its own implementation of `printf`). How can we support sharing of common code or data between multiple programs?


40

## Paging

41

## Paging


- "All problems in computer science can be solved by another level of indirection" (David Wheeler)
- Partition the address space in to a collection of "pages"
- Translate between addresses in some idealized "virtual address space" and "physical addresses" to memory.



42

## Example


- Suppose that we partition our memory into 8 pages:



43

## Practical reality

- IA32 partitions the 32-bit, 4GB address space in to 4KB pages



- It also allows the address space to be viewed as 4MB “super pages”



- We need a table with  $2^{10}$  entries to translate virtual super page numbers in to physical page numbers
- With 4 bytes/entry, this table, called a **page directory**, takes  $2^{12}$  bytes - one 4K page!

44

## Paging with 4MB super pages

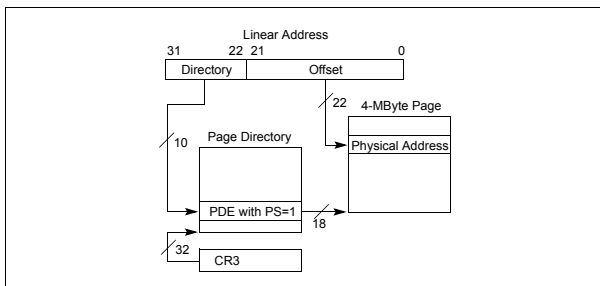


Figure 4-3. Linear-Address Translation to a 4-MByte Page using 32-Bit Paging

- The cr3 register points to the “current” page directory
- Individual page directory entries (PDEs) specify a 10 bit physical super page address plus some additional control bits

45

## Page tables


- A table describing translations for all 4KB pages would require  $2^{20}$  entries
- With four bytes per entry, a full page table would take 4MB
- Most programs are small, at least in comparison to the full address space
  - $\Rightarrow$  most address spaces are fairly sparse
- is there a more compact way to represent their page tables?

46

## Example

- Suppose that our memory is partitioned in to 64 pages
- But we are only use a small number of those pages...
- ... in fact, only a small number of the rows
- Then we can represent the full table more compactly as a tree:

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39
40	41	42	43	44	45	46	47
48	49	50	51	52	53	54	55
56	57	58	59	60	61	62	63



47

## Paging with 4KB pages

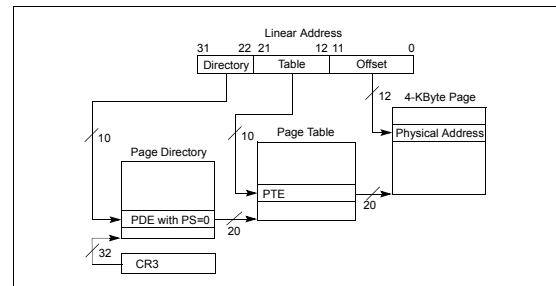


Figure 4-2. Linear-Address Translation to a 4-KByte Page using 32-Bit Paging

- A typical address space can now be described by a page directory plus one or two page tables (i.e., 4-12KB)
- Can mix pages and super pages for more flexibility

48



## CR3, PDEs, PTEs

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
Address of page directory <sup>1</sup>												Ignored				P C D	P W T	Ignored		CR3												
Bits 31:22 of address of 4MB page frame						Reserved (must be 0)			Bits 39:32 of address <sup>a</sup>		P A T	Ignored		G	1	D	A	P C D	P W T	U / S	R / W	1	PDE: 4MB page									
Address of page table												Ignored				0	1	A	P C D	P W T	U / S	R / W	1	PDE: page table								
Ignored																										0	PDE: not present					
Address of 4KB page frame												Ignored				G	P A T	D	A	P C D	P W T	U / S	R / W	1	PTE: 4KB page							
Ignored																										0	PTE: not present					

Figure 4-4. Formats of CR3 and Paging-Structure Entries with 32-Bit Paging

49

## Details

- Paging structures use physical addresses
- P(resent) bit 0 is used to mark valid entries (an OS can use the remaining “ignored” fields to store extra information)
- Hardware updates D(irty) and A(ccessed) bits to track usage
- R/W bits allow regions of memory to be marked “read only”
- S/U bits allow regions of memory to be restricted to “supervisor” access only (rather than general “user”)
- G(lobal) bit allows pages to be marked as appearing in every address space
- PCD and PWD bits control caching behavior

50

## The translation lookaside buffer (TLB)

- Recall that the IA32 tracks current segment base and limit values in hidden registers to allow for faster access
- A more sophisticated form of cache, called the translation lookaside buffer (TLB), is used to keep track of active mappings within the CPU's memory management unit
- Programmers typically ignore the TLB: “it just works”
- But not so in programs that modify page directories and page tables: extra steps are required to ensure that the TLB is updated to reflect changes in the page table
  - Loading a value in to CR3 will flush the TLB
  - the “invlpg addr” instruction removes TLB entries for a specific address

51

## Segmentation and paging

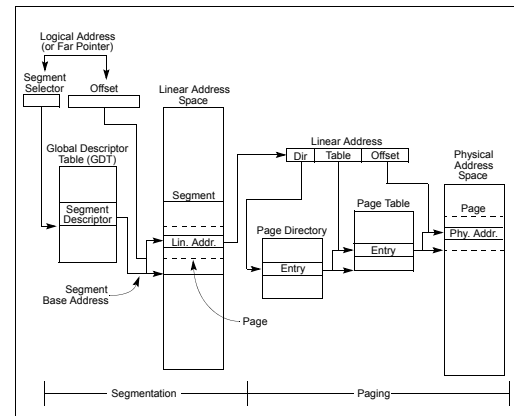


Figure 3-1. Segmentation and Paging

52

## Protection and address space layout

- A typical operating system adopts a virtual memory layout something like the following for all address spaces:




- The operating system is in every address space; it's pages are protected from user programs by limiting those parts of the page directory to “supervisor” access
- The OS portion of the page directory can take advantage of G(lobal) bits so that TLB entries for kernel space are retained when we switch between address spaces

53

## Protection and address space layout

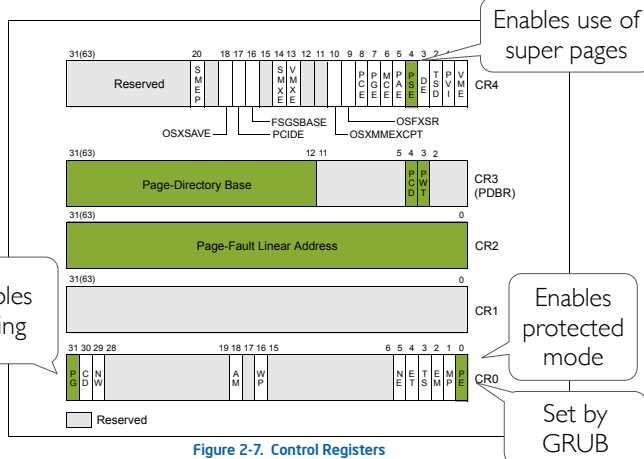
- A typical operating system adopts a virtual memory layout something like the following for all address spaces:



- User code and data mappings differ from one address space to the next
  - there is no way for one user program to access memory regions for another program ...
  - ... unless the OS provides the necessary mappings
  - user programs do not have a **capability** to access unauthorized regions of memory

54

## Control registers to enable paging



## Initialization

- How do we get from physical memory, after booting:




- to virtual address spaces with paging enabled?



- Two key steps
  - Create an initial page directory
  - Enable the CPU paging mechanisms

## Creating a 1:1 mapping

- While running at lower addresses, create an initial page directory that maps the lower 1GB of memory in two different regions of the virtual address space



- Turn on paging ...
- jump to an address in the upper 1GB of virtual memory ...
- and then proceed without the lower mapping ...

## Working with physical & virtual addresses

- It is convenient to work with page directories and page tables as regular data structures (virtual addresses):

```
struct Pdir { unsigned pde[1024]; };
struct Ptab { unsigned pte[1024]; };


/*
 * Return a pointer to the page table for the ith entry of the specified
 * pdir, or NULL if it is not present (0x1) or is a super page (0x80).
 */
static inline struct Ptab* getPagetab(struct Pdir* pdir, unsigned i) {
    return ((pdir->pde[i]&0x81)==0x1)
        ? fromPhys(struct Ptab*, align(pdir->pde[i], PAGE_SIZE)) : 0;
}
```

- But sometimes we have to work with physical addresses:

```
/*
 * Set the page directory control register to a specific value.
 */
static inline void setPdir(unsigned pdir) {
    asm(" movl %0, %%cr3\n" : "r"(pdir));
}
```

## From physical to virtual, and back again

- Because we map the top 1GB of virtual memory to the bottom 1GB of physical memory, it is easy to convert between virtual and physical addresses:



- (But how can we do this in a type safe language ... ?)

## Details (Part I)

- Constants to describe the virtual address space

```
KERNEL_SPACE = 0xc0000000 # Kernel space starts at 3GB
KERNEL_LOAD = 0x00100000 # Kernel loads at 1MB
```

- The kernel is configured to load at a low physical address but run at a high virtual address:

```
OUTPUT_FORMAT(elf32-i386)
ENTRY(physentry)

SECTIONS {
    physentry = entry - KERNEL_SPACE;
    . = KERNEL_LOAD + KERNEL_SPACE;

    .text ALIGN(0x1000) : AT(ADDR(.text) - KERNEL_SPACE) {
        _text_start = .; *(.text) *(.handlers) _text_end = .;
        *(.rodata*)
        *(.data)
        _start_bss = .; *(COMMON) *(.bss) _end_bss = .;
    }
}
```

## Details (Part 2)

- Reserve space for an initial page directory structure:

```
.data
.align (1<<PAGESIZE)
pdir: .space 4096          # Initial page directory
```

- Zero all entries in the table:

```
leal (pdir-KERNEL_SPACE), %edi
movl %edi, %esi          # save in %esi

movl $1024, %ecx         # Zero out complete page directory
movl $0, %eax
1: movl %eax, (%edi)
addl $4, %edi
decl %ecx
jnz 1b
```

61

## Details (Part 3)

- Install the lower and upper mappings in the initial page directory structure:

```
movl %esi, %edi          # Set up 1:1 and kernelspace mappings
movl $(PHYSMAP>>SUPERSIZE), %ecx
movl $(PERMS_KERNELSPACE), %eax

1: movl %eax, (%edi)
movl %eax, (4*(KERNEL_SPACE>>SUPERSIZE))(%edi)
addl $4, %edi            # move to next page dir slots
addl $(4<<20), %eax      # entry for next superpage to be mapped
decl %ecx
jnz 1b
```

- Load the CR3 register:

```
movl %esi, %cr3          # Set page directory

movl %cr4, %eax           # Enable super pages (CR4 bit 4)
orl $(1<<4), %eax
movl %eax, %cr4
```

62

## Details (Part 4)

- Turn on paging:

```
movl %cr0, %eax          # Turn on paging (1<<31)
orl $(1<<31), %eax        # and protection (1<<0)
movl %eax, %cr0

movl $high, %eax          # Make jump into kernel space
jmp *%eax

high: leal kernelstack, %esp # Now running at high addresses
      # Set up initial kernel stack
```

- And now that's out of the way, the kernel can get down to work ...

63

## Page faults

- If program tries to access an address that is either not mapped, or that it is not permitted to use, then a page fault exception (14) occurs
- The address triggering the exception is loaded in to CR2
- Details of the fault are in the error code in the context:


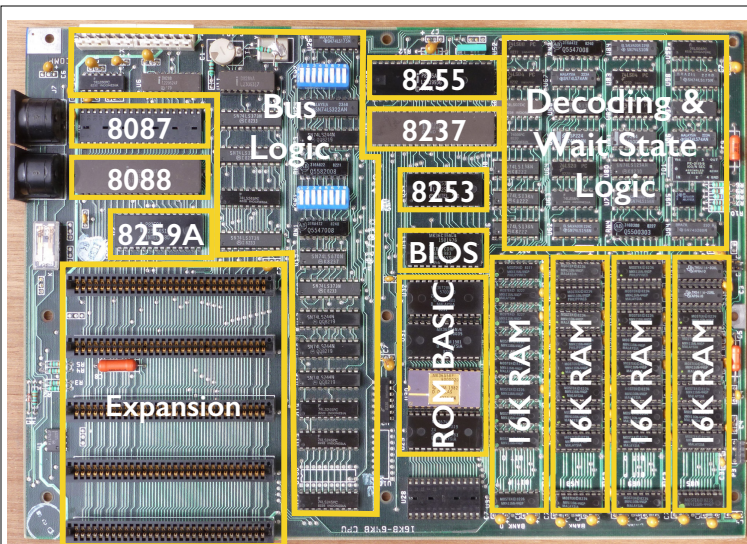


Figure 4-12. Page-Fault Error Code

64



[http://www.minuszerodegrees.net/5150/early/5150\\_early.htm](http://www.minuszerodegrees.net/5150/early/5150_early.htm)

65

Ok, kernel, over to you ...

66