



## CS 410/510


### Languages & Low-Level Programming

Mark P Jones  
Portland State University

Spring 2016

Week 10: Habit and HaL4

1




2

## The CEMLaBS Project

- “Using a Capability-Enhanced Microkernel as a Testbed for Language-Based Security”
- Started October 2014, Funded by The National Science Foundation
- Three main questions:
  - **Feasibility:** Is it possible to build an inherently “unsafe” system like seL4 in a “safe” language like Habit?
  - **Benefit:** What benefits might this have, for example, in reducing verification costs?
  - **Performance:** Is it possible to meet reasonable performance goals for this kind of system?

3

## Chipping away ...



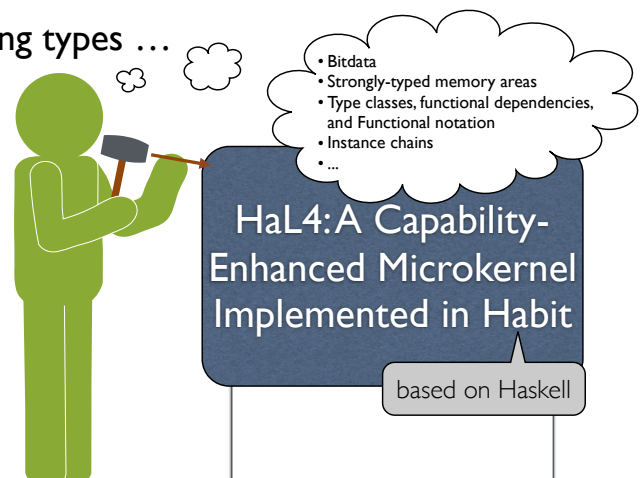
4

## Chipping away ...



5

## Using types ...



6

## Example: IA32 Paging Structures

Remember this?

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Address of page directory <sup>1</sup>																		Ignored				P C D	P W T	Ignored		CR3					
Bits 31:22 of address of 4MB page frame								Reserved (must be 0)				Bits 39:32 of address <sup>2</sup>				P A T	Ignored	G	1	D	A	P C D	U / S	R / W	1	PDE: 4MB page					
Address of page table																		Ignored				0	I g n	A	P C D	P W T	U / S	R / W	1	PDE: page table	
Ignored																										Q	PDE: not present				
Address of 4KB page frame																		Ignored				G	P A T	D	A	P C D	U / S	R / W	1	PTE: 4KB page	
Ignored																										Q	PTE: not present				

Figure 4-4. Formats of CR3 and Paging-Structure Entries with 32-Bit Paging

7

## Example: IA32 Paging Structures

Here is how we describe page directory entries in Habit:

```
bitdata PDE /WordSize -- Page Directory Entries
= UnmappedPDE [ unused=0 :: Bit 31 | B0 ] -- Unused entry (present bit reset)

| PageTablePDE [ ptab :: Phys PageTable -- physical address of page table
| unused=0 :: Bit 4
| B0 -- signals PageTablePDE
| attrs=readWrite :: PagingAttrs -- paging attributes
| B1 ] -- present bit set

| SuperPagePDE [ super :: Phys SuperPage -- physical address of superpage
| unused=0 :: Bit 13
| global=0 :: Bit 1 -- 1 => global translation (if cr4.pge=1)
| B1 -- signals SuperPagePDE
| attrs :: PagingAttrs -- paging attributes
| B1 ] -- present bit set

bitdata PagingAttrs /6
= PagingAttrs [ dirty = 0 :: Bit 1 -- Dirty; 1 => data written to page
| accessed = 0 :: Bit 1 -- Accessed; 1 => page accessed
| caching = Caching[] :: Caching
| us :: Bit 1 -- User/supervisor; 1 => user access allowed
| rw :: Bit 1 ] -- Read/write; 1 => write access allowed
```

8

## Example: IA32 Address Space Layout

Remember this?

0GB 1GB 2GB 3GB 4GB

user space kernel space

```
struct PageDir /4K {
  pdes :: Array UserSuperPages (Stored PDE)
  kpdes :: Array KernelSuperPages (Stored KPDE)
}

type UserSuperPages = 768 -- Number of superpages in the user space address range
type KernelSuperPages = 256 -- Number of superpages in the kernel space address range

bitdata KPDE /WordSize -- Kernel Page Directory Entries, a variant of SuperPagePDE
= KPDE [ B00 -- leading zeros
| ix :: Ix KernelSuperPages -- physical superpage number
| unused=0 :: Bit 13
| global=1 :: Bit 1 -- 1 => global translation (if cr4.pge=1)
| B1 -- indicates a form of SuperPagePDE
| attrs = kernelOnly :: PagingAttrs -- paging attributes
| B1 ]

instance Initable PageDir where
  initialize = PageDir [ pdes <- initArray (\ix -> initStored UnmappedPDE[])
| kpdes <- initArray (\ix -> initStored KPDE[ix]) ]
```

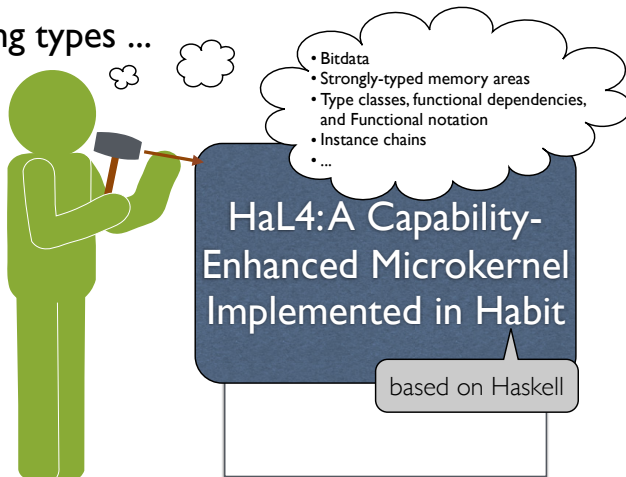
9

## Summary

- The code is not simple ...
- ... but this reflects the underlying data structures
- Correct usage is enforced by the type system:
  - The separation between user space and kernel space
  - The requirement for physical rather than virtual addresses in page directory entries
  - Invariants on data structure size
  - Initialization


10

## Using types ...



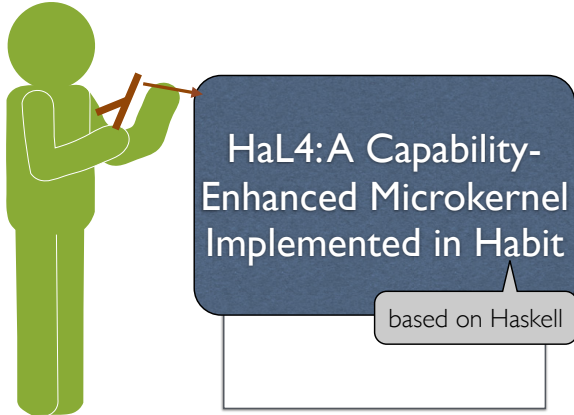
11

## Using lambda ...




12

## Using lambda ...




13

## Using lambda ...



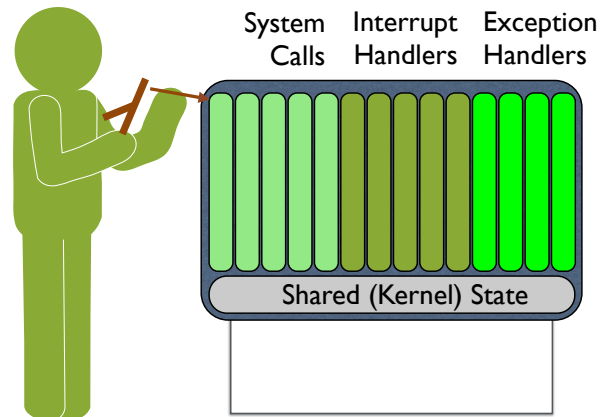
14

## Using lambda ...



15

## Using lambda ...



16

Validate Parameters

Perform Action

Return with error code

Return with success code

```
syscallMapPageDir :: (KE k, KW k) => k a
syscallMapPageDir
= do curr <- getCurrent
    asidIdx <- getReg asidCapReg curr
    case<- lookupCapAll curr.cspace asidIdx of
      Ref asidCap ->
        case<- get asidCap.objptr of
          ASIDTableObj[] ->
            range <- getCapdata asidCap
            offset <- getReg offsetReg
            case offset `inRange` range of
              Just asid ->
                let slot = asidTable @@ asid
                count <- get slot.count
                if count==0 then
                  pdirIdx <- getReg pdirCapReg curr
                  case<- lookupCapAll curr.cspace pdirIdx curr of
                    Ref pdirCap ->
                      case<- get pdirCap.objptr of
                        PageDirObj[pdir] ->
                          case<- getCapdata pdirCap of
                            UnmappedPD[] ->
                              set slot.pdir (Ref pdir)
                              set slot.count 1
                              setCapdata pdcap MappedPD[asid]
                              success curr
                              -> mappedErr curr
                            -> invalidCapabilityErr curr
                            Null -> invalidCapabilityErr curr
                            else mappedErr curr
                            Nothing -> rangeErrorErr curr
                            -> invalidCapabilityErr curr
                            Null -> invalidCapabilityErr curr
```

Parameter Validation

Action

Error Reporting

17

18

## Programming with continuations

- Traditional control flow structure

```
do f <- openFile "file.txt"
  l1 <- readLine f
  l2 <- readLine f
  out (l1, l2)
  closeFile f
```

- How to deal with errors?

- Make functions return error codes (and hope that callers will check those codes)?
- Add the ability to throw and catch exceptions?
- Use continuations ...

19

## Programming with continuations

- Instead of

```
openFile :: String -> IO FileHandle
```

- Try:

```
openFile :: String
  -> (ErrorCode -> IO a)
  -> (FileHandle -> IO a)
  -> IO a
```

higher-order, or  
first-class functions

- It's as if we've given openFile two return addresses: one to use when an error occurs, and one to use when the call is successful.

20

## Programming with continuations

- Our original program using continuations:

```
openFile "file.txt"
  (\error -> ...)
  (\f -> do l1 <- readLine f
           l2 <- readLine f
           out (l1, l2)
           closeFile f)
```

lambda  
expressions:  
\var -> expr

- Could we do the same for readLine?

21

## Programming with continuations

- Our original program using continuations:

```
openFile "file.txt"
  (\error -> ...)
  (\f -> readLine f
        (\error -> ...)
        (\l1 -> readLine f
              (\error -> ...)
              (\l2 -> do out (l1, l2)
                        closeFile f))))
```

- Hmm, not so pretty ...

22

## Programming with continuations

- Name the error handlers:

```
openFile "file.txt"
  err1
  (\f -> readLine f
        err2
        (\l1 -> readLine f
              err3
              (\l2 -> do out (l1, l2)
                        closeFile f))))
```

23

## Programming with continuations

- Reformat:

```
openFile "file.txt" err1 (\f ->
  readLine f          err2 (\l1 ->
  readLine f          err3 (\l2 ->
    do out (l1, l2)
    closeFile f)))
```

- Looking better ...

24

## Programming with continuations

- Add an infix operator: `f $ x = f x`

```
openFile "file.txt" err1 $ \f ->
readLine f          err2 $ \l1 ->
readLine f          err3 $ \l2 ->
do out (l1, l2)
  closeFile f
```

- Fewer parentheses ...
- Easier to add or remove individual lines ...
- ... still a little cluttered by error handling behavior

25

## Programming with continuations

- Integrate the error handlers in to the main operations:

```
openFile "file.txt" $ \f ->
readLine f          $ \l1 ->
readLine f          $ \l2 ->
do out (l1, l2)
  closeFile f
```

- Not always applicable ...
- ... but a good choice for HaL4 where the response to a particular type of invalid parameter is always the same (typically, returning an error code to the caller)
- ... and this also ensures a consistent API

26

## “Validators”

The implementation of HaL4 includes a small library of validator functions:

```
getCurrent      :: KR k => (TCBRef -> k a) -> k a
getRegCap       :: KE k => #r -> TCBRef
                -> (CapRef -> k a) -> k a
emptyCapability :: KE k => TCBRef -> CapRef -> k a -> k a
cdtLeaf         :: KE k => TCBRef -> CapRef -> k a -> k a
notMaxDepth     :: KE k => TCBRef -> CapRef -> k a -> k a
untypedCapability :: KE k => TCBRef -> CapRef
                -> (UntypedRef -> k a) -> k a
pageDirCapability :: KE k => TCBRef -> CapRef
                -> (PageDirRef -> PMapData -> k a) -> k a
pageTableCapability :: KE k => TCBRef -> CapRef
                -> (PageTableRef -> MapData -> k a) -> k a
```

27

## “Validators”

- In effect, we have built an embedded domain specific language, just for validating parameters in HaL4

- Benefits include:

- Ease of reuse
- Consistency
- Clarity
- Ability to pass multiple results on to continuation

28

```
syscallMapPageDir :: (KE k, KW k) => k a
syscallMapPageDir
```

<code>getCurrent</code>	<code>\$ \curr -&gt;</code>
<code>getMapPageDirASIDTab curr</code>	<code>\$ \asidcap -&gt;</code>
<code>asidTableCapability curr asidcap</code>	<code>\$ \range -&gt;</code>
<code>getMapPageDirOffset curr</code>	<code>\$ \offset -&gt;</code>
<code>asidInRange curr offset range</code>	<code>\$ \asid -&gt;</code>
<code>asidNotUsed curr asid</code>	<code>\$ \slot -&gt;</code>
<code>getMapPageDirPDir curr</code>	<code>\$ \pdcap -&gt;</code>
<code>pageDirCapability curr pdcap</code>	<code>\$ \pdir pdmd -&gt;</code>
<code>unmappedPD curr pdmd</code>	<code>\$</code>

**Validators**

```
do set slot.pdir (Ref pdir)
  set slot.count 1
  setCapdata pdcap MappedPD[asid]
  success curr
```

**Action**

29

```
syscallMapPageDir :: (KE k, KW k) => k a
```

```
syscallMapPageDir
= getCurrent          $ \curr ->
  getMapPageDirASIDTab curr $ \asidcap ->
    asidTableCapability curr asidcap $ \range ->

  getMapPageDirOffset curr $ \offset ->
    asidInRange curr offset range $ \asid ->
      asidNotUsed curr asid $ \slot ->

  getMapPageDirPDir curr $ \pdcap ->
    pageDirCapability curr pdcap $ \pdir pdmd ->
      unmappedPD curr pdmd $
```

```
do set slot.pdir (Ref pdir)
  set slot.count 1
  setCapdata pdcap MappedPD[asid]
  success curr
```

**clear and concise**

30

```
syscallMapPageDir :: (KE k, KW k) => k a
syscallMapPageDir
```

```
= getCurrent          $ \curr ->
  getMapPageDirASIDTab curr $ \asidcap ->
  asidTableCapability curr asidcap $ \range ->

  getMapPageDirOffset curr $ \offset ->
  asidInRange curr offset range $ \asid ->
  asidNotUsed curr asid $ \slot ->

  getMapPageDirPDir curr $ \pdcap ->
  pageDirCapability curr pdcap $ \pdir pdmd ->
  unmappedPD curr pdmd $
```

```
do set slot.pdir (Ref pdir)
  set slot.count 1
  setCapdata pdcap MappedPD[asid]
  success curr
```

reusable

31

```
syscallMapPageDir :: (KE k, KW k) => k a
syscallMapPageDir
```

```
= getCurrent          $ \curr ->
  getMapPageDirASIDTab curr $ \asidcap ->
  asidTableCapability curr asidcap $ \range ->

  getMapPageDirOffset curr $ \offset ->
  asidInRange curr offset range $ \asid ->
  asidNotUsed curr asid $ \slot ->

  getMapPageDirPDir curr $ \pdcap ->
  pageDirCapability curr pdcap $ \pdir pdmd ->
  unmappedPD curr pdmd $
```

```
do set slot.pdir (Ref pdir)
  set slot.count 1
  setCapdata pdcap MappedPD[asid]
  success curr
```

high-level

32

```
syscallMapPageDir :: (KE k, KW k) => k a
syscallMapPageDir
```

```
= getCurrent          $ \curr ->
  getMapPageDirASIDTab curr $ \asidcap ->
  asidTableCapability curr asidcap $ \range ->

  getMapPageDirOffset curr $ \offset ->
  asidInRange curr offset range $ \asid ->
  asidNotUsed curr asid $ \slot ->

  getMapPageDirPDir curr $ \pdcap ->
  pageDirCapability curr pdcap $ \pdir pdmd ->
  unmappedPD curr pdmd $
```

```
do set slot.pdir (Ref pdir)
  set slot.count 1
  setCapdata pdcap MappedPD[asid]
  success curr
```

performance concerns?

33

## Other examples: reading/writing registers

```
syscallRetype :: (KE k, RW k) => k a
syscallRetype
```

```
= getCurrent          $ \curr ->
  getRootCap curr $ \root ->
  untypedCapability curr root $ \ut ->
  cdtLeaf curr root $
  notMaxDepth curr root $
  getCNodeCap curr $ \cncap ->
  cnodeCapability curr cncap $ \cnode ->
  getCapArray curr cnode $ \caps ->
  allEmptyCapabilities curr caps $
  validType curr ut caps $ \objs ->
  addChildren curr root caps objs
```

34

## Other examples: reading/writing registers

```
syscallReadRegs :: (KE k, RW k) => k a
syscallReadRegs
= getCurrent          $ \curr -> -- get current thread
  getDstCap curr      $ \cap -> -- get specified capability
  tcbCapability curr cap $ \tcb ps -> -- ... to a tcb object
  permitsRead curr ps.perms $ -- ... that we can read
  transferRegs tcb (>->) curr
```

```
syscallWriteRegs :: (KE k, RW k) => k a
syscallWriteRegs
= getCurrent          $ \curr -> -- get current thread
  getDstCap curr      $ \cap -> -- get specified capability
  tcbCapability curr cap $ \tcb ps -> -- ... to a tcb object
  permitsWrite curr ps.perms $ -- ... that we can write
  transferRegs tcb (<-<) curr
```

$r_1 \leftarrow r_2 = r_2 \rightarrow r_1$

$r_1 \rightarrow r_2$   
= do x <- readRef r<sub>1</sub>  
writeRef r<sub>2</sub> x

$r_1 \rightarrow r_2$   
= readRef r<sub>1</sub> >=> writeRef r<sub>2</sub>

35

## Other examples: transferring registers

```
transferRegs :: KE k => TCBRef
-> (RegRef -> RegRef -> k ())
-> TCBRef -> k a
```

```
transferRegs tcb xfer curr
= do let tcbIframe = tcb.context.iframe
  currRegs = curr.context.regs
  tcbIframe.eip `xfer` currRegs.esi
  tcbIframe.esp `xfer` currRegs.ecx
  case<- getIPCBuffer curr of
    Ref buf -> let tcbRegs = tcb.context.regs
      tcbRegs.eax `xfer` (buf.mrs @@ 3)
      tcbRegs.ebx `xfer` (buf.mrs @@ 4)
      tcbRegs.ecx `xfer` (buf.mrs @@ 5)
      tcbRegs.edx `xfer` (buf.mrs @@ 6)
      tcbRegs.esi `xfer` (buf.mrs @@ 7)
      tcbRegs.edi `xfer` (buf.mrs @@ 8)
      tcbRegs.ebp `xfer` (buf.mrs @@ 9)
  success curr
```

<-< or >->  
would work here!

36

## Other examples: reading/writing registers

```
syscallReadRegs :: (KE k, RW k) => k a
syscallReadRegs
= getThread          $ \curr -> -- get current thread
  getDstCap curr     $ \cap -> -- get specified capability
  tcbCapability curr cap $ \tcb ps -> -- ... to a tcb object
  permitsRead curr ps.perms $ -- ... that we can read
  transferRegs tcb (>->) curr

syscallWriteRegs :: (KE k, RW k) => k a
syscallWriteRegs
= getThread          $ \curr -> -- get current thread
  getDstCap curr     $ \cap -> -- get specified capability
  tcbCapability curr cap $ \tcb ps -> -- ... to a tcb object
  permitsWrite curr ps.perms $ -- ... that we can write
  transferRegs tcb (<-<) curr
```

37

## Other examples: reading/writing registers

```
syscallReadRegs :: (KE k, RW k) => k a
syscallReadRegs
= regAccess          $ \curr tcb ps -> -- find thread
  permitsRead curr ps $ -- check for read perm
  transferRegs tcb (>->) curr

syscallWriteRegs :: (KE k, RW k) => k a
syscallWriteRegs
= regAccess          $ \curr tcb ps -> -- find thread
  permitsWrite curr ps $ -- check for write perm
  transferRegs tcb (<-<) curr

regAccess :: (KE k) => (TCBRef -> TCBRef -> Perms -> k a) -> k a
regAccess k
= getThread          $ \curr -> -- get current thread
  getDstCap curr     $ \cap -> -- get specified capability
  tcbCapability curr cap $ \tcb ps -> -- ... to a tcb object
  k curr tcb ps.perms
```

38

## Representing function values

- How should we represent values of type `Int -> Int`?
  - There are many different values, including: `(\z -> x+z)`, `(\x -> x+1)`, `(\x -> x*2)`, `(\x -> f(g(x)))`, ...
  - ... any of which could be passed as arguments to other functions ...
  - ... so we need a uniform, but flexible way to represent them
- A common answer is to represent functions like these by a pointer to a “closure”, a heap allocated object that contains:
  - a code pointer (i.e., the code for the function)
  - the values of its free variables

39

## Closures

- Every function of type `Int -> Int` will be represented using the same basic structure:

codeptr	...
---------	-----

- The code pointer and list of variables vary from one function value to the next:

<code>(\z -&gt; x+z)</code>	codeptr <sub>1</sub>	x	
<code>(\x -&gt; x+1)</code>	codeptr <sub>2</sub>		
<code>(\x -&gt; x*2)</code>	codeptr <sub>3</sub>		
<code>(\x -&gt; f(g(x)))</code>	codeptr <sub>4</sub>	f	g

- To make a closure, allocate a suitably sized block of memory and save the required code pointer and variable values

40

## Calling unknown functions

- If `f` is a known function, then we call it by pushing its arguments on the stack and jumping directly to its code
- What if `f` is an unknown function, represented by a variable that points to a closure structure instead?
- The System V ABI doesn't cover this case, but we can make up our own convention:
  - push the arguments
  - push a pointer to the closure (for access to free variables)
  - call the code that is pointed to at the start of the closure
- This process is known as “entering a closure”

41

## Thunks

- A closure that can be entered without any arguments is sometimes referred to as a “thunk”
- Thunks represent suspended/delayed computations that can be executed (or “entered” or “invoked”) repeatedly at a later stage.
- Examples:
  - procedure values / monadic computations
 

```
r1 >-> r2 = do x <- readRef r1
              writeRef r2 x
```
  - interrupt/exception/system call handlers
  - ...

42

## Pros/cons of using higher order functions

- Pros:
  - high-level
  - clear and concise
  - facilitates reuse (less code, greater consistency)
- Cons:
  - what impact do all these closures have on performance?
    - it's hard to explore this question without a way to talk about these things more precisely ...
    - in other words, we need a *language* that we can use to express these ideas...

43

## A notation for closures and thunks

- Closures:  $k\{x_1, \dots, x_n\}$ 
  - code pointer:  $k$
  - stored fields:  $x_1, \dots, x_n$
  - If  $f$  is a closure, then we write  $f @ x$  for the result of entering  $f$  with argument  $x$
- Thunks:  $m[x_1, \dots, x_n]$ 
  - code pointer:  $m$
  - stored fields:  $x_1, \dots, x_n$
  - if  $t$  is a thunk, then we write `invoke t` for the result that is obtained by invoking  $t$

44

## Do notation and three address code

- Some uses of do notation look a lot like three address code:

```
do x <- f y z      x := y + z
  t <- g x y        t := x * y
  u <- h t          u := -t
  p x t            goto p
```

- Some optimization techniques are like algebraic properties:

```
(do x <- return y; c) = [y/x] c
```

- Is this a “monad law” or is it “copy propagation”?

45

## Comprehending Monads

Types	
$K^*$	$= K$
$(U \rightarrow V)^*$	$= (U^* \rightarrow M V^*)$
$(U, V)^*$	$= (U^*, V^*)$
Terms	
$x^*$	$= [x]^M$
$(\lambda x \rightarrow v)^*$	$= [(\lambda x \rightarrow v^*)]^M$
$(tu)^*$	$= [y \mid f \leftarrow t^*, x \leftarrow u^*, y \leftarrow (f x)]^M$
$(u, v)^*$	$= [(x, y) \mid x \leftarrow u^*, y \leftarrow v^*]^M$
$(fst t)^*$	$= [(fst z) \mid z \leftarrow t^*]^M$
Assumptions	
$(x_1 :: T_1, \dots, x_n :: T_n)^*$	$= x_1 :: T_1^*, \dots, x_n :: T_n^*$
Typings	
$(A \vdash t :: T)^*$	$= A^* \vdash t^* :: M T^*$

Figure 7: Call-by-value translation.

Moggi,  
Wadler,  
Kennedy,  
Benton,  
...  
Bailey

46

## From source terms ...

```
id      = \x -> x

compose = \f g x -> f (g x)

map     = \f xs ->
  case xs of
    Nil      -> Nil
    Cons y ys -> Cons (f y) (map f ys)
```

47

## ... to MIL programs

```
id      ← k0{ }
k0{ } x = b0(x)
b0(x)   = return x

map     ← k4{ }
k4{ } f = k5{f}
k5{f} xs = b2(f, xs)
b2(f, xs) = case xs of
  Nil() → b3()
  Cons(y, ys) → b4(f, y, ys)

b3()    = Nil()
b4(f, y, ys) = z ← f @ y
              m ← map @ f
              zs ← m @ ys
              Cons(z, zs)

compose ← k1{ }
k1{ } f = k2{f}
k2{f} g = k3{f, g}
k3{f, g} x = b1(f, g, x)
b1(f, g, x) = u ← g @ x
              f @ u
```

48



## MIL, a monadic intermediate language

```

a ::= v          -- variable
   | n          -- integer literal

t ::= return a   -- simple value
   | p((a1, ..., an)) -- primitive call
   | b(a1, ..., an) -- block call
   | C(a1, ..., an) -- data constructor
   | k{a1, ..., an} -- closure constructor
   | f @ a       -- enter closure
   | m[a1, ..., an] -- thunk constructor
   | invoke a    -- invoke closure

c ::= v <- t; c   -- monadic bind
   | t           -- (generalized) tail call
   | case v of alt1; ...; altn -- conditional branch

alt ::= C(v1, ..., vn) -> b(a1, ..., an) -- constructor match
      _                -> b(a1, ..., an) -- default match

def ::= b(v1, ..., vn) = c -- block entry point
      | k{v1, ..., vn} v = c -- closure definition
      | v <- t              -- top-level constant

```

49

## MIL, a monadic intermediate language

```

a ::= v          -- variable
   | n          -- integer literal

t ::= return a   -- simple value
   | p((a1, ..., an)) -- primitive call
   | b(a1, ..., an) -- block call
   | C(a1, ..., an) -- data constructor
   | k{a1, ..., an} -- closure constructor
   | f @ a       -- enter closure
   | m[a1, ..., an] -- thunk constructor
   | invoke a    -- invoke closure

c ::= v <- t; c   -- monadic bind
   | t           -- (generalized) tail call
   | case v of alt1; ...; altn -- conditional branch

alt ::= C(v1, ..., vn) -> b(a1, ..., an) -- constructor match
      _                -> b(a1, ..., an) -- default match

def ::= b(v1, ..., vn) = c -- block entry point
      | k{v1, ..., vn} v = c -- closure definition
      | v <- t              -- top-level constant

```

Definitions

50

## MIL, a monadic intermediate language

```

a ::= v          -- variable
   | n          -- integer literal

t ::= return a   -- simple value
   | p((a1, ..., an)) -- primitive call
   | b(a1, ..., an) -- block call
   | C(a1, ..., an) -- data constructor
   | k{a1, ..., an} -- closure constructor
   | f @ a       -- enter closure
   | m[a1, ..., an] -- thunk constructor
   | invoke a    -- invoke closure

c ::= v <- t; c   -- monadic bind
   | t           -- (generalized) tail call
   | case v of alt1; ...; altn -- conditional branch

alt ::= C(v1, ..., vn) -> b(a1, ..., an) -- constructor match
      _                -> b(a1, ..., an) -- default match

def ::= b(v1, ..., vn) = c -- block entry point
      | k{v1, ..., vn} v = c -- closure definition
      | v <- t              -- top-level constant

```

Code

51

## MIL, a monadic intermediate language

```

a ::= v          -- variable
   | n          -- integer literal

t ::= return a   -- simple value
   | p((a1, ..., an)) -- primitive call
   | b(a1, ..., an) -- block call
   | C(a1, ..., an) -- data constructor
   | k{a1, ..., an} -- closure constructor
   | f @ a       -- enter closure
   | m[a1, ..., an] -- thunk constructor
   | invoke a    -- invoke closure

c ::= v <- t; c   -- monadic bind
   | t           -- (generalized) tail call
   | case v of alt1; ...; altn -- conditional branch

alt ::= C(v1, ..., vn) -> b(a1, ..., an) -- constructor match
      _                -> b(a1, ..., an) -- default match

def ::= b(v1, ..., vn) = c -- block entry point
      | k{v1, ..., vn} v = c -- closure definition
      | v <- t              -- top-level constant

```

Tails

52

## MIL, a monadic intermediate language

```

a ::= v          -- variable
   | n          -- integer literal

t ::= return a   -- simple value
   | p((a1, ..., an)) -- primitive call
   | b(a1, ..., an) -- block call
   | C(a1, ..., an) -- data constructor
   | k{a1, ..., an} -- closure constructor
   | f @ a       -- enter closure
   | m[a1, ..., an] -- thunk constructor
   | invoke a    -- invoke closure

c ::= v <- t; c   -- monadic bind
   | t           -- (generalized) tail call
   | case v of alt1; ...; altn -- conditional branch

alt ::= C(v1, ..., vn) -> b(a1, ..., an) -- constructor match
      _                -> b(a1, ..., an) -- default match

def ::= b(v1, ..., vn) = c -- block entry point
      | k{v1, ..., vn} v = c -- closure definition
      | v <- t              -- top-level constant

```

Atoms

53

## ... to MIL programs

```

id      ← k0{}
k0{x}   = b0(x)
b0(x)   = return x

map      ← k4{f}
k4{f}   = k5{f}
k5{f} xs = b2(f, xs)
b2(f, xs) = case xs of
              Nil() → b3()
              Cons(y, ys) → b4(f, y, ys)

b3()    = Nil()
b4(f, y, ys) = z ← f @ y
              m ← map @ f
              zs ← m @ ys
              Cons(z, zs)

compose ← k1{f}
k1{f}   = k2{f}
k2{f} g = k3{f, g}
k3{f, g} x = b1(f, g, x)
b1(f, g, x) = u ← g @ x
              f @ u

```

unknown function call

known function call

54

## ... to MIL programs

```

id      ← k0{ }
k0{ } x = b0(x)
b0(x)   = return x

map      ← k4{ }
k4{ } f  = k5{ f }
k5{ f } xs = b2(f, xs)
b2(f, xs) = case xs of
  Nil() → b3()
  Cons(y, ys) → b4(f, y, ys)

b3() = Nil()
b4(f, y, ys) = z ← f @ y
               m ← map @ f
               zs ← m @ ys
               Cons(z, zs)

```

unknown function call

55

## ... to optimized MIL programs

```

id      ← k0{ }
k0{ } x = b0(x)
b0(x)   = return x

map      ← k4{ }
k4{ } f  = k5{ f }
k5{ f } xs = b2(f, xs)
b2(f, xs) = case xs of
  Nil() → b3()
  Cons(y, ys) → b4(f, y, ys)

b3() = Nil()
b4(f, y, ys) = z ← f @ y
               m ← k5{ f }
               zs ← m @ ys
               Cons(z, zs)

```

unknown function call

known function call

56

## ... to optimized MIL programs

```

id      ← k0{ }
k0{ } x = b0(x)
b0(x)   = return x

map      ← k4{ }
k4{ } f  = k5{ f }
k5{ f } xs = b2(f, xs)
b2(f, xs) = case xs of
  Nil() → b3()
  Cons(y, ys) → b4(f, y, ys)

b3() = Nil()
b4(f, y, ys) = z ← f @ y
               m ← k5{ f }
               zs ← m @ ys
               Cons(z, zs)

```

unknown function call

57

## ... to optimized MIL programs

```

id      ← k0{ }
k0{ } x = b0(x)
b0(x)   = return x

map      ← k4{ }
k4{ } f  = k5{ f }
k5{ f } xs = b2(f, xs)
b2(f, xs) = case xs of
  Nil() → b3()
  Cons(y, ys) → b4(f, y, ys)

b3() = Nil()
b4(f, y, ys) = z ← f @ y
               m ← k5{ f }
               zs ← b2(f, ys)
               Cons(z, zs)

```

pure, dead code

58

## ... to optimized MIL programs

```

id      ← k0{ }
k0{ } x = b0(x)
b0(x)   = return x

map      ← k4{ }
k4{ } f  = k5{ f }
k5{ f } xs = b2(f, xs)
b2(f, xs) = case xs of
  Nil() → b3()
  Cons(y, ys) → b4(f, y, ys)

b3() = Nil()
b4(f, y, ys) = z ← f @ y
               zs ← b2(f, ys)
               Cons(z, zs)

```

59

## Constant/copy propagation / left monad law:

$x \leftarrow \text{return } a; c \longrightarrow [a/x] c$

## Tail call introduction / right monad law:

$x \leftarrow t; \text{return } x \longrightarrow t$

60

**Prefix inlining:**  $b(x) = v_1 \leftarrow t_1; \dots; v_n \leftarrow t_n; t$   
(monad associativity law)

$v \leftarrow b(x); c \longrightarrow v_1 \leftarrow t_1; \dots; v_n \leftarrow t_n; v \leftarrow t; c$

**Suffix inlining:**

$v \leftarrow t_0; b(x) \longrightarrow v \leftarrow t_0; v_1 \leftarrow t_1; \dots; v_n \leftarrow t_n; t$

61

**Wildcard introduction:**

$v \leftarrow t; c \longrightarrow \_ \leftarrow t; c$  ( $v$  not free in  $c$ )

**Dead tail elimination:**

$\_ \leftarrow t; c \longrightarrow c$  ( $t$  pure)

62

**Identity laws:**

$or((x, 0)) \longrightarrow return\ x$

**Idempotence:**

$or((x, x)) \longrightarrow return\ x$

**Common subexpression elimination:**

$\{y = t\}$   
 $v \leftarrow t \longrightarrow v \leftarrow return\ y$

63

**Constant folding:**

( $M, N$  constants)

$\{x=M, y=N\}$   
 $v \leftarrow and((x, y)) \longrightarrow v \leftarrow return\ (M \& N)$

**Argument ordering:**

( $M, N$  constants)

$v \leftarrow and((M, x)) \longrightarrow v \leftarrow and((x, M))$

**Associative folding:**

$\{v=and((w, M))\}$   
 $x \leftarrow and((v, N)) \longrightarrow x \leftarrow and((w, M \& N))$

$(u \& M) \& N = u \& (M \& N)$

64

**Distributive folding:**

( $M, N, P$  constants)

$\{v=or((u, M))\}$   
 $x \leftarrow and((v, N)) \longrightarrow \begin{array}{l} t \leftarrow and((u, N)) \\ x \leftarrow or((t, M \& N)) \end{array}$  ( $t$  new)

$(u \mid M) \& N = (u \& N) \mid (M \& N)$

$\{v=or((u, M)), w=and((v, N))\}$   
 $x \leftarrow or((w, P)) \longrightarrow \begin{array}{l} t \leftarrow and((u, N)) \\ x \leftarrow or((w, (M \& N) \mid P)) \end{array}$

$((u \mid M) \& N) \mid P = (u \& N) \mid ((M \& N) \mid P)$

65

**Known closure:**

$\{f = k\{x\}\}$   
 $v \leftarrow f\ @\ y \longrightarrow v \leftarrow t$  if  $k\{x\}\ y = t$

**Known thunk:**

$\{t = m\{x\}\}$   
 $v \leftarrow invoke\ t \longrightarrow v \leftarrow m(x)$

**Known constructor:**

$\{v = C(x)\}$   
 $case\ v\ of$   
 $\quad C(x) \rightarrow b(x, y)$   
 $\quad \dots \longrightarrow b(x, y)$

66

## Derived blocks - trailing enter:

$v \leftarrow b(x); v @ a \rightarrow b1(x, a)$   
 $b(x) = \begin{matrix} v_1 \leftarrow t_1 \\ \dots \\ v_n \leftarrow t_n \\ t \end{matrix} \rightarrow b1(x, a) = \begin{matrix} v_1 \leftarrow t_1 \\ \dots \\ v_n \leftarrow t_n \\ v \leftarrow t \\ v @ a \end{matrix}$


67

## Derived blocks - known constructors:


$\{u = C(p), f = k\{x\}\} \rightarrow b1(p, x, r)$   
 $b(u, f, r) = c \rightarrow \begin{matrix} b1(p, x, r) \\ = u \leftarrow C(p) \\ f \leftarrow k\{x\} \\ c \end{matrix}$

68


## MapPage system call



69



70



71

## PrioSet

```

prioSet = \i prio -> do writeDef (at prioSet i) prio
                        writeDef (at prioSet prio) i


insertPriority = \prio -> do s <- readDef prioSetSize
                        writeDef prioSetSize (add s 1)
                        heapPairUp (mod s) prio

heapPairUp = \i prio ->
  case doc i of
    Nothing -> prioSet 0 prio
    Just j -> do parent <- ret (shiftR j 1)
                pprio <- readDef (at prioSet parent)
                if is pprio prio then
                  prioSet i pprio
                else
                  heapPairUp parent prio

removePriority = \prio ->
  do s <- readDef prioSetSize
  writeDef prioSetSize (sub s 1)
  rprio <- readDef (at prioSet (mod s 1))
  if neq prio rprio then
    i <- readDef (at prioSet prio)
    heapPairDown i rprio (mod s 2)
  rprio <- readDef (at prioSet i)
  heapPairUp i rprio

heapPairDown = \i prio last ->
  do let u = undefined // <- ret (undefined)
  case last (add (mod 2 u) 1) last of
    Nothing -> prioSet i prio // Look for a left child
    Just 1 -> do lprio <- readDef (at prioSet i)
                if is lprio prio then
                  prioSet i lprio // i has no children
                else
                  do rprio <- readDef (at prioSet i)
                      if is rprio prio then
                        prioSet i rprio // i has a left child
                      else
                        prioSet i prio // i has a right child
    Just 2 -> do // i has two children
                do rprio <- readDef (at prioSet r)
                  if ge prio lprio && ge prio rprio then
                    prioSet i prio
                  else if ge prio rprio then
                    prioSet i lprio // left is higher
                  else
                    prioSet i rprio // right is higher
  heapPairUp rprio last
  
```

main = \u -> do insertPriority u  
removePriority u



72



## Practical and Effective Higher-Order Optimizations

Lars Bergstrom  
Mozilla Research \*  
larsberg@mozilla.com

Matthew Fluet  
Matthew Le  
Rochester Institute of Technology  
{mtf,ml9951}@cs.rit.edu

John Reppy  
Nora Sandler  
University of Chicago  
{jhr,nlsandler}@cs.uchicago.edu

### Abstract

Inlining is an optimization that replaces a call to a function with that function's body. This optimization not only reduces the overhead of a function call, but can expose additional optimization opportunities to the compiler, such as removing redundant operations or unused conditional branches. Another optimization, copy propagation, replaces a redundant copy of a still-live variable with the original. Copy propagation can reduce the total number of live variables, reducing register pressure and memory usage, and possibly eliminating redundant memory-to-memory copies. In practice, both of these optimizations are implemented in nearly every modern compiler.

These two optimizations are practical to implement and effective in first-order languages, but in languages with lexically-scoped first-class functions (aka, closures), these optimizations are not available to code programmed in a higher-order style. With higher-order functions, the analysis challenge has been that the environment at the call site must be the same as at the closure capture location, up to the free variables, or the meaning of the program may change. Olin Shivers' 1991 dissertation called this family of optimizations *Since B*, and he proposed one analysis technique

**Categories and Subject Descriptors** D.3.0 [Programming Languages]: General; D.3.2 [Programming Languages]: Language Classifications—Applicative (functional) languages; D.3.4 [Programming Languages]: Processors—Optimization

**Keywords** control-flow analysis, inlining, optimization

### 1. Introduction

All high level programming languages rely on compiler optimizations to transform a language that is convenient for software developers into one that runs efficiently on target hardware. Two such common compiler optimizations are copy propagation and function inlining. Copy propagation in a language like ML is a simple substitution. Given a program of the form:

```
let val x = y
in
  x+2*y
end
let val x = y
```

We want to propagate the definition of *x* to its uses, resulting in:

```
let val x = y
```

## A "challenging example"

```
fun mk i =
  let
    fun g j = j + i
    fun f (h : int -> int, k) =
      (h (k * i))
  in
    (f, g)
  end
val (f1, g1) = mk 1
val (f2, g2) = mk 2
val res = f1 (g2, 3)
```

res <- return 4

"While this example is obviously contrived, this situation occurs regularly in idiomatic higher-order programs and the inability to handle the environment problem in general is a limit in most compilers, leading developers to avoid higher-order language features in performance-critical code."

80

## An earlier example in the paper

```
let
  fun emit x = print (Int.toString x)
  fun fact i m k =
    if i=0 then k m
    else fact (i-1) (m*i) k
in
  fact 6 1 emit
end
```

```
b2() =
  t46 <- intToStr((720))
  print((t46))

main <-
  b2()
```

```
let
  fun emit x = print (Int.toString x)
  fun fact i m k =
    if i=0 then emit m
    else fact (i-1) (m*i) emit
in
  fact 6 1 emit
end
```

81

## ... generalized

```
emit x = print (intToStr x)

fact i m k
  = if i=0
    then k m
    else fact (i-1) (m*i) k

main m n = fact m n emit
```

```
b7(x) =
  t14 <- intToStr((x))
  print((t14))

b10(m, i) =
  t9 <- add((i, -1))
  t12 <- mul((m, i))
  t15 <- eq((t9, 0))
  case t15 of
    True -> b7(t12)
    False -> b10(t12, t9)
```

```
b8(i, m) =
  t5 <- eq((i, 0))
  case t5 of
    True -> b7(m)
    False -> b10(m, i)
```

```
k36{m} n = b8(m, n)
```

```
k37{ } m = k36{m}
```

```
main <-
  k37{ }
```

+ Third arg for fact has gone!

+ Direct call to (inlined) emit

- Some duplication of code ...

82

## Very good optimization, for low effort?

- 50% chance that I just got lucky with my choice of examples ...
- 50% chance that this is a result of aggressive inlining and specialization
- 1% chance that there is something fundamentally new about this approach to optimization ...
- 101% chance that I don't understand percentages :-)

83

## Summary

- Back to the three main questions for CEMLaBS:
  - **Feasibility:** Still chipping away ... but getting closer!
  - **Benefit:** Good evidence that we will benefit from the use of functional language features
    - +Types
    - +Higher-order functions
  - **Performance:** acceptable performance may be within reach
    - +We can generate good quality code, even when lambdas are used in fundamental ways
    - +Some code duplication (but, so far, this is entirely tolerable for our specific use case ...)

84