CS 410/510

Languages & Low-Level Programming

Mark P Jones
Portland State University

Spring 2016

Week 1: Introduction, Assembly Language

Introduction and Goals

Origins

- For a long time, a group of us at PSU have been looking at the role that high-level programming languages can play in the construction of (very) low-level software.

- By using **high-level languages**, we can hope to increase programmer productivity, and improve software quality

- By focussing on very **low-level software**, we hope to provide strong foundations for the complete software stack
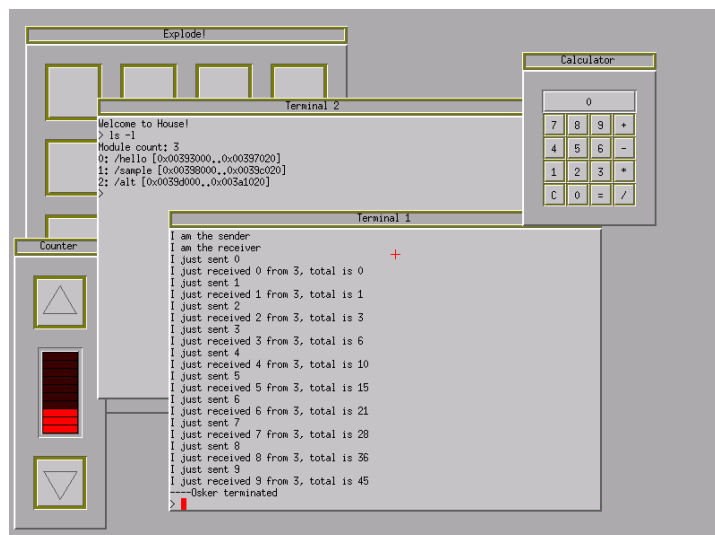
House (2005)

Kernel, GUI, drivers, network stack, and apps

Boots and runs in a bare metal environment

… all written in Haskell, a "purely functional" programming language

Why "House"?

"The Haskell User's Operating System Environment"

You are more secure in a house …

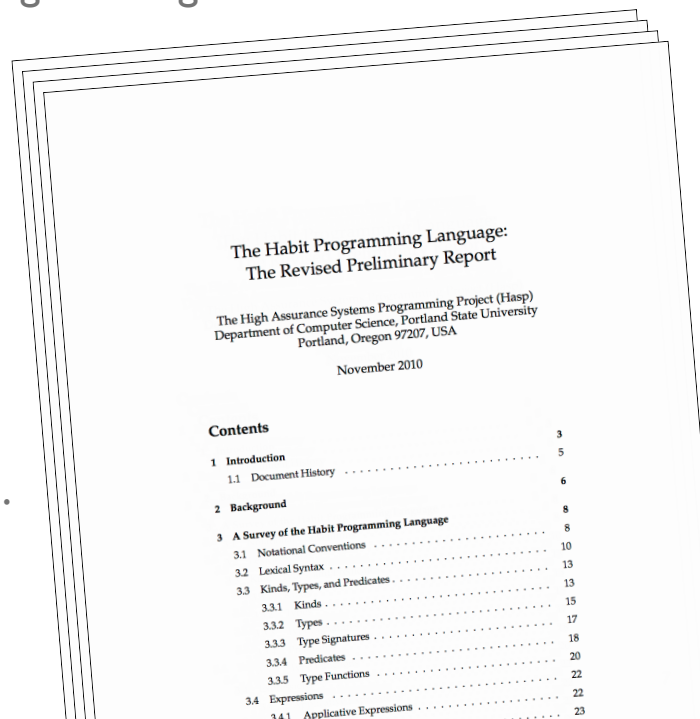than if you only have Windows …

Performance concerns

- By design, higher-level languages abstract away from the details of how the underlying machine works

- Can we obtain the levels of performance and predictability that are typically required/expected in the systems programming domain?

- Can we write good systems software in a language that intentionally distances users from details of memory layout, representation, instruction selection, alignment, caching, etc.?

- Traditional approaches to building system software resort to using old, low-level languages like assembly and C

- Do "modern" languages have anything to offer in this area?

The Habit programming language

- "a dialect of Haskell that is designed to meet the needs of high assurance systems programming"

- How do you design a programming language for a specific domain?

- Experiment with existing languages
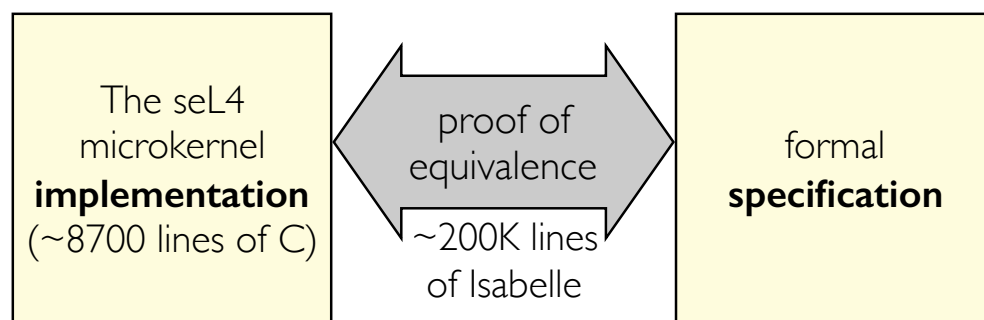
- Understand the domain …

The Habit Programming Language:
The Revised Preliminary Report

The High Assurance Systems Programming Project (Hasp)
Department of Computer Science, Portland State University
Portland, Oregon 97207, USA

November 2010

The seL4 experience

- In 2009, a group from NICTA, UNSW, and OK Labs in Australia announced seL4, as "the world's first operating system kernel with an end-to-end proof of implementation correctness and security enforcement."

The seL4 microkernel **implementation** (~8700 lines of C) — proof of equivalence ~200K lines of Isabelle — formal **specification**

- A landmark achievement for formal verification, and a strong foundation for building trustworthy systems

seL4 and capabilities

- Even without the verification result, the design of seL4 is interesting in its own right:

  - seL4 is a "capability enhanced" version of an earlier microkernel design called L4

  - The "capability" abstraction in seL4 provides facilities for implementing "least privilege" security policies and novel mechanisms for controlling resource usage

Proof vs performance?

- Security properties established in the seL4 verification include:

  - Absence of code injection attacks

  - Absence of buffer overflows

  - Absence of null pointer dereferences

  - …

- Many of these properties could be established "for free" if the implementation had been written in a "safer" language

- How might things be different if we built something like seL4 in Habit?

The CEMLaBS project

- "Using a Capability-Enhanced Microkernel as a Testbed for Language-Based Security"

- Started October 2014, funded by The National Science Foundation

- Three main questions:

  - **Feasibility**: Is it possible to build an inherently "unsafe" system like seL4 in a "safe" language like Habit?

  - **Benefit**: What benefits might this have, for example, in reducing verification costs?

  - **Performance**: Is it possible to meet reasonable performance goals for this kind of system?

CEMLaBS workplan

- The original plan was to use Habit to design and build HAL4, a "capability-enhanced" microkernel in the style of seL4 …

- At the time, neither the specification or the implementation of seL4 was publicly available … but some key aspects of its design had been revealed in published papers, dissertations, and talks

- In July 2014, seL4 was released as open source software!

- Now we are free to build HAL4 as an implementation of seL4, which will allow use to make more direct comparisons between the two implementations …

CEMLaBS status

- The original HAL4 implementation is incomplete (and will need to be reworked to match the seL4 specification)

- The Habit compiler is not complete (a full compiler pipeline has been implemented, but it is not feature complete, and it is not robust)

- We'll be working on addressing both of these shortcomings this coming summer

- In the meantime, there is still much to study …

Course description

- **An overview of conventional low-level programming techniques:**
  - Bare metal programming
  - Fundamental programmable hardware components

- **Case studies of practical microkernel implementations:**
  - OS abstractions (address spaces, threads, capabilities, …)
  - The L4 and seL4 microkernels

- **Reflections on the design of programming languages for use in this application domain:**
  - Assembly, C, Rust, Habit, domain specific languages, …

Course learning objectives

Upon the successful completion of this course, students will be able to:

1. Write simple programs that can run in a bare-metal environment using low-level programming languages.

2. Discuss common challenges in low-level systems software development, including debugging in a bare-metal environment.

3. Explain how conventional operating system features (multiple address spaces, context switching, protection, etc.) motivate the desire for (and benefit from) hardware support.

Course learning objectives, continued

4. Develop code to configure and use programmable hardware components such as a memory management unit (MMU), interrupt controller (PIC), and interval timer (PIT).

5. Describe the key steps in a typical boot process, including the role of a bootloader.

6. Describe the motivation, implementation, and application of microkernel abstractions for managing address spaces, threads, and interprocess communication (IPC).

7. Explain the use and implementation of capabilities in access control and resource management.

8. Develop programs using the capability abstraction provided by the seL4 microkernel.

Course learning objectives, continued

9. Illustrate the use of a range of domain specific languages in the development of systems software.

10. Use practical case studies to evaluate and compare language design proposals.

11. Describe features of modern, high-level programming languages—including abstract datatypes and higher-order functions—and show how they can be leveraged in the construction of low-level software.

12. Explain how the requirements of low-level systems programming motivate the desire for (and benefit from) language-based support.

The "programming languages" perspective

- We will survey and evaluate a range of programming languages during this course:
  - Low-level machine and assembly languages
  - Systems programming languages (e.g., C, Rust, …)
  - Object-oriented languages (e.g., the seL4 API)
  - Domain specific languages
  - Functional languages (e.g., Habit, Haskell, …)
- What are the driving needs of the systems domain?
- How can a programming language design best meet those needs?

Context

- Basic Platform: Generic "IBM PC" compatible
  - 32 bits … not 64
  - IA32 … not x86_64 or ARM
  - BIOS … not EFI or UEFI
  - `int` and `iret` … not `sysenter`/`sysexit`
  - PIC … not APIC
  - No PAE, PCI, ACPI, MMX, SSE, SMM, SMP, VTx, …
  - etc., …
- Already complicated enough for our purposes!
- Well supported by current hardware, emulators, and tools
- Underlying concepts still very broadly applicable

Development environment

- Ubuntu Linux
  - Week 1: using the LinuxLab machines (Mac OS also an option)
  - Weeks 2+: using a VirtualBox virtual machine, preconfigured with appropriate development tools (can be used on Linux, Mac OS, or Windows)
- Bare metal emulation using the QEMU emulator

Rough schedule

Week	Topic
1	Assembly language programming
2	Bare metal programming
3	Hardware support for OS abstractions
4	Memory management & protection
5	Case Study: L4 use & implementation
6	
7	Case Study 2: seL4 use & implementation
8	
9	Language design for low-level programming
10	

Lectures and labs

- Monday Lectures (UTS 210), Wednesday Labs (FAB 88-09)
  - Students are expected to attend all lectures and labs
  - Sets of guided practical exercises will be provided for each lab
  - Students will work on labs during lab sessions with instructor supervision
  - Students are expected to complete each set of exercises before the following week's lab session
  - We will use D2L forums for discussion of exercises outside lab sessions
  - Lab sessions may also be used for direct instruction, as necessary

Exams

- No midterm or final exams
  - Students are still required to attend the "final" on 6/9
  - This will be used for student presentations or demos, or else for other instruction

Portfolio assessment

- Final grades will be assigned on the basis of a "Portfolio" comprising the results of lab work and independent projects
- Based on previously listed collection of 12 course objectives
- General process:
  - Student uploads work item (typically a tar ball or zip file), accompanied by a write up explaining how the submitted work addresses the expectations of one or more objectives
  - Grading assigns a score of N/A, Below, Meets or Exceeds for each objective (with +/- versions)
  - Different final letter grades awarded depending on the combination of individual scores

Portfolio assessment, continued

- Students may request consideration of at most two objectives in any given week (Saturday-Friday)
    - discourages procrastination, balances grading load
- I will attempt to provide feedback on submitted items within a week
- Students may resubmit in an attempt to improve score for any given objective
    - (Each resubmission counts as a new submission)
- Completion of lab exercises for a given week can typically be used to satisfy the "Meets" condition for specific objectives
- Process starts in Week 2, runs until end of finals week

General approach and caveats

- This course will be experimental, open-ended, hands-on, and interactive
- Your flexibility, patience, and tolerance of informal or underspecified aspects of the course will be very much appreciated
- Questions, suggestions, and contributions are welcome at any point

An introduction to
IA32 assembly
language programming

What is IA32?

- We'll be using the IA32 (x86) architecture as our main target:

  - A "32-bit" instruction set

  - Broadly adopted by:

    - processors from Intel, AMD, Via, ...

    - laptops, desktops, servers, gaming consoles, ...

    - Linux, Mac OS X, Windows, …

- Arguably, a bit dated … but still very relevant, and a good platform for learning and exploration

- (… and one of the two architectures supported by seL4)

Other architectures:

- Not to be confused with:

  - x86-64/AMD64: a 64 bit architecture supported (in addition to IA32) by more recent AMD/Intel designs

  - IA64: a completely different 64-bit Intel architecture (Itanium)

  - ARM: widely used in phones, tablets, and more

  - IBM Power: used in Xbox 360, PS3, Wii, servers, and more

  - SPARC: used by some of the college's Unix servers

- Except for x86-64, you can't run IA32 code directly on a machine that uses one of these alternatives instruction sets!

Notes

- No prior or in-depth knowledge of IA32 programming will be assumed

- We will only use a small subset of the full instruction set

- If you're looking to become an expert on IA32 programming, you'll want to look for another class!

- We'll be using the *AT&T syntax* for IA32 assembly language rather than the *Intel syntax*. This is the default syntax used by the free GNU tools in Linux, MacOS, and DJGPP or Cygwin on Windows, and others

- For simplicity, I recommend: ssh yourid@linuxlab.cs.pdx.edu (or, on Windows, the equivalent using PuTTY)

A (greatly) simplified view of the IA32

CPU

instruction pointer

general purpose registers

8 general purpose 32-bit registers provide fast temporary storage for integers, pointers, ...

ALU

address / 32

data / 8

Memory

up to 4 GB (2^{32} bytes) for stored programs and data
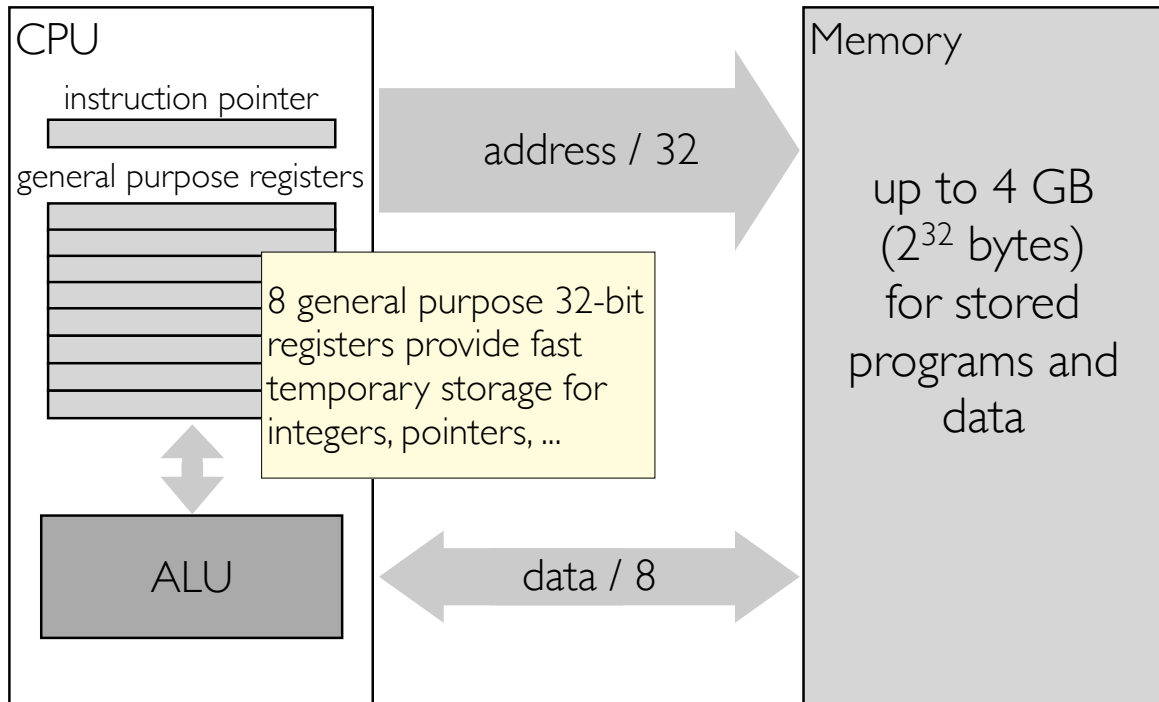
Programming for IA32

- In concrete terms, an IA32 program is just a collection of byte values (*machine code*)

- Once it has been loaded in to memory, the processor can *execute* a program by interpreting the byte values as *instructions* for the processor to act on

- For practical purposes, we will usually write IA32 programs in a textual format called *assembly language* that is easier to read than the raw byte values

- The program that translates assembly language programs in to machine code is called an *assembler*

The GNU assembler, as

- Assembly code goes in files with a .s suffix

- We will typically use gcc to invoke the assembler

  ```
  gcc –m32 —o output assemblyCode.s extras.c
  ```

- You can also invoke the assembler directly: detailed documentation is available from:
  http://sourceware.org/binutils/docs/as/
  For IA32 programming, look in particular at the section on "80386 Dependent Features"

An assembly code listing

```
        .globl  f
f:
        pushl   %ebp
        movl    %esp,%ebp
        pushl   %ebx
        movl    8(%ebp), %ebx

        movl    $0, %eax      # initialize length count in eax

        jmp     test
loop:   incl    %eax          # increment count
        addl    $4, %ebx      # and move to next array element

test:   movl    (%ebx), %ecx # load array element
        cmpl    $0, %ecx      # test for end of array
        jne     loop          # repeat if we're not done ...

        popl    %ebx
        movl    %ebp,%esp
        popl    %ebp
        ret
```

Assembly code

An assembly code listing

```
                        .globl  f
                 f:
0000 55                 pushl   %ebp
0001 89E5               movl    %esp,%ebp
0003 53                 pushl   %ebx
0004 8B5D08             movl    8(%ebp), %ebx

0007 B8000000           movl    $0, %eax       # initialize length count in eax
     00
000c EB04               jmp     test
000e 40         loop:   incl    %eax           # increment count
000f 83C304             addl    $4, %ebx       # and move to next array element

0012 8B0B      test:    movl    (%ebx), %ecx # load array element
0014 83F900             cmpl    $0, %ecx       # test for end of array
0017 75F5               jne     loop           # repeat if we're not done ...

0019 5B                 popl    %ebx
001a 89EC               movl    %ebp,%esp
001c 5D                 popl    %ebp
001d C3                 ret
```

Machine code   Assembly code

35

---

addresses
/offsets   labels

```
                        .globl  f  directive
                 f:
0000 55                 pushl   %ebp
0001 89E5               movl    %esp,%ebp         comments
0003 53                 pushl   %ebx
0004 8B5D08             movl    8(%ebp), %ebx

0007 B8000000           movl    $0, %eax       # initialize length count in eax
     00
000c EB04               jmp     test
000e 40         loop:   incl    %eax           # increment count
000f 83C304             addl    $4, %ebx       # and move to next array element

0012 8B0B      test:    movl    (%ebx), %ecx # load array element
0014 83F900             cmpl    $0, %ecx       # test for end of array
0017 75F5               jne     loop           # repeat if we're not done ...

0019 5B                 popl    %ebx
001a 89EC               movl    %ebp,%esp
001c 5D                 popl    %ebp
001d C3                 ret
```
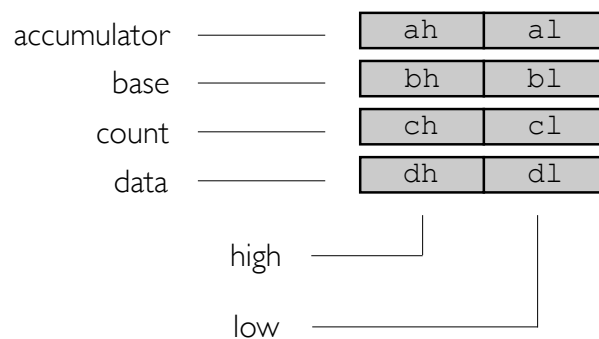
machine   instructions
code

36

IA32 registers

8-bit registers (holding a single byte, 0-255)

accumulator ——————	ah	al
base ——————	bh	bl
count ——————	ch	cl
data ——————	dh	dl

high —————————

low —————————

Introduced in 1978 as part of the 8086 architecture

16-bit registers ("word")

accumulator	ax	ah	al
base	bx	bh	bl
count	cx	ch	cl
data	dx	dh	dl
source index	si		
destination index	di		
base pointer	bp		
stack pointer	sp		

Introduced in 1978 as part of the 8086 architecture

32-bit registers ("double word")

accumulator	eax	ax	ah	al
base	ebx	bx	bh	bl
count	ecx	cx	ch	cl
data	edx	dx	dh	dl
source index	esi	si		
destination index	edi	di		
base pointer	ebp	bp		
stack pointer	esp	sp		

"e" for extended

sometimes referred to as "long word"s

Introduced in 1985 as part of the 80386 architecture

Special vs. general purpose registers

- `eip`: the instruction pointer register

- `esp`: the stack pointer register

- `eflags`: the flags register, stores information about the results of the most recent arithmetic or logic instruction

- Other registers can typically be used for any purpose (although some instructions—division, for example—work only with specific registers)
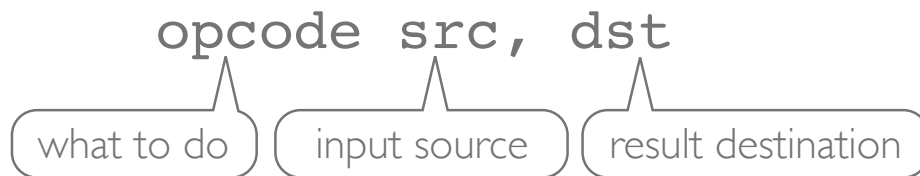
IA32 instructions

Instruction format

- A typical IA32 instruction has the form:

$$\text{opcode src, dst}$$

(what to do) (input source) (result destination)

- A suffix on the opcode indicates the size of the data that is being operated on:

  - 32-bit values use the suffix **l**(ong)

  - 16-bit values use the suffix **w**(ord)

  - 8-bit values use the suffix **b**(yte)

43

Addressing modes

- **Register access**, reg:
  - `%eax`: the value in register `eax`
  - Can typically use any registers except `eip` and `eflags`
- **Memory access**, mem:
  - `var`: the value in the memory location at address var
  - `(%eax)`: the value in memory at the address in `eax`
  - `8(%eax)`: the value in memory at the address given by adding 8 to the value in `eax`
- **Immediate**, immed:
  - `$42`: the constant value 42 (decimal; use `$0x2A` for hex)
  - `$var`: the address of memory location var

44

Directives for "declaring" variables

```
        .data           # put variables in the "data" section
                        # (code usually goes in .text)

        .align  4       # make sure address is multiple of 4
myvar:  .long   42      # Simple variable, initialized to 42

        .global days    # A globally accessible array of ints
days:   .long    31, 28, 31, 30, 30, 30
        .long    31, 31, 30, 31, 30, 31

scratch:.space  4*100   # reserve uninitialized space

medium: .long   123     # a 32-bit integer (takes 4 bytes)
regular:.short  123     # a 16-bit integer (takes 2 bytes)
small:  .byte   123     # an 8-bit integer (takes 1 byte)
```
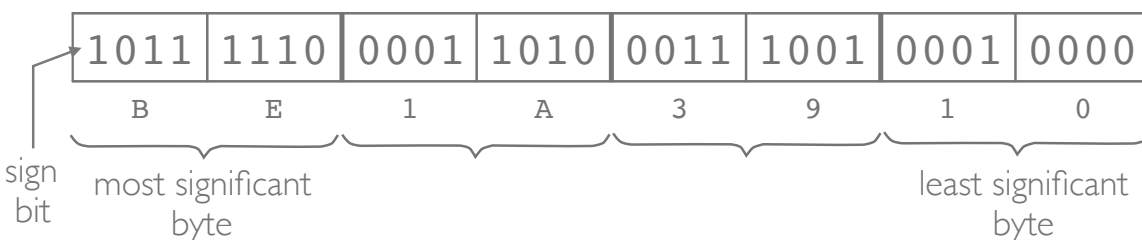
How values are stored in memory

- A double word holds 32 binary digits ("bits") (i.e., 4 bytes)

1011	1110	0001	1010	0011	1001	0001	0000
B	E	1	A	3	9	1	0

sign bit — most significant byte — least significant byte

- 0xBE1A3910 can be interpreted as -1,105,577,712 (signed) or 3,189,389,584 (unsigned)

- Stored in memory with the least significant byte at the lowest address ("little endian"):

stored byte	0x10	0x39	0x1A	0xBE
address	400	401	402	403

IA32 instructions:
data movement

Move instructions

- Copy data from a source to a destination (where `X` is one of the size suffixes: `b,w,l`):

$$\texttt{movX src, dst}$$

- Any of the following combinations of arguments is allowed:

    `movX` reg, (reg | mem)

    `movX` mem, reg

    `movX` immed, (reg | mem)

- Note that you can't move mem to mem in one instruction

Examples

Suppose that the memory (starting at address 0) contains the following (four byte) values:

8	6	2	8	0	2	4	1	7	3	4	5	6
0	4	8	12	16	20	24	28	32	36	40	44	48

Then

instruction	contents of eax
movl $12, %eax	12
movl (%eax), %eax	8
movl 8(%eax), %eax	0

Zero and sign-extension

- Suppose we want to copy a value from a 16-bit register in to a 32-bit register

```
stu ....... xyz          ???????????????? stu ....... xyz
     ax                              eax
```
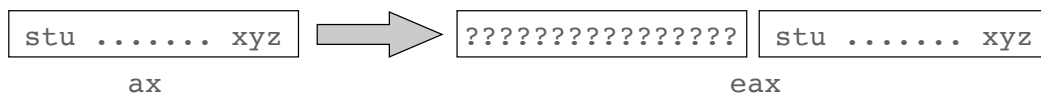
- Two common strategies:

  - Zero extension: for unsigned values

```
stu ....... xyz          0000000000000000 stu ....... xyz
     ax                                eax
```

  - Sign extension: for signed values

```
stu ....... xyz          ssssssssssssssss stu ....... xyz
     ax                                eax
```

Move with sign, move with zero extension

- Copy from source to larger destination with sign extension:

$$\texttt{movsFT src, dst}$$

- Copy from source to larger destination with zero extension:

$$\texttt{movzFT src, dst}$$

- `F` and `T` are the "from" and "to" sizes (either b, w, or l)

- Valid combinations: bw, bl, or wl

- Examples:

```
movsbw %al, %dx     # byte to word
movzwl %ax, %edx    # word to long
```

Scaled indexed addressing

- [base]([reg_1],reg_2 [,index])

  a memory operand whose address is the value in reg_1, *plus* the specified base constant, *plus* the value of reg_2 *times* the index (which must be 1, 2, 4, or 8)

- Any of the parts in [...] can be omitted

- Examples:

  `(eax,ebx,4)` the ebx[th] element in the array of 32-bit words starting at the address in eax

  `days(,ebx,4)` the ebx[th] element in the array of 32-bit words starting at the address days

More examples

Suppose that the memory (starting at address 0) contains the following (four byte) values:

8	6	2	8	0	2	4	1	7	3	4	5	6
0	4	8	12	16	20	24	28	32	36	40	44	48

Then

instruction	eax	ebx
movl $12, %eax	12	
movl 8(%eax), %ebx	12	2
movl 12(%eax,%ebx,4), %eax	7	2

The `lea` (load effective address) instruction

- Load the address of the source operand (must be memory) to a destination (where X is one of the size suffixes: b,w,l):

$$\texttt{leaX src, dst}$$

- Can also be used to co-opt the addressing mode circuitry into performing arithmetic operations:

```
leal 4(%eax),%eax           # eax += 4
leal 1(%eax,%eax,2),%eax # eax = 3*eax + 1
leal 1(%eax,%eax), %eax  # eax = 2*eax + 1
leal 4(,%eax,8), %eax    # eax = 8*eax + 4
```

- These instructions just do an address calculation and do not attempt to read the data at that address.

The exchange instruction

- Exchange data between two locations

$$\texttt{xchgX} \ (\text{reg} \mid \text{mem})\texttt{, reg}$$

- Consider the following instructions in a high-level language:

```
int tmp = x;
x       = y;
y       = tmp;
```

- If `x` and `y` are held in registers, then a "clever enough" compiler can translate this code into a single `xchgl` instruction

The instruction pointer, `eip`

- The `eip` register holds the address of the next instruction to be executed

- As the processor reads each instruction, it increments the value in `eip` by the appropriate number of bytes to point to the following instruction

- This mechanism allows the processor to execute a sequence of instructions stored in contiguous locations in memory

- What would happen if we "move" a different value in to `eip`?

Jumping and labels

- We can transfer control and start executing instructions at address addr by using a jump instruction

$$\texttt{jmp addr}$$

- Labels can be attached to instructions in an assembly language program:

```
            jmp b
a:          jmp c
b:          jmp a
c:          ...
```

- Modern, pipelined machines work well with sequences of instructions that appear in consecutive locations. Jumps can be expensive: one of the goals of an optimizing compiler is to avoid unnecessary jumps.

IA32 instructions:
arithmetic and logic operations

Arithmetic instructions

- Combine a given `src` with a given `dst` value and leave the result in `dst`:

```
▶  addX   src, dst
   subX   src, dst   ⎫ integer arithmetic
   imulX  src, dst   ⎭ (signed)
   andX   src, dst   ⎫
   orX    src, dst   ⎬ bitwise arithmetic
   xorX   src, dst   ⎭
```

- Similar to `dst += src, dst -= src`, etc.. in C/C++

Examples

- To compute $x^2 + y^2$ and store the result in `z`:

```
▶  movl  x, %eax
   imull %eax, %eax
   movl  y, %ebx
   imull %ebx, %ebx
   addl  %ebx, %eax
   movl  %eax, z
```

register	contents
eax	x^2+y^2
ebx	y^2

```
           .data
   x:   .long   4
   y:   .long   3
   z:   .long   0
```
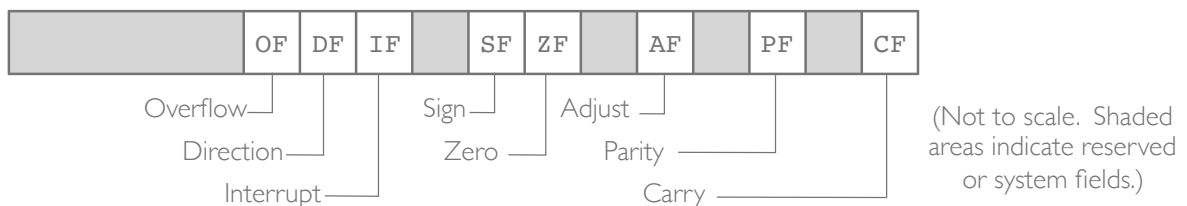
IA32 instructions:
conditional execution

Flags

- In addition to performing the required operation, arithmetic instructions also change bits in the `eflags` register

				OF	DF	IF		SF	ZF		AF		PF		CF

Overflow —
Direction —
Interrupt —
Sign —
Zero —
Adjust —
Parity —
Carry —

(Not to scale. Shaded areas indicate reserved or system fields.)

- The flags record details about the last operation, such as:

  - Was the result zero?

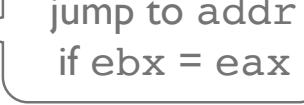 - Was the result positive?
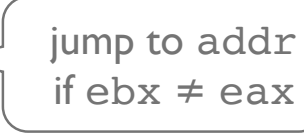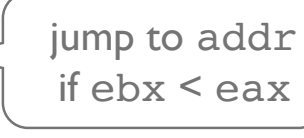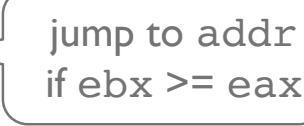
  - Did a carry occur?

  - etc...

Conditional jumps, jCC

We can test these flags in *conditional jump* instructions

```
jz  addr    (jump to addr if the zero flag is set)
jnz addr    (jump to addr if the zero flag is not set)
je  addr    (jump to addr if equal; same as jz)
jne addr    (jump to addr if not equal; same as jnz)
jl  addr    (jump to addr if less than)
jnl addr    (jump to addr if not less than)
jg  addr    (jump to addr if greater than)
jng addr    (jump to addr if not greater than)
...
```

(signed)

63

Examples

```
subl    %eax,%ebx
jz      addr
```
jump to addr
if ebx = eax

```
subl    %eax,%ebx
jnz     addr
```
jump to addr
if ebx ≠ eax

```
subl    %eax,%ebx
jl      addr
```
jump to addr
if ebx < eax

```
subl    %eax,%ebx
jnl     addr
```
jump to addr
if ebx >= eax

If the specified condition does not apply, then execution just continues with the next instruction ...

64

The compare instruction

- The `cmpX` instruction behaves like `subX` except that the result is not saved; only the flags are changed

- For example:
```
cmpl  %eax,%ebx
jl    addr
```

  will jump to `addr` if the value in `ebx` is less than the value in `eax`, but it will <u>not</u> change the values in either register

Other conditional instructions

- There are some other instructions that perform an action based on the conditional flags without the cost of a jump

- `setCC` reg8 sets the value in a specified 8-bit register to 0 or 1, based on the condition specified by `CC`:

```
cmpl    %ecx,%ebx    # set eax to 1 if
setl    %al          # ebx < ecx, or
movzbl  %al,%eax     # else to 0
```

- `cmovCC` src, dst copies data from the specified src to dst, but only if the condition specified by `CC` holds:

```
cmpl    %ebx,%eax    # set eax to the max of
cmovl   %ebx,%eax    # eax and ebx
```

condition code; no size suffix here!

IA32 instructions:
more arithmetic

Unary operations

- The following arithmetic operations have only one argument (which serves as both source and destination)

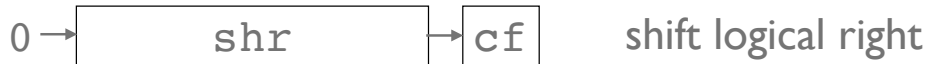|		
▶ `negX` (reg | mem)	negate
`notX` (reg | mem)	complement
`incX` (reg | mem)	increment
`decX` (reg | mem)	decrement

- Like the binary operators, these instructions also set the flags for subsequent testing

Bitwise shift operations

- Shift operations are handled using instructions of the form:

$$\texttt{op} \quad \texttt{count,} \; (\text{reg} \mid \text{mem})$$

cf ← shl/sal ← 0   shift (logical/arithmetic) left

0 → shr → cf   shift logical right

sar → cf   shift arithmetic right

- count is either a constant or else the `%cl` register.

- In all cases, the count value will be masked to 31 bits.

Example

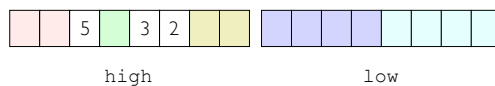- Given two 32 bit input values:

  - base:

  - limit: 0 0 0

  Each box is one nibble (4 bits), least significant bits on the right

- Calculate a 64 bit descriptor:

  5   3 2

  high   low

- (Needed for the calculation of "GDT entries")

Example

```
movl    base, %eax
movl    limit, %ebx

mov     %eax, %edx
shl     $16, %eax
mov     %bx, %ax
movl    %eax, low

shr     $16, %edx
mov     %edx, %ecx
andl    $0xff, %ecx
xorl    %ecx, %edx
shl     $16,%edx
orl     %ecx, %edx
andl    $0xf0000, %ebx
orl     %ebx, %edx
orl     $0x503200, %edx
movl    %edx, high
```


71

Bitwise rotate operations

• Rotate operations use the same instruction format:

rol   rotate left

rcl   rotate left with carry

ror   rotate right

rcr   rotate right with carry
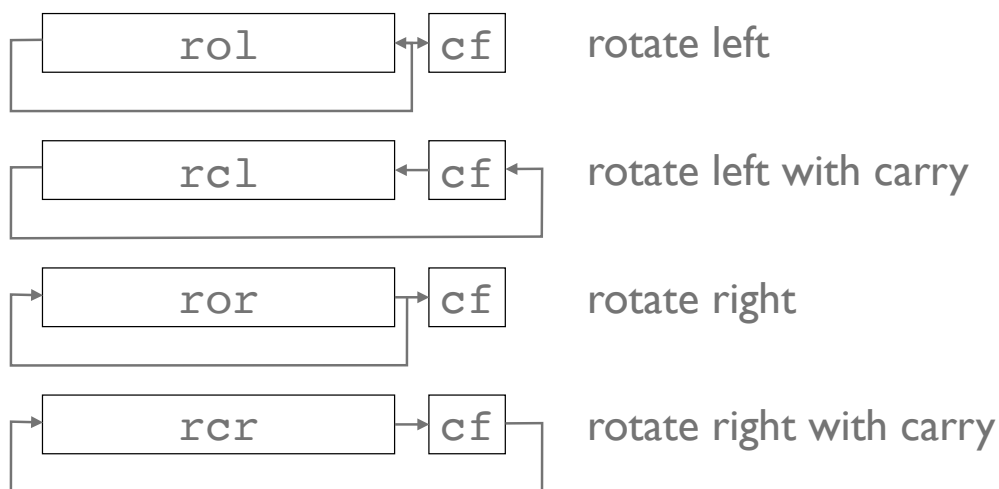
• [Aside: Curiously, "higher level" languages often include shift operators, but not rotates, even though the latter have more interesting/uniform behavior …]

72

Division

- Divide implicit destination (`edx:eax`) (a 64-bit quantity) by a specified argument with result in `eax` and remainder in `edx`

$$\texttt{idivl} \quad (\text{reg} \,|\, \text{mem})$$

- Often used in conjunction with the `cltd` instruction ("convert long to double", a.k.a. `cdq`), which converts a signed 32-bit value in `eax` into the corresponding signed 64-bit value in `edx:eax`.
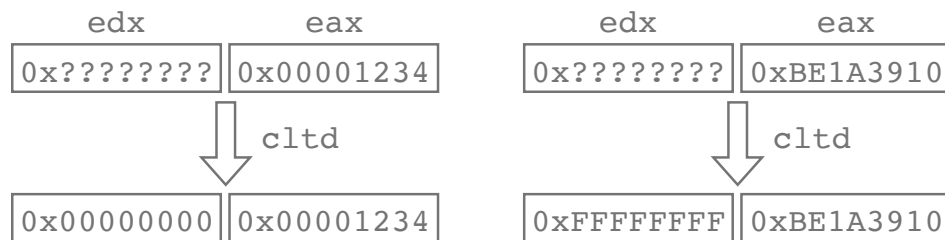
	edx	eax
	0x????????	0x00001234

⬇ cltd

| 0x00000000 | 0x00001234 |

	edx	eax
	0x????????	0xBE1A3910

⬇ cltd

| 0xFFFFFFFF | 0xBE1A3910 |

Example 1

Divide 4,660 (i.e., 0x1234) by 25:

```
▶    movl   $0x1234, %eax
     cltd
     movl   $25, %ecx
     idivl  %ecx
```

Results: eax = 0xBA (186)

edx = 0xA (10)

Sure enough: 186*25 + 10 = 4,660

Example 2

Divide -1,105,577,712 (i.e., 0xBE1A3910) by 256

```
   ▶      movl    $0xBE1A3910, %eax
          cltd
          movl    $256, %ecx
          idivl   %ecx
```

Results: eax = 0xFFBE1A3A (-4,318,662)

         edx = 0xffffff10 (-240)

Sure enough: -4,318,662 * 256 - 240 = -1,105,577,712

Complications of division

- Division produces multiple results: a quotient and a remainder

- Division uses special registers: we'd better not store any other values in eax or edx if there's a chance that a division instruction might be executed

- Doesn't set flags: requires separate tests, for example, to determine whether quotient or remainder was zero

- Division can raise an exception if the src is zero (or -1)

IA32 instructions:
using the stack

Stack

- The IA32 includes features that allow the programmer to use a region of memory as a simple stack:

  - the `esp` (stack pointer) register

  - special instructions like `push, pop, call, ret, ...`

- There is no obligation for the programmer to use these features, but it is often convenient to do so:

  - for temporary/scratch storage when a calculation needs more storage than the CPU registers can provide

  - to support calling and returning from functions

A typical memory layout

- A typical operating system reserves an area of scratch memory for each program, and sets the `esp` register to point to the end of this region when the program begins

program	data	stack

<center>esp ↑</center>

- The stack pointer moves

  - down (decreases) as values are pushed on to the stack

  - up (increases) as values are popped off of the stack

- So long as they never overlap, the data and stack areas can grow or shrink as necessary as the program runs

Stack operations

- Push a value onto the stack

$$\texttt{pushl} \quad (\text{reg} \mid \text{mem} \mid \text{immed})$$

- Pop a value of the stack

$$\texttt{popl} \quad (\text{reg} \mid \text{mem})$$

- Roughly speaking:

```
pushl src   =   subl $4, %esp;   movl src, (%esp)

popl dst    =   movl (%esp), dst;   addl $4, %esp
```

Spilling temporaries on the stack

- The stack is often used for saving the contents of a register on the stack ("spilling") so that the register can be used, temporarily, for some other reason

- For example:
  ```
  pushl  %eax
  pushl  %edx
  ... code that changes eax and/or edx ...
  popl   %edx
  popl   %eax
  ```
  > pop values in reverse order that was used to push them!

- Note that values on the stack can still be accessed, from memory, using `(%esp), 4(%esp), 8(%esp), 12(%esp),` ...

Call and return

- There is a special instruction for calling a function

  ```
  call addr       ≃            pushl $lab
                               jmp   addr
                      lab:...
  ```

- And a special instruction for returning from a function

  ```
  ret             ≃            popl   %eax
                               jmp    *%eax
  ```
  > assuming eax isn't being used for something else ...

  > special syntax: jump to the address given by the contents of eax

- In practice, additional instructions are often needed to deal with parameter passing, etc. ...

Functions
and the System V ABI

Implementing functions

• How do we pass arguments to a function?

• How does a function return a result?

• How do we handle local variables?

• In principle, especially in a bare metal setting, we can implement these features any way we like, using the basic tools that the IA32 instruction set provides

• But there are some existing standards we can follow, notably the "System V IA32 Application Binary Interface (ABI)":

  http://www.sco.com/developers/devspecs/abi386-4.pdf
  particularly Section 3-9

Stack frames

The code for any given function/procedure call runs in the context of a stack frame of the form:

	l_m	...	l_1	old	retn	a_1	...	a_n	...	
esp		...	-4	ebp	4	8	12	...		
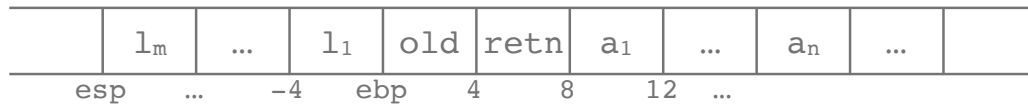
- Frame (base) pointer: ebp points to the stack frame; the caller's frame pointer is stored in old (i.e., (%ebp))
- Return address: retn is the return address
- Actual parameters: $a_1, ... , a_n$ are the function's arguments. We can access a_1 as 8(%ebp), etc...
- Local variables: $l_1, ... , l_m$ are the function's local variables. We can access l_1 as -4(%ebp), etc...

Building the stack frame … in the caller

-

		...	
	esp	ebp	

- The caller starts by pushing the arguments:

	a_1	...	a_n	...	
	esp		ebp		

- Then it executes a call instruction, which pushes the return address:

	retn	a_1	...	a_n	...	
	esp			ebp		

- … and jumps to the code for the callee …

Building the stack frame … in the callee

-

			retn	a_1	…	a_n	…	

esp   ebp

- The <u>callee</u> saves the old frame pointer, and sets a new value:

			old	retn	a_1	…	a_n	…	

ebp=esp   4   8   12   …
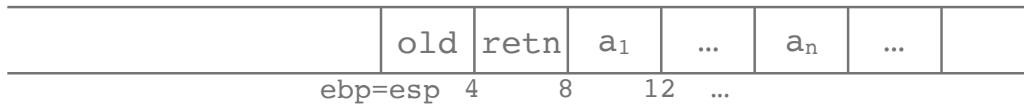
- Then it decrements the stack pointer to reserve space for any local variables:

	l_m	…	l_1	old	retn	a_1	…	a_n	…	

esp   …   −4   ebp   4   8   12   …
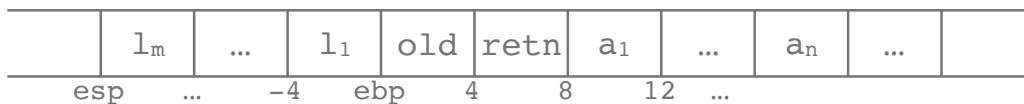
- … and now the callee can start work …

Function prologue

- The code that builds the stack frame at the start of a function body is called the <u>prologue</u>:

  - At the beginning of a function body, the parameters and return address have already been pushed on to the stack. We need to:
    ```
    pushl %ebp        # save old frame pointer
    movl  %esp, %ebp # and set new value
    ```

  - If local variables taking M bytes of storage are required, then we need to reserve space for them:
    ```
    subl   $M, %esp  # allocate space for
                     # locals (skip if M=0)
    ```

Function epilogue

- When a function completes, we must dismantle the stack frame and return the machine to the state it was in before the call. The code to do this is called the epilogue:

  - Running the previous process in reverse:
    ```
    movl    %ebp, %esp  # discard locals/temps
    popl    %ebp        # restore frame pointer
    ret                 # return to caller
    ```

  - The first two instructions here can be replaced with the more efficient, but otherwise equivalent `leave` instruction

Removing the parameters

- Once we return to the caller, the result of the function is in `eax`, but the parameters are still on the stack:

		a_1	...	a_n	...	

  esp   ebp

- We restore the stack pointer to its original value by adding on the number of bytes that are used by the parameters:

  ```
  addl   $N, %esp
  ```

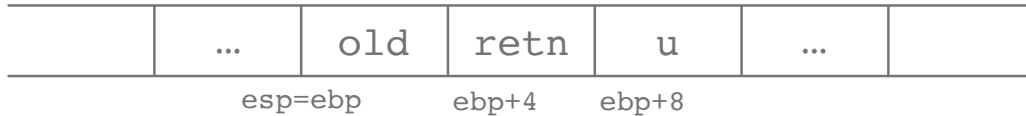- If no parameters were passed, then this step can be omitted

Example: a leaf function

```
int g(int u) {
  return u*u;
}
```

```
g: pushl %ebp
   movl  %esp, %ebp
   movl  8(%ebp), %eax
   imull %eax, %eax
   movl  %ebp, %esp
   popl  %ebp
   ret
```
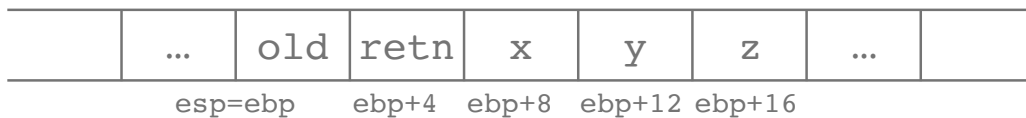
	…	old	retn	u	…	

esp=ebp   ebp+4   ebp+8

Example: multiple parameters + call

```
int f(int x,
      int y,
      int z) {
  return g(x+y);
}
```

```
f: pushl %ebp
   movl  %esp, %ebp
   movl  8(%ebp), %eax
   addl  12(%ebp), %eax
   pushl %eax
   call  g
   addl  $4, %esp
   movl  %ebp, %esp
   popl  %ebp
   ret
```

	…	old	retn	x	y	z	…	

esp=ebp   ebp+4   ebp+8   ebp+12 ebp+16

Example: spilling

```
int h(int x,              h: pushl %ebp
      int y,                 movl  %esp, %ebp
      int z) {               pushl 8(%ebp)
   return g(x)+g(y);         call  g
}                            addl  $4,%esp
                             pushl %eax -- spill
                             pushl 12(%ebp)
                             call  g
                             addl  $4, %esp
                             popl  %ecx -- unspill
                             addl  %ecx, %eax
                             movl  %ebp, %esp
                             popl  %ebp
                             ret
```

	spill	old	retn	x	y	z	
	esp	ebp	ebp+4	ebp+8	ebp+12	ebp+16	

93

Observations

• There is a four instruction overhead for each function that uses the frame pointer

 • Increases execution time

 • Prevents use of ebp as a general purpose register

• For larger functions, the four instruction overhead is less of an issue

• For small functions, we would prefer to inline rather than copy

• Nevertheless, it is common to produce code that doesn't use ebp as a frame pointer (e.g., -fomit-frame-pointer in gcc)

94

Caller and callee saves

We (System V) can designate some registers as:

- **caller saves** (eax, ecx, and edx)
  - can be freely used by the callee
  - the caller is responsible for saving (and later restoring) the value of a caller save register before a call

- **callee saves** (ebp, ebx, esi, and edi)
  - can be freely used by the caller
  - the callee is responsible for saving (and later restoring) the value of a callee saves register before using it to store temporary values

Revisiting the previous example: h

```
int h(int x,
      int y,
      int z) {
    return g(x)+g(y);
}
```

```
h: pushl %ebp
   movl  %esp, %ebp

   pushl 8(%ebp)  -- x
   call  g
   addl  $4,%esp
   pushl %eax -- spill
   pushl 12(%ebp) -- y
   call  g
   addl  $4, %esp
   popl  %ecx -- unspill
   addl  %ecx, %eax

   movl  %ebp, %esp
   popl  %ebp
   ret
```

instead of having the compiler save this value on the stack ...

Revisiting the previous example: h

```
int h(int x,
      int y,
      int z) {
    return g(x)+g(y);
}
```

```
h: pushl %ebp
   movl  %esp, %ebp

   pushl 8(%ebp)  -- x
   call  g
   addl  $4,%esp
   movl  %eax, %esi
   pushl 12(%ebp) -- y
   call  g
   addl  $4, %esp

   addl  %esi, %eax

   movl  %ebp, %esp
   popl  %ebp
   ret
```

… we can move it to a callee saves register, esi

g will preserve the value in esi, if necessary

so it will still contain the correct value here…

97

Revisiting the previous example: h

```
int h(int x,
      int y,
      int z) {
    return g(x)+g(y);
}
```

```
h: pushl %ebp
   movl  %esp, %ebp
   pushl %esi
   pushl 8(%ebp)  -- x
   call  g
   addl  $4,%esp
   movl  %eax, %esi
   pushl 12(%ebp) -- y
   call  g
   addl  $4, %esp

   addl  %esi, %eax
   popl  %esi
   movl  %ebp, %esp
   popl  %ebp
   ret
```

… now h has to save the value in register, esi

one save in h is better than one saves in each of two calls to g

empirically, more than 50% of calls are to leaf functions

98

Closing thoughts

Assembly "Language"?

• Highly imperative, primitive instructions, no expressions

• No high-level abstractions, but all the building blocks:

　• No arrays, records, variants, objects, closures, …

　• No loops, switch statements, functions, local variables, …

• Type System?

　• Values classified by size (e.g., 8 vs 32 bits) and storage class (e.g., memory, flag, integer register, floating point register, …)

　• Limited protection against common programming mistakes

　• Programmer has full control over data representation

Macros

macro name — *parameters* — *defaults*

```
        .macro  idtcalc handler, slot, dpl=0, type=IDT_INTR, seg=KERN_CS
        mov     $\seg, %ax                 # eax =   ? # seg
        shl     $16, %eax                  # eax = seg #   0
        movl    $\handler, %edx            # edx = hhi # hlo
        mov     %dx, %ax                   # eax = seg # hlo
        mov     $(0x8e00 | (\dpl<<13) | \type), %dx
        movl    %eax, idt + (     8*\slot)
        movl    %edx, idt + (4 + 8*\slot)
        .endm

initIDT:.irp    num, 0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,16,17,18,19
        idtcalc exc\num, slot=\num
        .endr
```

loop constructs

```
        idtcalc handler=syscall, slot=0x80, dpl=3
```

macro invocation

```
        lidt    idtptr
        ret
```

Macros support compile-time code generation:

- Macros are like "used-defined instructions"
- Also supported: conditional code generation, loops, …
- One way to compensate for language weaknesses?

101

Summary

- IA32 provides a very basic programming language:

  - A fixed set of registers
  - Instructions for moving and operating on data
  - Instructions for testing and control transfer

- In programming language terms:

  - Low-level, primitive instructions, loosely typed
  - No high-level abstractions, but all the building blocks
  - Very close to the metal, low-level control, "predictable" performance

- Let's write some programs!

102